



# Intel® Intelligent Power Node Manager 2.0

External Interface Specification Using IPMI

---

*April 2012*

Revision 1.3

**Intel Confidential**

## Disclaimers

**Notice:** This document contains information on products in the design phase of development. The information here is subject to change without notice. Do not finalize a design with this information.

INFORMATION IN THIS DOCUMENT IS PROVIDED IN CONNECTION WITH INTEL PRODUCTS. NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. EXCEPT AS PROVIDED IN INTEL'S TERMS AND CONDITIONS OF SALE FOR SUCH PRODUCTS, INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO SALE AND/OR USE OF INTEL PRODUCTS INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT OR OTHER INTELLECTUAL PROPERTY RIGHT.

A "Mission Critical Application" is any application in which failure of the Intel Product could result, directly or indirectly, in personal injury or death. SHOULD YOU PURCHASE OR USE INTEL'S PRODUCTS FOR ANY SUCH MISSION CRITICAL APPLICATION, YOU SHALL INDEMNIFY AND HOLD INTEL AND ITS SUBSIDIARIES, SUBCONTRACTORS AND AFFILIATES, AND THE DIRECTORS, OFFICERS, AND EMPLOYEES OF EACH, HARMLESS AGAINST ALL CLAIMS COSTS, DAMAGES, AND EXPENSES AND REASONABLE ATTORNEYS' FEES ARISING OUT OF, DIRECTLY OR INDIRECTLY, ANY CLAIM OF PRODUCT LIABILITY, PERSONAL INJURY, OR DEATH ARISING IN ANY WAY OUT OF SUCH MISSION CRITICAL APPLICATION, WHETHER OR NOT INTEL OR ITS SUBCONTRACTOR WAS NEGLIGENT IN THE DESIGN, MANUFACTURE, OR WARNING OF THE INTEL PRODUCT OR ANY OF ITS PARTS.

Intel may make changes to specifications and product descriptions at any time, without notice. Designers must not rely on the absence or characteristics of any features or instructions marked "reserved" or "undefined." Intel reserves these for future definition and shall have no responsibility whatsoever for conflicts or incompatibilities arising from future changes to them.

Intel software products are copyrighted by and shall remain the property of Intel Corporation. Use, duplication, or disclosure is subject to restrictions stated in Intel's Software License Agreement, or in the case of software delivered to the government, in accordance with the software license agreement as defined in FAR 52.227-7013.

The products described in this document may contain design defects or errors known as errata which may cause the product to deviate from published specifications. Current characterized errata are available on request.

The code names presented in this document are only for use by Intel to identify products, technologies, or services in development, that have not been made commercially available to the public, i.e., announced, launched or shipped. They are not "commercial" names for products or services and are not intended to function as trademarks.

Contact your local Intel sales office or your distributor to obtain the latest specifications and before placing your product order.

Copies of documents which have an order number and are referenced in this document, or other Intel literature, may be obtained by calling 1-800-548-4725, or by visiting [Intel's Web Site](#).

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States and other countries.

Copyright © 2010-2012, Intel Corporation. All rights reserved.

\* Other brands and names may be claimed as the property of others.



# Contents

|        |  |    |
|--------|--|----|
| 1.     | Introduction.....  | 6  |
| 1.1    | Purpose of this Document.....  | 6  |
| 1.2    | Scope .....  | 6  |
| 1.3    | General Conventions .....  | 6  |
| 1.4    | System States and Power Management.....                              | 6  |
| 1.5    | Reference Documents .....  | 6  |
| 2.     | Intel® Management Engine IPMI Interface.....                         | 8  |
| 2.1    | SEL Device Commands.....   | 8  |
| 2.2    | IPMI Device “Global” Commands.....                                   | 8  |
| 2.3    | Sensor Device Commands.....  | 11 |
| 2.4    | ME Firmware Debug Event.....   | 13 |
| 2.4.1  | Debug SEL Entry Definition .....                                     | 13 |
| 2.5    | IPMI OEM Device Commands.....  | 14 |
| 2.6    | IPMI Device “Global” Sensors.....                                    | 15 |
| 2.6.1  | ME Power State sensor .....  | 16 |
| 2.6.2  | ME Firmware Health sensor .....                                      | 16 |
| 2.6.3  | PCH Thermal Sensor.....  | 16 |
| 2.6.4  | Event Messages Definition.....                                       | 17 |
| 2.7    | IPMI OEM ME Firmware Update Commands .....                           | 19 |
| 2.7.1  | Online Update Flow .....   | 23 |
| 2.7.2  | Backward Compatibility Mode .....                                    | 25 |
| 2.7.3  | IPMI OEM Image Inventory Command .....                               | 29 |
| 2.7.4  | Optimizing Online Upgrade Performance .....                          | 30 |
| 2.8    | IPMI Commands Supported by Recovery Boot Loader .....                | 30 |
| 2.9    | IPMI OEM PECI Proxy Commands.....                                    | 30 |
| 2.10   | IPMI OEM PECI Proxy Sensors .....                                    | 38 |
| 2.10.1 | CPU Thermal Status Sensors .....                                     | 40 |
| 2.10.2 | CPU Thermal Control Circuit Activation Sensors.....                  | 40 |
| 2.10.3 | CPU T-Control Sensors.....   | 40 |
| 2.10.4 | CPU Tjmax Sensors .....  | 40 |
| 2.10.5 | Memory Throttling Status Sensors.....                                | 40 |
| 2.11   | IPMI Standard Completion Codes .....                                 | 41 |
| 2.12   | Generic Event/Reading Type Codes.....                                | 42 |
| 3.     | ME NM IPMI Interface.....  | 43 |
| 3.1    | External Node Manager Configuration and Control Commands .....       | 43 |
| 3.1.1  | Get Node Manager Statistics Current Value Field .....                | 60 |
| 3.1.2  | Get Node Manager Statistics Minimum Value Field.....                 | 60 |
| 3.1.3  | Get Node Manager Statistics Maximum Value Field .....                | 61 |
| 3.1.4  | Get Node Manager Statistics Average Value Field .....                | 61 |
| 3.2    | Local Platform Node Manager Configuration and Control Commands ..... | 61 |
| 3.3    | External DCMI Power Management Commands.....                         | 67 |
| 3.4    | IPMI Sensors Implemented by Node Manager FW.....                     | 71 |
| 3.4.1  | Typical Platform power consumption in Sx state.....                  | 74 |
| 3.4.2  | NM Exception Event Sensor .....                                      | 74 |
| 3.4.3  | NM Alert Threshold Exceeded Sensor .....                             | 74 |
| 3.4.4  | NM Health Sensor.....  | 74 |
| 3.4.5  | NM Operational Capabilities Change Sensor .....                      | 75 |
| 3.4.6  | PSU Status Sensors .....   | 75 |
| 3.4.7  | Event Messages Definition.....                                       | 76 |
| 3.5    | Error Conditions.....  | 80 |

|       |   |    |
|-------|---|----|
| 4.    | BMC IPMI Interface .....                              | 82 |
| 4.1   | IPMI Device "Global" Commands .....                   | 82 |
| 4.2   | Sensor Device Commands .....                          | 82 |
| 4.3   | Alert Immediate Command .....                         | 83 |
| 4.4   | OEM Commands Implemented by BMC .....                 | 83 |
| 4.4.1 | Power Consumption Readings .....                      | 83 |
| 4.4.2 | Inlet Air Temperature readings .....                  | 84 |
| 4.4.3 | OEM Management Engine Power State Change .....        | 84 |
| 4.4.4 | OEM Command Definition .....                          | 85 |
| 4.4.5 | Summary of Options .....                              | 86 |
| 4.5   | BMC requirements for Node Manager Discovery .....     | 86 |
| 4.6   | Alerts .....  | 87 |
| 4.7   | Command Passing via BMC .....                         | 88 |
| 4.8   | IPMB Reset Scenarios .....                            | 90 |
| 5.    | IPMI OEM Command Summary .....                        | 91 |
| 6.    | Appendix NM2.0 Platform Changes .....                 | 93 |
| 6.1   | IPMI Commands Not Supported on NM2.0 Platform .....   | 93 |
| 6.2   | IPMI Sensors Implemented by Node Manager FW .....     | 93 |
| 6.2.1 | Event Generation Control .....                        | 93 |
| 6.3   | IPMI Platform Event Messages Generated by NM FW ..... | 96 |



## Revision History

| Revision Number | Description   | Date          |
|-----------------|---|---------------|
| 0.7             | Initial release   | October 2009  |
| 0.8             | Updates to Specification  | January 2010  |
| 0.9             | Modification after external review  | June 2010     |
| 1.0             | Cleanup before final release  | November 2010 |
| 1.1             | New IPMI message Get ME Factory Presets Signature defined<br>Several editorial changes and clarification. | August 2011   |
| 1.2             | -Changes for Denlow<br>-DCMI version corrected<br>-Added clarification to statistics                      | March 2012    |
| 1.3             | PSU Status Sensor clarification   | April 2012    |

§

# 1. Introduction

---

## 1.1 Purpose of this Document

This document is prepared to assist BMC vendors and external management software vendor in supporting Intel® Intelligent Power Node Manager 2.0. It documents the IPMI commands that can be sent to the Intel® Management Engine (Intel® ME) present in the Cougar Point chipset and in the Patsburg chipset. This document describes also IPMI sensors implemented in order to ensure the correct and reliable operation of the platform.

## 1.2 Scope

This document describes the details on the IPMI commands used by Node Manager firmware components running on an Intel® platform.

## 1.3 General Conventions

By default all the values occupying more than one byte are LS Byte first encoded if not specified differently in their descriptions.

## 1.4 System States and Power Management

| Acronym or Term | Definition   |
|-----------------|--|
| S0              | A system state where power is applied to all HW devices and system is running normally.                    |
| S1, S2, S3      | A system state where the host CPU is not running however power is connected to the memory system.          |
| S4              | A system state where the host CPU and memory is not active   |
| S5              | A system state where all power to the host system is off however power cord is still connected.            |
| Sx              | All S states that are different than S0/S1   |
| OS Hibernate    | OS state where the OS state is saved on the hard drive.  |
| Standby         | OS state where the OS state is saved on memory and resumed from the memory when mouse/keyboard is clicked. |
| Shut Down       | All power is off for the host machine however the power cord is still connected.                           |

## 1.5 Reference Documents

| Ref    | Document Name   | File/Location   |
|--------|---|---|
| [Addr] | IPMB v1.0 Address Allocation, 1998.   | <a href="http://www.intel.com/design/servers/ipmi/spec.htm">http://www.intel.com/design/servers/ipmi/spec.htm</a> |
| [IPMI] | Intelligent Platform Management Interface Specification, version 2.0, 2004. | <a href="http://www.intel.com/design/servers/ipmi/spec.htm">http://www.intel.com/design/servers/ipmi/spec.htm</a> |



| Ref    | Document Name  | File/Location   |
|--------|--|---|
| [Addr] | IPMB v1.0 Address Allocation, 1998.  | <a href="http://www.intel.com/design/servers/ipmi/spec.htm">http://www.intel.com/design/servers/ipmi/spec.htm</a>                           |
| [IPMB] | Intelligent Platform Management Bus Specification, version 1.0, 1999.          | <a href="http://www.intel.com/design/servers/ipmi/spec.htm">http://www.intel.com/design/servers/ipmi/spec.htm</a>                           |
| [PET]  | IPMI Platform Event Trap Format Specification, version 1.0, 1999.              | <a href="http://www.intel.com/design/servers/ipmi/spec.htm">http://www.intel.com/design/servers/ipmi/spec.htm</a>                           |
| [DCMI] | Data Center Management Interface Specification version 1.5, draft revision 0.7 | <a href="http://www.intel.com/technology/product/dcmi/specification.htm">http://www.intel.com/technology/product/dcmi/specification.htm</a> |
| [EMTS] | TBD  | TBD   |

## 2. Intel® Management Engine IPMI Interface

This section contains IPMI commands and sensor devices provided by Intel® ME. BMC shall use these commands and sensors to control Node Manager firmware running on Intel® ME. All commands listed in chapter 2 are mandatory and will be implemented by NM firmware.

### 2.1 SEL Device Commands

| Net Function = Storage (0Ah) |              |   |  |
|------------------------------|--------------|---|--|
| Code                         | Command      | Request, Response Data  | Description  |
| 48h                          | Get SEL Time | <b>Request</b><br>None  | This is standard IPMI 2.0 command.<br><br>Node Manager firmware responds to this command returning internal clock value. |
|                              |              | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11)<br>Bytes 2:5 - Present Timestamp value. |  |

### 2.2 IPMI Device “Global” Commands

| Net Function = App (06h) |            |   |   |
|--------------------------|------------|---|---|
| Code                     | Command    | Request, Response Data  | Description   |
| 02h                      | Cold Reset | <b>Request</b><br>None  | This is standard IPMI 2.0 command.<br><br>Reboots ME without resetting host platform. |
|                          |            | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11) |   |

| Net Function = App (06h) |               |  |                                    |
|--------------------------|---------------|--|------------------------------------|
| Code                     | Command       | Request, Response Data   | Description                        |
| 01h                      | Get Device ID | <b>Request</b><br>None   | This is standard IPMI 2.0 command. |
|                          |               | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion |                                    |





| Net Function = App (06h) |         |   |             |
|--------------------------|---------|---|-------------|
| Code                     | Command | Request, Response Data  | Description |
|                          |         | <p>Codes are shown in section 2.11)</p> <p>Byte 2 – Device ID</p> <p>=50h - Intel® Management Engine (Intel® ME)</p> <p>Byte 3 – Device Revision</p> <p>[7] = 1 - For DM and DNM firmware SKU device provides Device SDRs</p> <p>= 0 - For NM, Silicon Enabling and Recovery boot-loader device does not provide Device SDRs</p> <p>[6:4] reserved. Return as 000b.</p> <p>[3:0] Device Revision, binary encoded. = 1</p> <p>Byte 4 Firmware Revision 1</p> <p>[7] Device available:</p> <p>= 0 - normal operation,</p> <p>= 1 - device firmware update or self-initialization in progress or firmware in the recovery boot-loader mode</p> <p>[6:0] Major Firmware Revision, binary encoded = 2</p> <p>Byte 5 - Firmware Revision 2: Minor Firmware Revision. BCD encoded.</p> <p>Byte 6 - IPMI Version. Holds IPMI Command Specification Version. BCD encoded.</p> <p>00h = reserved.</p> <p>Bits 7:4 hold the Least Significant digit of the revision, while</p> <p>Bits 3:0 hold the most significant digits.</p> <p>=02h to indicate revision 2.0.</p> <p>Byte 7 - Additional Device Support. Lists the IPMI 'logical device' commands and functions that the controller supports that are in addition to the mandatory IPM and Application commands.</p> <p>For NM and Silicon enabling SKU byte 7 is set to:</p> <p>[7] = 0 Not a chassis Device</p> <p>[6] = 0 Not a Bridge</p> <p>[5] = 1 IPMB Event Generator</p> <p>[4] = 0 Not a IPMB Event Receiver</p> <p>[3] = 0 Not a FRU Inventory Device</p> <p>[2] = 0 Not a SEL Device</p> <p>[1] = 0 Not a SDR Repository Device</p> <p>[0] = 1 Sensor Device</p> <p>For DM and DNM SKU byte 7 is set to:</p> <p>[7] = 1 Chassis device Supported</p> <p>[6] = 0 Bridge not supported</p> <p>[5] = 0 IPMB event generator not supported</p> <p>[4] = 0 IPMB event receiver not supported</p> <p>[3] = 1 FRU inventory device supported</p> <p>[2] = 1 SEL device supported</p> <p>[1] = 1 SDR repository device supported</p> <p>[0] = 1 Sensor device supported</p> <p>If Recovery boot-loader image is loaded byte 7 is set to:</p> <p>[7] = 0 Not a chassis Device</p> <p>[6] = 0 Not a Bridge</p> <p>[5] = 1 IPMB Event Generator</p> <p>[4] = 0 Not a IPMB Event Receiver</p> <p>[3] = 0 Not a FRU Inventory Device</p> |             |

| Net Function = App (06h) |                       |  |                                    |
|--------------------------|-----------------------|--|------------------------------------|
| Code                     | Command               | Request, Response Data   | Description                        |
|                          |                       | <p>[2] = 0 Not a SEL Device<br/> [1] = 0 Not a SDR Repository Device<br/> [0] = 0 Not a Sensor Device</p> <p>Bytes 8:10 – Manufacturer ID = 57h, 01h, 00h.</p> <p>Byte 11 – Product ID Minor Version<br/> = 00h – Tylersburg platform<br/> = 01h – Bromolow platform<br/> = 02h – Romley platform<br/> = 03h – Denlow platform<br/> Byte 12 – Product ID Major Version = 0Bh</p> <p>Bytes 13:16 – Auxiliary Firmware Revision Information<br/> Byte 13 – (Binary encoded) Implemented version of Firmware<br/> [7:4] Implemented DCMI version<br/> =0 – DCMI not implemented/enabled<br/> =1 – DCMI Revision 1.0<br/> =2 – DCMI Revision 1.1<br/> =3 – DCMI Revision 1.5</p> <p>[3:0] Implemented Node Manager IPMI interface version<br/> =0 – NM not implemented/enabled<br/> =1 – NM Revision 1.5<br/> =2 – NM Revision 2.0</p> <p>Note: For Silicon Enabling SKU byte 13 will be set to all zeros.</p> <p>Byte 14 –Firmware build number BCD encoded = A.B<br/> Byte 15 –The last digit of firmware build number and patch number BCD encoded = C.PATCH<br/> Byte 16 – Image flags<br/> [7:2] – reserved. Return as 000000b<br/> [1:0] – image type<br/> = 00b – recovery image<br/> = 01b – operational image 1<br/> = 10b – operational image 2<br/> = 11b – unspecified: flash error indication</p> <p>Note: Full version number is: "Major Firmware Revision. Minor Firmware Revision.ABC.PATCH" where ABC is Firmware build number.</p> |                                    |
| 04h                      | Get Self Test Results | <b>Request</b><br>None   | This is standard IPMI 2.0 command. |
|                          |                       | <b>Response</b><br><p>Byte 1 – Completion Code<br/> =00h – Success (Remaining standard Completion Codes are shown in section 2.11)<br/> =D5h – Returned if self tests is not finished yet.</p> <p>Byte 2 –<br/> =55h – No error. All Self Tests Passed.<br/> =56h – Self Test function not implemented in this controller.<br/> =57h – Corrupted or inaccessible data or devices.<br/> =58h – Fatal hardware error (system should consider BMC inoperative). This will indicate that the controller hardware (including associated devices such as sensor hardware or RAM) may need to be repaired or replaced or multiple software exceptions occurred.</p>   |                                    |



| Net Function = App (06h) |         |  |             |
|--------------------------|---------|--|-------------|
| Code                     | Command | Request, Response Data   | Description |
|                          |         | <p>=80h – PSU Monitoring service error see Byte 3 for error description only if ME Firmware directly monitors PMBUS PSU. PMBUS PSU Monitoring service will return the current status of all defined PSUs on 'Get Self Test Results' call.</p> <p><b>Note:</b> The error code is continuously updated in runtime in S0/S1 host power states by the Monitoring Service. Additionally, the test will be performed in any host power state if Manufacturing Test On Command is issued.</p> <p>=81h – Firmware entered Recovery boot-loader mode</p> <p>=FFh – reserved.</p> <p>Byte 3 –</p> <p>For byte 2 = 55h, 56h, FFh:</p> <p>=00h</p> <p>For byte 2 = 58h: Exception type.</p> <p>For byte 2 = 57h: Self-test error bit field.</p> <p>[7] – Factory Presets checksum error.</p> <p>[6] – SDR access error.</p> <p>[5] – FRU access error.</p> <p>[4] – SEL access error.</p> <p>[3] – PIA access error.</p> <p>[2] – SDR repository empty.</p> <p>[1] – Firmware boot error.</p> <p>[0] – Reserved.</p> <p>For byte 2 = 80h: PSU monitoring error bit field, where each bit corresponds to one of the PSU's in order. If bit[N] is set to 1b PSU[N] is not responding. PSU order is set by factory presets.</p> <p>For byte 2 = 81h: byte 3 contains reason</p> <p>= 00h – recovery entered due to recovery jumper being asserted</p> <p>= 01h – recovery entered due to Security strap override jumper being asserted</p> <p>= 02h – recovery mode entered by IPMI command "Force ME Recovery"</p> |             |

## 2.3 Sensor Device Commands

| Net Function = S/E (04h) |                    |   |  |
|--------------------------|--------------------|---|--|
| Code                     | Command            | Request, Response Data  | Description  |
| 00h                      | Set Event Receiver | <p><b>Request</b></p> <p>Byte 1 – Event Receiver Slave Address.</p> <p>=0FFh disables Event Message Generation, Otherwise:</p> <p>[7:1] – IPMB (I<sup>2</sup>C) Slave Address</p> <p>[0] – always 0b when [7:1] hold I<sup>2</sup>C slave address</p> <p>Byte 2 – [7:2] – reserved. Write as 000000b.</p> <p>[1:0] – Event Receiver LUN</p> <p><i>Note: Depending on the Factory preset: "Default Event Receiver Address":</i></p> <p>- if 00h is set in the factory presets ME Firmware will not send any event until Set Event Receiver command will be sent by BMC on platform startup from G3 or on Global Platform Reset (see definition in PCH EDS)</p> | <p><i>Note: Value set by Set Event Receiver command is not stored in the persistent storage so it should be send on any platform startup from G3 and on Global Platform Reset (see definition in PCH EDS).</i></p> |

| Net Function = S/E (04h) |                    |   |             |
|--------------------------|--------------------|---|-------------|
| Code                     | Command            | Request, Response Data  | Description |
|                          |                    | - if 20h is set in the factory presets ME Firmware will not wait for BMC to send Set Event Receiver command before starting events generation. BMC can still use the command to regenerate all the active events."  |             |
|                          |                    | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11)   |             |
| 01h                      | Get Event Receiver | <b>Request</b><br>None  |             |
|                          |                    | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11)<br>Byte 2 – Event Receiver Slave Address. 0FFh indicates Event Message Generation has been disabled. Otherwise:<br>[7:1] – IPMB (I <sup>2</sup> C) Slave Address.<br>[0] – Always 0b when [7:1] hold I <sup>2</sup> C slave address.<br>Byte 3 –<br>[7:2] – Reserved. Return as 000000b.<br>[1:0] – Event Receiver LUN. |             |

| Net Function = S/E (04h) |                         |                                    |                                    |
|--------------------------|-------------------------|------------------------------------|------------------------------------|
| Code                     | Command                 | Request, Response Data             | Description                        |
| 26h                      | Set Sensor Thresholds   | For command description see [IPMI] | This is standard IPMI 2.0 command. |
| 27h                      | Get Sensor Thresholds   | For command description see [IPMI] | This is standard IPMI 2.0 command. |
| 28h                      | Set Sensor Event Enable | For command description see [IPMI] | This is standard IPMI 2.0 command. |
| 29h                      | Get Sensor Event Enable | For command description see [IPMI] | This is standard IPMI 2.0 command. |
| 2Ah                      | Re-arm Sensor Events    | For command description see [IPMI] | This is standard IPMI 2.0 command. |
| 2Bh                      | Get Sensor Event Status | For command description see [IPMI] | This is standard IPMI 2.0 command. |



| Net Function = S/E (04h) |                    |                                    |  |
|--------------------------|--------------------|------------------------------------|--|
| Code                     | Command            | Request, Response Data             | Description  |
| 2Dh                      | Get Sensor Reading | For command description see [IPMI] | This is standard IPMI 2.0 command.<br><br><i>Note: If sensor scanning is disabled for example using Flash Image Tool Get Sensor Reading command will return:</i><br>- Completion Code 00h<br>- last reading or 00h if there was not reading<br>- and bit [6] of byte 3 set to 1. |

## 2.4 ME Firmware Debug Event

After recovery from an ME system error ME Firmware generates “Add SEL Entry” command with debug information about the error. Purpose of this message is solely in field debugging, it is not intended to replace health sensor events. It is recommended that BMC shall support storing the SEL records passed in the “Add SEL Entry” command. By assumption debug event does not contain any sensitive information. Generation of this event may be disabled by setting Enable Debug SEL Entries parameter in Flash Image Tool. It is recommended to disable debug event generation in the firmware loaded on the platforms shipped to the end customers.

Event is repeated every 5s until NM receives response with Success Completion Code.

### 2.4.1 Debug SEL Entry Definition

Debugging information in the message is intended to be used only by Intel Corporation engineering team. Since internal structure definition may be subject of changes any support request should contain both message content and firmware revision (it is not embedded in message due to space constraints).

| Net Function = Storage(0Ah) |               |  |  |
|-----------------------------|---------------|--|--|
| Code                        | Command       | Request, Response Data   | Description  |
| 44h                         | Add SEL Entry | <b>Request</b><br>Byte 1:2 – Record ID<br>=1..N<br><br>Byte 3 – Record Type<br>=EEh – OEM non-timestamped.<br><br>Byte 4:16 – Debugging information. | This type of record contains extended information about errors detected by ME. Purpose of it is solely in field debug. Generation of this event may be disabled by flash image configuration tool. |
|                             |               | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11)                        |  |

## 2.5 IPMI OEM Device Commands

| Net Function = 2Eh-2Fh |                    |  |   |
|------------------------|--------------------|--|---|
| Code                   | Command            | Request, Response Data   | Description   |
| DCh                    | Set ME Power State | <b>Request</b><br><br>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first.<br><br>Byte 4 – New power state:<br>=00 – Turn off ME power.<br><br><i>Note: After turning-off ME will wake-up on one of the following events:</i> <ul style="list-style-type: none"><li>Automatically when Host CPU goes to S0</li><li>Write to any SMT device on the I2C address defined in softstraps. In that case only ME wakes-up.</li></ul><br><i>ME wake-up SMBUS message definition:</i> <ul style="list-style-type: none"><li>SMBUS Host Address = as defined in for I2C address in softstraps.</li><li>Device Address = 01h.</li><li>Data Byte Low = 01h.</li><li>Data Byte High = 03h .</li></ul> |   |
|                        |                    | <b>Response</b><br><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11)<br><i>Note: The ME may be turned on and off by the BMC in Sx Host CPU states. In S0/S1 Host CPU state command will return D5h error code.</i><br><br>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first.  |   |
| Net Function = 2Eh-2Fh |                    |  |   |
| DFh                    | Force ME Recovery  | <b>Request</b><br><br>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first.<br><br>Byte 4 – Command<br>=01h Restart using Recovery Firmware (ME FW configuration is not restored to factory defaults)<br>=02h Restore Factory Default Variable values and restart the ME FW   | With parameter Command = 01h ME FW resets and prevents loading the regular operational FW code – ME FW stops in Recovery Boot Loader (note: in DM/DNM FW, ME FW becomes inaccessible using DCMI over HECI or LAN protocols; the only access is from host – e.g. to perform FW update from BIOS). After issuing this command the direct FW update is available even after End-of-POST reception.<br><br>With parameter Command = 02h ME FW restores all variables stored in non volatile memory to its factory defaults (as set using FIT in the factory). This requires two ME FW resets: ME FW first resets and temporarily stays in the recovery boot loader code while the factory defaults are restored, and then another FW reset happens and ME FW comes back with the factory settings |
|                        |                    | <b>Response</b><br><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11)<br>81h – Unsupported Command parameter value in the Byte 4 of the request.<br><br>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first.   |   |



| Net Function = 2Eh-2Fh |                                  |  |   |
|------------------------|----------------------------------|--|---|
| Code                   | Command                          | Request, Response Data   | Description   |
|                        |                                  |  | restored.<br><br>Note: ME FW will always stay in recovery boot loader for other reasons, such as recovery jumper asserted or security strap override asserted.  |
| E0h                    | Get ME Factory Presets Signature | <b>Request</b><br><br>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first.   | The signature length is ME Firmware implementation specific and fixed for given ME firmware release, but may be different in subsequent releases.<br><br>Bytes 5:N exist in response frame only if Completion Code (Byte 1) = 00h |
|                        |                                  | <b>Response</b><br><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11)<br><br>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first.<br><br>Byte 5:8 – Vendor Label<br>4 bytes long data assigned by platform vendor at ME image creation time.<br><br>Byte 9:N – Factory defaults signature of length between 1..32 bytes. |   |

## 2.6 IPMI Device “Global” Sensors

Two IPMI sensors are always exposed by ME FW:

- ME Power State
- ME Firmware Health

Detailed description of these sensors is provided in the subsections.

Reading Availability column specifies when the sensor reading is available:

- A – always when ME is On
- H – when HOST CPU is On
- O – after reception of END\_OF\_POST notification
- E – No reading available (Event Only)

Defaults Configurable in FIT column defines whether the default configuration of the sensors can be set using Flash Image Tool. The default configuration includes:

- Thresholds
- Event Enable Mask
- Scanning Periods
- Scanning Enable Flag
- Per-sensor Event Enable Flag

| Sensor Number | Description                 | Firmware SKU Availability | Reading Availability | Defaults Configurable in Flash Image Tool | Notes   |
|---------------|-----------------------------|---------------------------|----------------------|---|---|
| 22            | ME Power State <sup>1</sup> | A E Y                     |                      | es  | OEM Event only sensor<br>"Command illegal for specified sensor or record type (CDh)" error code is returned in response to the following commands: Get Sensor Reading, Set/Get Sensor Thresholds, Re-Arm Sensor Events, Set/Get Sensor Event Enable |
| 23            | ME Firmware Health          | A                         | E                    | No  | OEM Event only sensor<br>"Command illegal for specified sensor or record type (CDh)" error code is returned in response to the following commands: Get Sensor Reading, Set/Get Sensor Thresholds, Re-Arm Sensor Events, Set/Get Sensor Event Enable |
| 8             | PCH Thermal Sensor          | A                         | A                    | Yes                                       | Temperature threshold sensor  |

### 2.6.1 ME Power State sensor

The sensor is used to send Platform Event messages to BMC when ME power state is changing. The sensor uses Generic Event Reading code 0Ah. It supports only the offsets:

00h – Transition to Running – ME is started

02h – Transition to Power Off – ME is to be powered down

Note: Optionally, instead or in addition to the event, ME Firmware may send an IPMI command with power state change notification as defined in section 4.4.3. Using Factory presets OEM may choose to use an event or OEM command or both.

### 2.6.2 ME Firmware Health sensor

The sensor is used in Platform Event messages to BMC containing health information including but not limited to FW Upgrade and application errors.

### 2.6.3 PCH Thermal Sensor

This sensor provides PCH die temperature sensor value in Celsius degrees.

---

<sup>1</sup> *Optionally, instead of the event NM Firmware may send an IPMI command with information as defined above. Using Factory presets OEM may choose to use an event and/or OEM command. User can decide to disable the sensor and not disable the OEM command generation.*





## 2.6.4 Event Messages Definition

| Net Function = S/E (04h) |  |   |  |
|--------------------------|--|---|--|
| Code                     | Command  | Request, Response Data  | Description  |
| 02h                      | Platform Event Message<br>ME Power State           | <b>Request</b><br>Byte 1 – EvMRev<br>=04h (IPMI2.0 format)<br>Byte 2 – Sensor Type<br>=16h (microcontroller)<br>Byte 3 – Sensor Number<br>=22 – ME Power State<br>Byte 4 – Event Dir   Event Type<br>[7] – Event Dir<br>=0 – Assertion Event.<br>[6:0] – Event Type<br>=0Ah – Availability Status.<br>Byte 5 – Event Data 1<br>[7:6]=00b – unspecified byte 2<br>[5:4]=00b – unspecified byte 3<br>[3:0] – offset from event type code:<br>=00h – Transition to Running<br>=02h – Transition to Power Off   | <i>Note: OEM can select using Flash Image Tool that the notification is sent using OEM command instead of Platform Event Message.</i><br><br>See section OEM Management Engine Power State Change for definition of the OEM command.                             |
|                          |  | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11)   |  |
| 02h                      | Platform Event Message<br>ME Firmware Health Event | <b>Request</b><br>Byte 1 – EvMRev<br>=04h (IPMI2.0 format)<br>Byte 2 – Sensor Type<br>=DCh (OEM)<br>Byte 3 – Sensor Number<br>=23 – ME Firmware Health<br>Byte 4 – Event Dir   Event Type<br>[7] – Event Dir<br>=0 – Assertion Event.<br>[6:0] – Event Type<br>=75h (OEM)<br>Byte 5 – Event Data 1<br>[7:6]=10b – OEM code in byte 2<br>[5:4]=10b – OEM code in byte 3<br>[3:0] – Health Event Type<br>=00h – Firmware Status.<br>Byte 6 – Event Data 2<br>=0 – Recovery GPIO forced . Recovery Image loaded due to recovery MGPIO pin asserted. Pin number is configurable in factory presets, Default recovery pin is MGPIO1.<br><b>Repair action:</b> Deassert MGPIO1 and reset the ME<br>=1 – Image execution failed. Recovery Image or backup operational image loaded because operational image is corrupted. This may be either caused by Flash device corruption or failed upgrade procedure. | This platform event provides a run-time status of general Firmware status. Recovery from the errors may require ME reset or even FW upgrade or HW repair if the error is persistent.<br><br><i>Note: This sensor cannot be disabled using Factory Image Tool</i> |

| Net Function = S/E (04h) |         |  |             |
|--------------------------|---------|--|-------------|
| Code                     | Command | Request, Response Data   | Description |
|                          |         | <p><b>Repair action:</b> Either the Flash device must be replaced (if error is persistent) or the upgrade procedure must be started again.</p> <p>=2 – Flash erase error. Error during Flash erasure procedure probably due to Flash part corruption.<br/> <b>Repair action:</b> The Flash device must be replaced.</p> <p>=3 – Flash state information..<br/> <b>Repair action:</b> Check extended info byte in Event Data 3 (byte 7) whether this is wear-out protection causing this event. If so just wait until wear-out protection expires, otherwise probably the flash device must be replaced (if error is persistent).</p> <p>=4 – Internal error. Error during firmware execution – FW Watchdog Timeout.<br/> <b>Repair action:</b> Operational image shall be updated to other version or hardware board repair is needed (if error is persistent).</p> <p>=5 – BMC did not respond to cold reset request and ME rebooted the platform.<br/> <b>Repair action:</b> Verify the Node Manager configuration.</p> <p>=6 – Direct Flash update requested by the BIOS. ME Firmware will switch to recovery mode to perform full update from BIOS.<br/> <b>Repair action:</b> This is transient state. ME Firmware should return to operational mode after successful image update performed by the BIOS.</p> <p>=7 – Manufacturing error. Wrong manufacturing configuration detected by ME Firmware.<br/> <b>Repair action:</b> The Flash device must be replaced (if error is persistent).</p> <p>=8 – Persistent storage integrity error. Flash file system error detected.<br/> <b>Repair action:</b> If error is persistent restore factory presets using “Force ME Recovery” IPMI command or by doing AC power cycle with Recovery jumper asserted.</p> <p>=9..255 – Reserved</p> <p>Byte 7 – Event Data 3<br/> Extended error info. Should be used when reporting an error to the support. The following values can be interpreted directly: Event Data 2 (byte 6) equal to 3<br/> = 00 - recovery bootloader image or Factory presets image corrupted = 01 - flash erase limit has been reached<br/> = 02 - flash write limit has been reached, writing to flash has been disabled.<br/> = 03 - writing to the flash has been enabled</p> <p>Event Data 2 (byte 6) equal to 7<br/> = 04 - ME FW configuration is inconsistent or out of range</p> |             |
|                          |         | <b>Response</b>  |             |



| Net Function = S/E (04h) |         |  |             |
|--------------------------|---------|--|-------------|
| Code                     | Command | Request, Response Data   | Description |
|                          |         | Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11) |             |

## 2.7 IPMI OEM ME Firmware Update Commands

Server Node Manager Firmware supports online firmware update only in dual operational image with the recovery boot-loader image flash configuration. In that case the same set of commands for image upgrade is available from any operational image. During upgrade all enabled functionalities will work. A rollback scenario is supported.

In single operational image mode the online firmware update functionality is not supported.

To finish the operational image upgrade after successful image registering “Cold reset” command should be sent to in order to boot the new operational image.

Online update of Operational image upgrades only the code. Settings, such as any configuration information including DCMI and NM configured NM policies, factory presets and recovery firmware remain unchanged after the operational code upgrade.

Runtime data sent by BIOS and/or BMC is preserved between cold resets of Intel ME.

All the below Firmware Update commands are not available for 5 seconds after platform or CPU reset. In such case, NM firmware returns D5h (Command Not Supported in Present Stats) Completion Code.

The length of the supported IPMI frames is extended to 80 bytes so up to 68 bytes of payload can be passed in one IPMI request.

| Net Function = 2Eh-2Fh |                                  |  |  |
|------------------------|----------------------------------|--|--|
| Code                   | Command                          | Request, Response Data   | Description  |
| A0h                    | Online Update Prepare For Update | <b>Request</b><br>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first  | This is the first command that should be sent to initiate FW upgrade.  |
|                        |                                  | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11)<br>=80h – Operation refused (too many requests)<br>=81h – Flash error<br>=82h – Operation in progress (flash erase)<br>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first | This command restarts the FW upgrade to the initial state.<br><br>Upon success command <b>Online Update Get Status</b> returns “update requested” status. Upon failure “update init failed” is returned. |
| A1h                    | Online Update Open Area          | <b>Request</b><br>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first<br><br>Byte 4 – Area type<br>=00h – Reserved   | Only update of operational code will be supported<br><br><i>Note: Online Update Open Area invalidates</i>  |

| Net Function = 2Eh-2Fh |                          |  |  |
|------------------------|--------------------------|--|--|
| Code                   | Command                  | Request, Response Data   | Description  |
|                        |                          | =01h – Operational code<br>=02h – PIA<br>=03h – SDR<br><br>Byte 5 – Area flags<br>[0:7] – Reserved. Write as 00000000b.  | the rollback image partition. Only a successful image upload allows switching to a rollback image.<br><br>Upon success command <b>Online Update Get Status</b> returns “update failed” status in case of flash error or “update in progress” status in case of success.<br><br>This command performs access to the flash so the response will be sent after completing the operation and may take longer than 250ms. |
|                        |                          | <b>Response</b><br><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11)<br>80h – Operation refused. After this error code the FW upgrade should be reinitialized by sending <b>Online Update Prepare For Update</b> command.<br>81h – Flash error. After this error code the FW upgrade should be reinitialized by sending <b>Online Update Prepare For Update</b> command.<br><br>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first                                       |  |
| A2h                    | Online Update Write Area | <b>Request</b><br><br>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first<br><br>Byte 4 – Sequence number – Must start with 0 value<br><br>Byte 5:N – Area data, where <N> should not exceed 73. Maximum supported number of data bytes is 68  | This command writes data into opened area.<br><br>Upon success command <b>Online Update Get Status</b> returns “update failed” status in case of flash error or “update in progress” status in case of success.  |
|                        |                          | <b>Response</b><br><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11)<br>=80h – Operation refused. Wrong CRC or image length mismatch. After this error code the FW upgrade should be reinitialized by sending <b>Online Update Prepare For Update</b> command.<br>=81h – Flash error. After this error code the FW upgrade should be reinitialized by sending <b>Online Update Prepare For Update</b> command.<br><br>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first | This command performs access to the flash so the response will be sent after completing the operation and may take longer than 250 ms.   |
| A3h                    | Online Update Close Area | <b>Request</b><br><br>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first<br><br>Byte 4:7 – Area size in bytes<br><br>Byte 8:9 – Area checksum. Byte 9 should be set to 0. Byte 8 is a CRC8 ATM HEC (based on $x^8 + x^2 + x + 1$ polynomial) calculated over the operational binary image directly from the distribution package i.e. NMOperational.bin file.   | This command verifies the image checksum and size.<br><br>Upon success command <b>Online Update Get Status</b> returns “update failed” status in case of checksum verification failure or “update requested” status otherwise.   |
|                        |                          | <b>Response</b><br><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard  |  |



| Net Function = 2Eh-2Fh |                               |   |  |
|------------------------|-------------------------------|---|--|
| Code                   | Command                       | Request, Response Data  | Description  |
|                        |                               | <p>Completion Codes are shown in section 2.11)<br/>           =80h – Operation refused. Wrong CRC or image length mismatch. After this error code the FW upgrade should be reinitialized by sending <b>Online Update Prepare For Update</b> command.<br/>           =81h – Flash Error. After this error code the FW upgrade should be reinitialized by sending <b>Online Update Prepare For Update</b> command.<br/>           =82h – In Progress. Not used on Nehalem-EP platform.</p> <p>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first</p>   |  |
| A4h                    | Online Update Register Update | <p><b>Request</b></p> <p>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first</p> <p>Byte 4 – update type<br/>           =0 – Reserved<br/>           =1 – Normal update. Use the new image for the next boot. Use this update type to verify and to switch to a newly uploaded image – after a successful <b>Online Update Close Area</b> command. Upon success command <b>Online Update Get Status</b> returns “update failed” status in case of image verification error or “update success” status otherwise.<br/>           =2 – Reserved<br/>           =3 – Manual rollback. Use this update type to cancel the switch to a new image and to use the current code – when executed right after <b>Online Update Prepare For Update</b> command. Upon success command <b>Online Update Get Status</b> returns “update failed” status in case of image verification error or “update rolled back” status otherwise.<br/>           =4 – Abort update. Exits upgrade. Upon success command <b>Online Update Get Status</b> returns “update aborted” status.</p> <p>Byte 5 – dependent flags<br/>           [0:7] – Reserved. Write as 00000000b.</p> <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>           =00h – Success (Remaining standard Completion Codes are shown in section 2.11)<br/>           =80h – Operation refused.<br/>           =81h – Flash Error. After this error code the FW upgrade should be reinitialized by sending <b>Online Update Prepare For Update</b> command.<br/>           =CCh – Invalid field</p> <p>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first</p> | <p>Instructs the controller that all areas to be updated have been sent and schedules an update to occur on the next ME FW reset. End user is expected to perform system power cycle to reset the ME FW and host.</p> <p>Upon success command <b>Online Update Get Status</b> returns status as described on the Figure below.</p> <p>This command performs access to the flash so the response will be sent after completing the operation and may take longer than 250 ms.</p> |
| A6h                    | Online Update Get Status      | <p><b>Request</b></p> <p>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first</p>  | Returns the current status of the FW Update  |

| Net Function = 2Eh-2Fh |                                |   |                                   |
|------------------------|--------------------------------|---|-----------------------------------|
| Code                   | Command                        | Request, Response Data  | Description                       |
|                        |                                | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11)<br>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first<br>Byte 5 – Image status<br>[0] – Reserved. Return as 0b.<br>[1] – Staging image (new)<br>=1 – Image valid<br>[2] – Rollback image<br>=1 – Image valid<br>[3:4] – Running image area<br>=0 – CODE (Recovery mode)<br>=1 – COD1<br>=2 – COD2<br>[5:7] – Reserved. Return as 000b.<br>Byte 6 – Update state<br>=0 – Idle (no update in progress)<br>=1 – Update requested<br>=2 – Update in progress<br>=3 – Update success<br>=4 – Update failed<br>=5 – Update rolled back<br>=6 – Update aborted<br>=7 – Update initialization failed<br>=8-255 – Reserved<br>Byte 7 – Update Attempt Status<br>=0-255 – Reserved. Return as 0.<br>Byte 8 – Rollback Attempt Status<br>=0-255 – Reserved. Return as 0.<br>Byte 9 – Update Type<br>=0-255 – Reserved. Return as 0.<br>Byte 10 – Dependent Flags<br>[0:7] – Reserved. Return as 0.<br>Byte 11:14 – Free Area Size in bytes. |                                   |
| A7h                    | Online Update Get Capabilities | <b>Request</b><br>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first<br><b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11)<br>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first<br>Byte 5 – Areas supported<br>[0] – Reserved. Return as 0b.<br>[1] – Operational code<br>= 1 – OpCode supported<br>[2] – PIA<br>= 1 – PIA supported<br>[3] – SDR<br>= 1 – SDR supported  | PIA and SDR updates not supported |



| Net Function = 2Eh-2Fh |         |   |             |
|------------------------|---------|---|-------------|
| Code                   | Command | Request, Response Data  | Description |
|                        |         | [4:7] – Reserved. Return as 0000b.<br>Byte 6 – Special capabilities<br>[0] – Rollback<br>= 1 – Rollback supported<br>[1] – Recovery<br>= 1 – Recovery supported<br>[2:7] – Reserved. Return as 000000b. |             |

## 2.7.1 Online Update Flow

The sequence of events that occurs during a full update with rollback is as follows.

1. The online update application sends an *Online Update Prepare for Update* command, causing the ME to re-initialize the staging image and enter the update requested state. Update state returned by **Online Update Get Status** command should return value 1 – update requested or 7 – update init failed in a case of flash error.
2. The online update application sends an *Online Update Open Area* command, telling the ME to create an area in the staging image and prepare to receive download data of the specified type. Update state returned by **Online Update Get Status** should return value 2 – update in progress unless problem occurs with the area preparation. In that case update will return value 4 – update failed.
3. The online update application sends multiple *Online Update Write Area* commands to fill the area and then sends an *Online Update Close Area* command to indicate the area is finished. The *Online Update Close Area* command contains length and checksum information that allows the ME to validate that all data was successfully received. Update state returned by **Online Update Get Status** should return value 2 – update in progress in a case of flash write success until *Online Update Close Area* command will be send. From that time update state should return value 1 – update requested or 4 – update failed depending on the CRC verification.
4. The online update application repeats the *Online Update Open Area*, *Online Update Write Area*, and *Online Update Close Area* sequence until all areas that are to be updated have been downloaded. Update state returned by **Online Update Get Status** should return value 2 – update in progress in a case of flash write success until *Online Update Close Area* command will be sent. From that time update state should return value 1 – update requested or 4 – update failed depending on the checksum verification for details see **Online Update Close Area** command description.
5. The online update application sends an *Online Update Register Update* command to indicate the download is complete and is ready to be enacted. The ME verifies the validity of the staging area image and set the status accordingly. When ME is running NM operational image during the update procedure, only basic image verification is done. The verification includes checking image CRC, but does not include checking image signature (i.e. whether the image was correctly signed by Intel). The full image verification is performed, when ME is

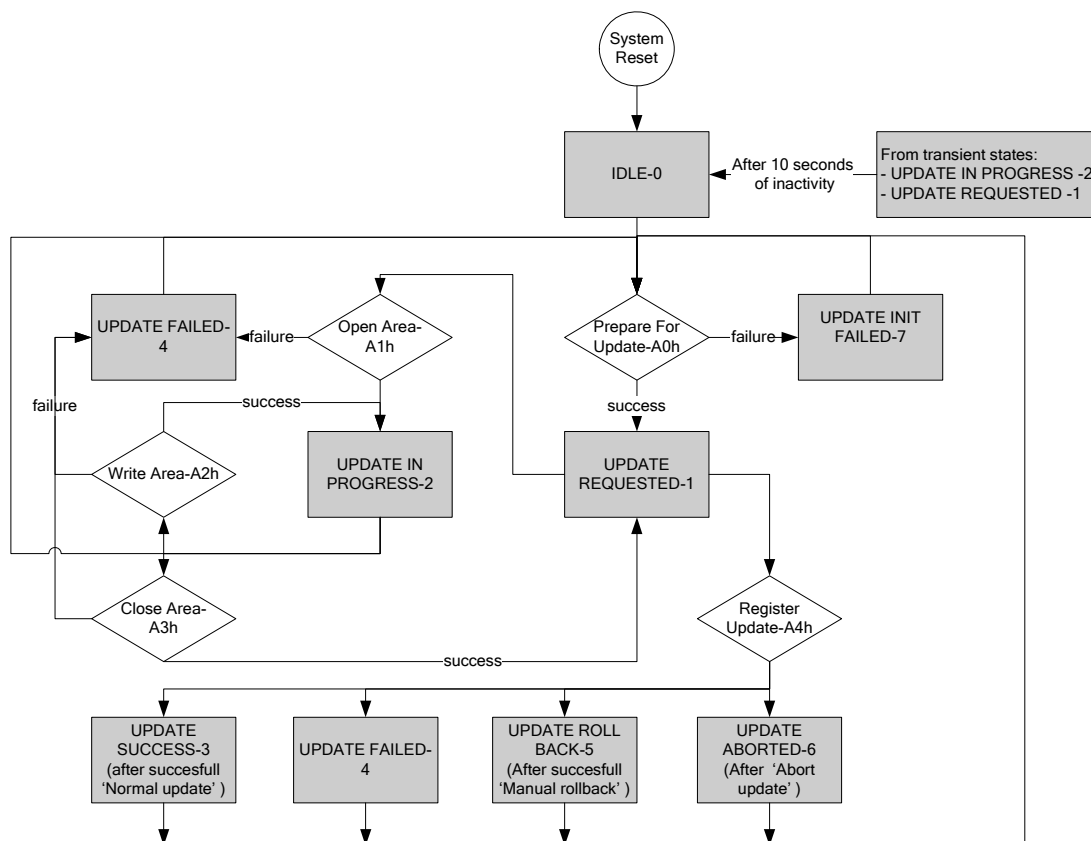
loading the image. Update state returned by **Online Update Get Status** should return value 3 – update success (in a case of successful verification) or 4 – update failed (in a case of verification failure) unless *Online Update Register Update* command was sent without finalizing update with *Online Update Close Area* command. In that case update state will return value 6 – update aborted. Value 5 – update rollback will be returned after successful manual rollback.

6. A system reset occurs (potentially much later).
7. The ME hard resets itself to allow the boot code to run.
8. The boot code verifies the status of the new downloaded operational image, and transfers control to and executes the new image if verification succeeds.
9. If verification status of the update image is bad, the last-known-good operational code bank is automatically selected for execution to prevent the server from being inoperable or degraded. Update state returned by **Online Update Get Status** should return value 0 – idle.

System power cycles are treated the same as system resets. If A/C power is lost before the actual update copying process starts, the registered update is discarded. When A/C power is reapplied, the ME comes up in the idle condition. If A/C power is lost during or after the copy process is started, the copy is resumed after A/C power is restored.

There can be one and only one area of each type in a single update sequence. Partial updates are not permitted.





**Figure 1. Management Engine Online Update Flow**

## 2.7.2 Backward Compatibility Mode

The NM firmware supports backward compatibility mode with the NM1.5 Firmware for the online upgrade commands. So online upgrade commands are accessible also on the Net Function 0x30 with the commands codes A0h to A7h.

Up to 68 bytes of payload can be passed in one IPMI request. Update flow is the same as specified in section 2.7.1.

| Net Function = SDK General Application (0x30) |                                  |   |   |
|---|----------------------------------|---|---|
| Code  | Command                          | Request, Response Data  | Description   |
| A0h   | Online Update Prepare For Update | <b>Request</b><br>None  | This is the first command that should be sent to initiate FW upgrade.   |
|   |                                  | <b>Response</b><br><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11)<br>=80h – Operation refused (too many requests)<br>=81h – Flash error<br>=82h – Operation in progress (flash erase) | This command restarts the FW upgrade to the initial state.<br><br>Upon success command <b>Online Update Get Status</b> returns "update requested" status. Upon failure "update init failed" |

| Net Function = SDK General Application (0x30) |                          |  |   |
|---|--------------------------|--|---|
| Code  | Command                  | Request, Response Data   | Description   |
|   |                          |  | is returned.  |
| A1h   | Online Update Open Area  | <b>Request</b><br>Byte 1 – Area type<br>=00h – Reserved<br>=01h – Operational code<br>=02h – PIA<br>=03h – SDR<br><br>Byte 2 – Area flags<br>[0:7] – Reserved. Write as 00000000b.   | Only update of operational code will be supported<br><br><i>Note: <b>Online Update Open Area</b> invalidates the rollback image partition. Only a successful image upload allows switching to a rollback image.</i>   |
|   |                          | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11)<br>=80h – Operation refused. After this error code the FW upgrade should be reinitialized by sending <b>Online Update Prepare For Update</b> command.<br>=81h – Flash error. After this error code the FW upgrade should be reinitialized by sending <b>Online Update Prepare For Update</b> command.  | Upon success command <b>Online Update Get Status</b> returns “update failed” status in case of flash error or “update in progress” status in case of success.<br><br>This command performs access to the flash so the response will be sent after completing the operation and may take longer than 250 ms. |
| A2h   | Online Update Write Area | <b>Request</b><br>Byte 1 – Sequence number – Must start with 0 value<br><br>Byte 2:N – Area data, where <N> should not exceed 70. Maximum supported number of data bytes is 68.  | This command writes data into opened area.<br><br>Upon success command <b>Online Update Get Status</b> returns “update failed” status in case of flash error or “update in progress” status in case of success.   |
|   |                          | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br>=80h – Operation refused. After this error code the FW upgrade should be reinitialized by sending <b>Online Update Prepare For Update</b> command.<br>=81h – Flash error. After this error code the FW upgrade should be reinitialized by sending <b>Online Update Prepare For Update</b> command. | This command performs access to the flash so the response will be sent after completing the operation and may take longer than 250 ms.  |
| A3h   | Online Update Close Area | <b>Request</b><br>Bytes 1:4 – Area size in bytes.<br><br>Bytes 5:6 – Area checksum. Byte 6 should be set to 0. Byte 5 is a CRC8 ATM HEC (based on $x^8 + x^2 + x + 1$ polynomial) calculated over the operational binary image directly from the distribution package i.e. SPSoOperational.bin file.   | This command verifies the image checksum and size.<br><br>Upon success command <b>Online Update Get Status</b> returns “update failed” status in case of checksum verification failure or “update requested” status otherwise.  |
|   |                          | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).   |   |



| Net Function = SDK General Application (0x30) |                               |  |   |
|---|-------------------------------|--|---|
| Code  | Command                       | Request, Response Data   | Description   |
|   |                               | =80h – Operation refused. Wrong CRC or image length mismatch. After this error code the FW upgrade should be reinitialized by sending <b>Online Update Prepare For Update</b> command.<br>=81h – Flash Error. After this error code the FW upgrade should be reinitialized by sending <b>Online Update Prepare For Update</b> command.<br>=82h – In Progress. Not used on Turley.  |   |
| A4h   | Online Update Register Update | <b>Request</b><br>Byte 1 – Update type<br>=0 – Reserved<br>=1 – Normal update. Use the new image for the next boot. Use this update type to verify and to switch to a newly uploaded image – after a successful <b>Online Update Close Area</b> command. Upon success command <b>Online Update Get Status</b> returns “update failed” status in case of image verification error or “update success” status otherwise.<br>=2 – Reserved<br>=3 – Manual rollback. Use this update type to cancel the switch to a new image and to use the current code – when executed right after <b>Online Update Prepare For Update</b> command. Upon success command <b>Online Update Get Status</b> returns “update failed” status in case of image verification error or “update rolled back” status otherwise.<br>=4 – Abort update. Exits upgrade. Upon success command <b>Online Update Get Status</b> returns “update aborted” status.<br>Byte 2 – Dependent flags<br>[0:7] – Reserved. Write as 00000000b. | Instructs the controller that all areas to be updated have been sent and schedules an update to occur on the next system reset or system power cycle.<br><br>Upon success command <b>Online Update Get Status</b> returns status as described on the Figure below.<br><br>This command performs access to the flash so the response will be sent after completing the operation and may take longer than 250ms. |
|   |                               | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br>=80h – Operation refused.<br>=81h – Flash error. After this error code the FW upgrade should be reinitialized by sending <b>Online Update Prepare For Update</b> command.<br>=CCh – Invalid field.   |   |
| A6h   | Online Update Get Status      | <b>Request</b><br>None   | Returns the current status of the FW Update   |
|   |                               | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br>Byte 2 – Image status<br>[0] – Reserved. Return as 0b.<br>[1] – Staging image (new)<br>=1 – Image valid.<br>[2] – rollback image<br>=1 – Image valid.<br>[3:4] – Running image area  |   |

| Net Function = SDK General Application (0x30) |                                |  |                                   |
|---|--------------------------------|--|-----------------------------------|
| Code  | Command                        | Request, Response Data   | Description                       |
|   |                                | =0 – CODE (Recovery mode)<br>=1 – COD1<br>=2 – COD2<br>[5:7] – Reserved. Return as 000b.<br><br>Byte 3 – Update state<br>=0 – Idle (no update in progress).<br>=1 – Update requested.<br>=2 – Update in progress.<br>=3 – Update success.<br>=4 – Update failed.<br>=5 – Update rolled back.<br>=6 – Update aborted.<br>=7 – Update initialization failed.<br>=8-255 – Reserved.<br><br>Byte 4 – UpdateAttemptStatus<br>=0-255 – Reserved. Return as 0.<br><br>Byte 5 – RollbackAttemptStatus<br>=0-255 – Reserved. Return as 0.<br><br>Byte 6 – Update Type<br>=0-255 – Reserved. Return as 0.<br><br>Byte 7 – DependentFlags<br>[0:7] – Reserved. Return as 0.<br><br>Byte 8:11 – Free Area Size in bytes. |                                   |
| A7h   | Online Update Get Capabilities | <b>Request</b><br>None   | PIA and SDR updates not supported |
|   |                                | <b>Response</b><br><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br><br>Byte 2 – Areas supported<br>[0] – Reserved. Return as 0b.<br>[1] – Operational code<br>= 1 – OpCode supported.<br>[2] – PIA<br>= 1 – PIA supported.<br>[3] – SDR<br>= 1 – SDR supported.<br>[4:7] – Reserved. Return as 0000b.<br><br>Byte 3 – Special capabilities<br>[0] – Rollback<br>= 1 – Rollback supported.<br>[1] – Recovery<br>= 1 – Recovery supported.<br>[2:7] – Reserved. Return as 000000b.  |                                   |



## 2.7.3 IPMI OEM Image Inventory Command

| Net Function = 2Eh-2Fh |                                   |   |   |
|------------------------|-----------------------------------|---|---|
| Code                   | Command                           | Request, Response Data  | Description   |
| A8h                    | Online Update Get Image Inventory | <b>Request</b><br>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first<br><br>Byte 4 = Image ID<br>=0 – Provide information for the Recovery image.<br>=1 – Provide information about 1 <sup>st</sup> operational image.<br>=2 – Provide information about 2 <sup>nd</sup> operational image.<br>=3-255 – Reserved for future use.   | This is the command to query image inventory.<br><br>This command can be used to learn what images are present. |
|                        |                                   | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br>=80h – Operation refused (too many requests).<br>=81h – Flash error.<br>=82h – Operation in progress (flash erase).<br>=C9h – Image ID out of range.<br><br>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first<br><br>Byte 5 – Firmware Revision 1<br>[7] – Device available<br>=0 – Normal operation.<br>=1 – Device firmware update or self-initialization in progress.<br><br>[6:0] – Major Firmware Revision, binary encoded<br>=2<br><br>Byte 6 – Firmware Revision 2: Minor Firmware Revision. BCD encoded.<br><br>Byte 7:10 – Auxiliary Firmware Revision Information<br>Byte 7 – Implemented version of NM Firmware IPMI command specification BCD encoded = 2.0<br>Byte 8 – NM Firmware build number BCD encoded = A.B<br>Byte 9 – NM Firmware last digit of build number and patch number BCD encoded = C.PATCH<br>Byte 10 – Image flags<br>[7:3] – Reserved. Return as 00000b.<br>[2] – Image is currently running<br>=0 – Image not running.<br>=1 – Image is currently a running image.<br>[1:0] – Image type<br>=00b – Recovery image.<br>=01b – Operational image 1.<br>=10b – Operational image 2.<br>=11b – Unspecified: flash error indication.<br><i>Note: Full version number is: "Major Firmware Revision. Minor Firmware Revision.ABC.PATCH" where ABC is Firmware build number.</i> |   |

## 2.7.4 Optimizing Online Upgrade Performance

NM image can take up to 2MB of space on the flash so in order to speed up the upgrade process application responsible for upgrading should implement the following requirements:

- To limit the number of IPMI requests the application should send maximum payload size of 64 to 68 bytes in each IPMI request.
- NM Firmware upgrade is capable to maintain a receive window of 16 IPMI requests so the application should send up to 16 **Online Update Write Area** request without waiting for IPMI message response.
- NM Firmware supports out-of-order IMPI messages in a window size up to 16. It switches to that mode right after verifying that initial part of the image (~16kB) is valid ME firmware load. This allows remote console to automatically resend only the missing messages.
- NM Firmware will try to send responses to **Online Update Write Area** requests in-order. In case a lost message is detected NM firmware will send acknowledge to next successfully received message so the application may choose to resend just the missing request (optimal) or the whole window.
- Flash erase requests will be made asynchronously to the transfer but flash erase may cause delays in upgrade processing. To avoid timeouts and retransmissions on the SMBUS line if the buffer for the IPMI upgrade messages is exhausted NM firmware will respond with Node Busy Completion Code.

## 2.8 IPMI Commands Supported by Recovery Boot Loader

The NM firmware supports only the commands for IPMI over IPMB interface listed below when running in recovery boot loader mode:

- **Get Device ID** (Net Function 06h, Command Code 01h)
- **Cold Reset** (Net Function 06h, Command Code 02h)
- **Get Self Test Results** (Net Function 06h, Command Code 04h)
- **Online Update Get Image Inventory** (Net function 2Eh, Command Code A8h)
- **Force ME Recovery** (Net function 2Eh, Command code DFh)

In recovery mode only local update from the BIOS is possible.

## 2.9 IPMI OEM PECI Proxy Commands

IPMI OEM PECI Proxy Completion Codes used in the commands response:

| Code  | Definition  |
|---|---|
| <b>IPMI OEM PECI Proxy Completion Codes</b> |   |
| 00h   | Command Completed Normally.   |
| A0h   | Partial success (only few first responses are provided; the remaining responses did not fit in the IPMI response message as the response message would exceed maximum IPMI message size supported) – only for Aggregated Send RAW PECI. |
| A1h   | Wrong CPU number.   |



| Code | Definition   |
|------|--|
| A2h  | Command response timeout, retry may be needed.   |
| A4h  | Bad read FCS in the response (even after the retry).   |
| A5h  | Bad write FCS field in the response or Abort FCS in the response (even after the retry).                               |
| A6h  | Wrong (unsupported) write length in IPMI request.  |
| A7h  | Wrong (unsupported) read length in IPMI request.   |
| ABh  | Wrong (unknown/invalid/illegal) command code, request not understood by CPU.   |
| ACH  | CPU not present, this error code is returned if no response from PECI client (client device is not responding at all). |
| D5h  | Command not supported in present state – Platform not in S0/S1 state.  |
| FFh  | Other error encountered (code returned for all other unexpected errors).   |

IPMI OEM PECI Proxy commands are product specific.

| Net Function = 2Eh-2Fh |               |  |  |
|------------------------|---------------|--|--|
| 40h                    | Send Raw PECI | <b>Request</b><br><br>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first.<br><br>Byte 4 – PECI Client Address (part of PECI standard header – values shall be in the range from 0x30 through 0x37).<br><br>Byte 5 – Write Length (part of PECI standard header); this field shall be set to the proper value for this PECI command as if there was AWFCs byte provided but ME FW does not verify if the length matches the PECI protocol specification.<br><br>Byte 6 – Read Length (part of PECI standard header); this field shall be set to the proper value for this PECI command but ME FW does not verify if the length matches the PECI protocol specification.<br><br>Byte 7:M – The remaining part of PECI command following the Read Length field (if any – this field does not exist for PECI Ping command); only write data bytes shall be put here, excluding AWFCs bytes (AWFCs will be added by ME FW); note that the retry bit shall normally be set to zero and the command code byte shall be one of the codes understood by ME FW (0x01, 0xF7, 0xA1, 0xA5, 0xB1, 0xB5, 0x61, 0x65, 0xE1, 0xE5; note that only Domain 0 codes are supported). | The command initiates a single PECI transaction on PECI bus.<br><br><b>Only PECI 3.0 command set is supported.</b><br><br><i>Note: In order to use this command OEM needs to sign “INTEL LICENSE AGREEMENT TO PLATFORM ENVIRONMENT CONTROL INTERFACE SPECIFICATION”.</i> |
|                        |               | <b>Response</b><br><br>Byte 1 – Completion Code (Remaining standard  |  |

| Net Function = 2Eh-2Fh |                          |  |   |
|------------------------|--------------------------|--|---|
|                        |                          | <p>Completion Codes are shown in section 2.11)</p> <p>=00h – PECI response successfully returned (see PECI response Completion Code for detailed response from PECI client, which may be not fully successful).</p> <p>=A4h – Bad read FSC in the response.</p> <p>=A5h – Bad write FCS field in the response.</p> <p>=ABh – Wrong command code.</p> <p>=ACh – CPU not present.</p> <p>=D5h – Platform not in S0/S1 state.</p> <p>=FFh – Other error encountered.</p> <p>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Bytes 5:N – PECI response data (if any – no data is returned for Ping command or for Completion Code in Byte#1 other than 00h); data following the Write FSC field are put here exactly as received from PECI client during Read transaction phase, excluding the Write FCS and Read FCS bytes.</p> <p><b>Retries:</b></p> <ul style="list-style-type: none"> <li>- For no response (all zeros), ME FW performs one retry attempt on the PECI bus before the IPMI response is sent back to BMC</li> <li>- For commands other than Ping(), GetTemp(0x01) and GetDIB(0xF7): <ul style="list-style-type: none"> <li>o ME FW verifies the checksums in the PECI response from client and it performs one retry on the PECI bus for Bad FSC (equivalent to Abort FCS, or no response).</li> <li>o If checksum is valid, ME FW interprets the PECI Completion Code byte and performs one retry on PECI bus for Completion Codes of type 0x8X – in this case the retry bit is set in the PECI request and the AW FCS field is updated for the following commands: WrPkgConfig (0xA5), WrIAMS (0xB5), WrPCICfg (0x65), and WrPCICfgLocal(0xE5).</li> </ul> </li> </ul> <p>After 3 retries on PECI bus, the subsequent retries of the same command are not attempted. The response is returned to BMC even in case of an error (no further retries).</p> |   |
| 41h                    | Aggregated Send Raw PECI | <p><b>Request</b></p> <p>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Bytes 4:M – Raw PECI command bytes formatted according to the same rules as bytes 4 to M in Send Raw PECI.</p> <p>Bytes M+1:N – Next RAW PECI command (if any).</p>   | <p>The command initiates multiple PECI transactions on PECI bus.</p> <p>One IPMI request can contain a bunch of PECI RAW transactions.</p> <p><b>The only PECI 3.0 command set is fully</b></p> |





| Net Function = 2Eh-2Fh |                                |   |  |
|------------------------|--------------------------------|---|--|
|                        |                                | <p><b>Response</b></p> <p>Byte 1 – Completion Code related to overall IPMI request (Remaining standard Completion Codes are shown in section 2.11).<br/>           =00h – Success; it means that all the PECI raw requests have been processed and the responses fit in the IPMI response message (not necessarily that all the responses have completed with success). The completion codes for the particular PECI raw transactions are included in the appropriate parts of this response frame.<br/>           =A0h – Partial response (all the PECI commands have been executed but only few first responses are provided; the remaining responses did not fit in the IPMI response message as the response message would exceed maximum IPMI message size supported see 4.7).<br/>           =D5h – Platform not in S0/S1 state.</p> <p>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 5 – Completion Code for this particular transaction – same as byte#1 in Send Raw PECI command response.</p> <p>Byte 6:6+N – the first PECI response data received from PECI client during Read transaction phase (if any), formatted in the same way as response in Send Raw PECI.</p> <p>Byte N+1:M+1– next PECI transaction: Completion Code byte + response (if any); individual responses are returned in the order they were in the IPMI request.</p> <p><b>Note</b> – It is the responsibility of the BMC to ensure that the responses can fit in the IPMI response message (knowing the max IPMI response frame length supported by ME FW). If some responses do not fit into the IPMI response message, they are not returned (but the execution of the corresponding PECI commands is attempted).</p> <p><b>Note</b> – ME FW interprets each raw PECI request based on the Write Length field in each PECI request. Invalid Write Length in a malformed raw PECI request will likely cause all subsequent requests to be interpreted by ME FW starting from a wrong offset and will likely result in C7h error code for the whole IPMI request.</p> | <p><b>supported.</b></p> <p>The PECI 3.0 command suite retains only the Ping(), GetDIB() and GetTemp() PECI 2.0 commands. Other PECI 2.0 commands ARE NOT SUPPORTED.</p> <p><i>Note: In order to use this command OEM needs to sign “INTEL LICENSE AGREEMENT TO PLATFORM ENVIRONMENT CONTROL INTERFACE SPECIFICATION”.</i></p> |
| 42h                    | CPU Package Configuration Read | <p><b>Request</b></p> <p>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 4 – CPU Number<br/>           [7:2] – Reserved.<br/>           [1:0] – CPU number (starting from 0).</p> <p>Byte 5 – PCS Index</p> <p>Byte 6:7 – Parameter Number (WORD)<br/>           Byte 6 – Parameter [7..0]<br/>           Byte 7 – Parameter [15..8]</p>   | <p>This command provides read access to the "package Configuration Space" that is maintained by the CPU.</p>   |

| Net Function = 2Eh-2Fh |                                 |   |  |
|------------------------|---------------------------------|---|--|
|                        |                                 | <p>Byte 8 – Read Length – number of bytes to read<br/>[7:2] – Reserved.<br/>[1:0] – Read Length – number of bytes to read:<br/>=0 – Reserved – shouldn't be used.<br/>=1 – 1 byte.<br/>=2 – 2 bytes (word).<br/>=3 – 4 bytes (double word).</p> <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>=A1h – Wrong CPU number.<br/>=A2h – Command response timeout.<br/>=A4h – Bad read FSC in the response.<br/>=A5h – Bad write FCS field in the response.<br/>=A7h – Wrong read length.<br/>=ABh – Wrong command code.<br/>=ACh – CPU not present.<br/>=D5h – Platform not in S0/S1 state.<br/>=FFh – Other error encountered.</p> <p>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 5:N – Configuration Data returned by CPU. Size of this field depends on Read Length parameter specified in the request.</p>  |  |
| 43h                    | CPU Package Configuration Write | <p><b>Request</b></p> <p>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 4 – CPU Number<br/>[7:2] – Reserved.<br/>[1:0] – CPU number (starting from 0).</p> <p>Byte 5 – PCS Index.</p> <p>Byte 6:7 – Parameter Number (WORD)<br/>Byte 6 – Parameter [7..0].<br/>Byte 7 – Parameter [15..8].</p> <p>Byte 8 – Write Length – number of bytes to write<br/>[7:2] – Reserved.<br/>[1:0] – Write Length – number of bytes to write<br/>=0 – Reserved – shouldn't be used.<br/>=1 – 1 byte.<br/>=2 – 2 bytes (word).<br/>=3 – 4 bytes (double word).</p> <p>Byte 9:N – Data to be written to CPU. Length of this data (1B, 2B, 4B) depends on Write Length value included in Byte 8 of this request.</p> <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>=A1h – Wrong CPU number.<br/>=A2h – Command response timeout.<br/>=A4h – Bad read FSC in the response.<br/>=A5h – Bad write FCS field in the response.<br/>=A6h – Wrong write length.<br/>=ABh – Wrong command code.<br/>=ACh – CPU not present.<br/>=D5h – Platform not in S0/S1 state.</p> | This command provides write access to the "package Configuration Space" that is maintained by the CPU. |



| Net Function = 2Eh-2Fh |                             |   |   |
|------------------------|-----------------------------|---|---|
|                        |                             | <p>=FFh – Other error encountered.</p> <p>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first.</p>  |   |
| 44h                    | CPU PCI Configuration Read  | <p><b>Request</b></p> <p>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 4 – CPU Number</p> <p>[7] – Reserved.</p> <p>[6] = 1b – PCI local space;</p> <p>The RdPCISConfigLocal() command will be used that provides read access to the PCI configuration space that resides on the processor itself (named here - "local" PCI space). Accessing the local PCI space is possible before BIOS has enumerated the systems buses.</p> <p>[5:2] – Reserved.</p> <p>[1:0] – CPU number (starting from 0).</p> <p>Byte 5:8 – PCI Address</p> <p>[31:28] – Reserved.</p> <p>[27:20] – Bus Number.</p> <p>[19:15] – Device Number.</p> <p>[14:12] – Function Number.</p> <p>[11:0] – Register Address.</p> <p>Byte 9 – Read Length – number of bytes to read</p> <p>[7:2] – Reserved.</p> <p>[1:0] – Read Length – number of bytes to read</p> <p>=0 – Reserved – shouldn't be used.</p> <p>=1 – 1 byte.</p> <p>=2 – 2 bytes (word).</p> <p>=3 – 4 bytes (double word).</p> | <p>The command reads from PCI configuration space of selected CPU.</p> <p>This command allows BMC to read the configuration from the local PCI configuration space as well.</p> |
|                        |                             | <p><b>Response</b></p> <p>Byte 1 – Completion Code</p> <p>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).</p> <p>=A1h – Wrong CPU Number.</p> <p>=A2h – Command response timeout.</p> <p>=A4h – Bad read FSC in the response.</p> <p>=A5h – Bad write FCS field in the response.</p> <p>=A7h – Wrong read length.</p> <p>=ABh – Wrong command code.</p> <p>=ACh – CPU not present.</p> <p>=D5h – Platform not in S0/S1 state.</p> <p>=FFh – Other error encountered.</p> <p>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 5:N – Register Value returned by CPU. Size of this field depends on Read Length parameter specified in the request.</p>  |   |
| 45h                    | CPU PCI Configuration Write | <p><b>Request</b></p> <p>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 4 – PCI type &amp; CPU Number</p> <p>[7] – Reserved.</p> <p>[6] = 1b – PCI local space;</p> <p>The RdPCISConfigLocal() command will be used</p>   | <p>The command writes a value to PCI configuration space of selected CPU.</p> <p>This command allows BMC to write the configuration to the local PCI configuration space</p>    |

| Net Function = 2Eh-2Fh |                 |  |   |
|------------------------|-----------------|--|---|
|                        |                 | <p>that provides write access to the PCI configuration space that resides on the processor itself (named here - "local" PCI space). Accessing The local PCI space is possible before BIOS has enumerated the systems buses.</p> <p>[5:2] – Reserved.<br/>[1:0] – CPU number (starting from 0).</p> <p>Byte 5:8 – PCI Address<br/>[31:28] – Reserved.<br/>[27:20] – Bus Number.<br/>[19:15] – Device Number.<br/>[14:12] – Function Number.<br/>[11:0] – Register Address.</p> <p>Byte 9 – Write Length<br/>[7:2] – Reserved.<br/>[1:0] – Write Length – number of bytes to write<br/>=0 – Reserved – shouldn't be used.<br/>=1 – 1 byte.<br/>=2 – 2 bytes (word).<br/>=3 – 4 bytes (double word).</p> <p>Byte 10:N – Register Value to be written to CPU. Length of this data (1B, 2B, 4B) depends on Write Length value included in Byte 9 of this request.</p> | as well.  |
|                        |                 | <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>=A1h – Wrong CPU Number.<br/>=A2h – Command response timeout.<br/>=A4h – Bad read FSC in the response.<br/>=A5h – Bad write FCS field in the response.<br/>=A6h – Wrong write length.<br/>=ABh – Wrong command code.<br/>=ACh – CPU not present.<br/>=D5h – Platform not in S0/S1 state.<br/>=FFh – Other error encountered.</p> <p>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first.</p>  |   |
| 46h                    | CPU IA MSR Read | <p><b>Request</b></p> <p>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 4 – CPU Number<br/>[7:2] – Reserved.<br/>[1:0] – CPU number (starting from 0).</p> <p>Byte 5 – Thread ID.</p> <p>Byte 6:7 – MSR Address<br/>Byte 6 – MSR Address [7..0]<br/>Byte 7 – MSR Address [15..8]</p> <p>Byte 8 – Read Length – number of bytes to read<br/>[7:3] – Reserved.<br/>[2:0] – Read Length – number of bytes to read<br/>=0 – Reserved – shouldn't be used.<br/>=1 – 1 byte.<br/>=2 – 2 bytes (word).<br/>=3 – 4 bytes (double word).<br/>=4 – 8 bytes (quad word).<br/>=5-7 – Reserved – illegal value in the current</p>   | <p>This command provides access to the IA MSR space (core and uncore). Specific processors might limit accessibility to certain areas of the MSR space. Please refer to the appropriate processor documentation for a description of the accessibility limitations.</p> |



| Net Function = 2Eh-2Fh |                                |  |   |
|------------------------|--------------------------------|--|---|
|                        |                                | version.   |   |
|                        |                                | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br>=A1h – Wrong CPU number.<br>=A2h – Command response timeout.<br>=A4h – Bad read FSC in the response.<br>=A5h – Bad write FCS field in the response.<br>=A7h – Wrong read length.<br>=ABh – Wrong command code.<br>=ACh – CPU not present.<br>=D5h – Platform not in S0/S1 state.<br>=FFh – Other error encountered.<br><br>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first.<br><br>Byte 5:N – Data returned by CPU. Size of this field depends on Read Length parameter specified in the request.  |   |
| 4Bh                    | Get CPU and Memory Temperature | <b>Request</b><br>Byte 1:3 = Intel Manufacturer ID – 000157h, LS byte first.<br><br>Byte 4 – CPUs for which temperature readings are requested.<br>[0] – Bit set indicates CPU#0 readings are requested; bit clear indicates that readings are not requested.<br>[1] – Bit set indicates CPU#1 readings are requested; bit clear indicates that readings are not requested.<br>[2] – Bit set indicates CPU#2 readings are requested; bit clear indicates that readings are not requested.<br>[3] – Bit set indicates CPU#3 (PECI readings are requested; bit clear indicates that readings are not requested.<br>[7:4] – Reserved.<br><br>Byte 5:6 – 16 bits for CPU#0 indicating memory channels and DIMMs for which temperature readings are requested (4x4 bitmask):<br><br>Byte 5 [0] – CHANNEL#0, DIMM#0.<br>Byte 5 [1] – CHANNEL#0, DIMM#1.<br>Byte 5 [2] – CHANNEL#0, DIMM#2.<br>Byte 5 [3] – CHANNEL#0, DIMM#3.<br><br>Byte 5 [4] – CHANNEL#1, DIMM#0.<br>Byte 5 [5] – CHANNEL#1, DIMM#1.<br>Byte 5 [6] – CHANNEL#1, DIMM#2.<br>Byte 5 [7] – CHANNEL#1, DIMM#3.<br><br>Byte 6 [0] – CHANNEL#2, DIMM#0.<br>Byte 6 [1] – CHANNEL#2, DIMM#1.<br>Byte 6 [2] – CHANNEL#2, DIMM#2.<br>Byte 6 [3] – CHANNEL#2, DIMM#3.<br><br>Byte 6 [4] – CHANNEL#3, DIMM#0.<br>Byte 6 [5] – CHANNEL#3, DIMM#1.<br>Byte 6 [6] – CHANNEL#3, DIMM#2. | The command returns CPU and all Memory DIMMs temperature value for a selected CPU(s) and DIMM(s). |

| Net Function = 2Eh-2Fh |  |  |  |
|------------------------|--|--|--|
|                        |  | <p>Byte 6 [7] – CHANNEL#3, DIMM#3.</p> <p>Byte 7:8 – 16 bits for CPU#1 indicating memory channels and DIMMs for which temperature readings are requested (4x4 bitmask) – the format is the same as for CPU#0.</p> <p>Byte 9:10 – 16 bits for CPU#2 indicating memory channels and DIMMs for which temperature readings are requested (4x4 bitmask) – the format is the same as for CPU#0.</p> <p>Byte 11:12 – 16 bits for CPU#3 indicating memory channels and DIMMs for which temperature readings are requested (4x4 bitmask) – the format is the same as for CPU#0.</p>   |  |
|                        |  | <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>           =00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>           =A1h – Wrong CPU number.<br/>           =D5h – Platform not in S0/S1 state.<br/>           =ADh – Response cannot be delivered because its length is not supported for underlying transport.<br/>           =FFh – Other error encountered.</p> <p>When byte 1 indicates success, the remaining bytes contain the thermal status information for requested CPUs and memory DIMMs. Information bytes are not returned for remaining CPUs or memory DIMMs (the length of the response depends on the number of requested CPUs or DIMMs). In other words, the order of bytes returning the readings is the same as listed in this specification in the request, skipping items that are not requested.</p> <p>Byte 2:4 = Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 5:N – CPU temperatures (up to 4 bytes – only bytes for the requested CPUs are returned); the data is returned as an unsigned integer; it is representing the number of degrees of Celsius below the Thermal Control Circuit Activation temperature, with some values are reserved to provide error indication, as specified below.</p> <p>Byte N+1:M – Memory DIMM temperatures (up to 64 bytes – only bytes for the requested DIMMs are returned); each byte shall be interpreted as an unsigned value containing the absolute temperature expressed in degrees of Celsius with the following values reserved to provide error indication:<br/>           =0xFF – Sensor or device not present.<br/>           =0xFE – Reserved.<br/>           =0xFD – Data unavailable due to sensor or interface failure.</p> |  |

## 2.10 IPMI OEM PECI Proxy Sensors

The below table and subsection summarize the sensors exposed by PECI Proxy functionality in ME FW. Sensor readings are not available when system is in low power state.



Note that the default configuration of the sensors can be set using Flash Image Tool. The default configuration includes:

- Thresholds
- Event Enable Mask
- Scanning Periods
- Scanning Enable Flag
- Per-sensor Event Enable Flag

| Sensor Number | Description                                     | Notes   |
|---------------|---|---|
| 28            | CPU#0 Thermal Status                            | Discrete sensor.  |
| 29            | CPU#1 Thermal Status                            | Discrete sensor.  |
| 30            | CPU#2 Thermal Status                            | Discrete sensor.  |
| 31            | CPU#3 Thermal Status                            | Discrete sensor.  |
| 32            | CPU#0 Thermal Control Circuit Activation Sensor | Threshold sensor.   |
| 33            | CPU#1 Thermal Control Circuit Activation Sensor | Threshold sensor.   |
| 34            | CPU#2 Thermal Control Circuit Activation Sensor | Threshold sensor.   |
| 35            | CPU#3 Thermal Control Circuit Activation Sensor | Threshold sensor.   |
| 36            | CPU#0 Tcontrol Sensor                           | OEM sensor that presents T-Control parameter of the CPU. The value of this sensor is constant and may only change upon HW configuration change. |
| 37 CPU#1      | Tcontrol Sensor                                 |   |
| 38 CPU#2      | Tcontrol Sensor                                 |   |
| 39 CPU#3      | Tcontrol Sensor                                 |   |
| 48            | CPU#0 Tjmax Sensor                              | OEM sensor that presents Tjmax parameter of the CPU. The value of this sensor is constant and may only change upon HW configuration change.     |
| 49            | CPU#1 Tjmax Sensor                              |   |
| 50            | CPU#2 Tjmax Sensor                              |   |
| 51            | CPU#3 Tjmax Sensor                              |   |
| 52            | Memory Throttling for CPU#0                     | Threshold sensor.   |
| 53            | Memory Throttling for CPU#1                     | Threshold sensor.   |
| 54            | Memory Throttling for CPU#2                     | Threshold sensor.   |

| Sensor Number | Description                 | Notes             |
|---------------|-----------------------------|-------------------|
| 55            | Memory Throttling for CPU#3 | Threshold sensor. |

### 2.10.1 CPU Thermal Status Sensors

The sensors are discrete sensors presenting various states associated with CPU thermal status.

| Bit Offset | Name                     | Description  |
|------------|--------------------------|--|
| 0          | CPU Critical Temperature | Indicates whether CPU temperature is above critical temperature point.   |
| 1          | PROCHOT# Assertions      | Indicates whether PROCHOT# signal is asserted.   |
| 2          | TCC Activation           | Indicates whether CPU thermal throttling functionality is activated due to CPU temperature being above Thermal Circuit Control Activation point. |

The value of this sensor is updated by ME FW every 250 ms.

### 2.10.2 CPU Thermal Control Circuit Activation Sensors

The sensors are threshold-based sensors presenting the percentage of time the processor has been operating at a lowered performance due to TCC activation. It does not include the TCC activation time as a result of an external assertion of PROCHOT# signal.

The value of the sensor is updated every 250 ms but the sensor returns the average over the last 6 seconds (24 samples).

### 2.10.3 CPU T-Control Sensors

These sensors are presenting T-Control parameter for the processors. T-Control value is fan speed control reference temperature. For detailed description of the value refer to [EMTS] document.

Fan Temperature target offset (T-Control) indicates the relative offset from CPU Tjmax temperature at which fans should be engaged. For detailed description of the Tjmax value refer to [EMTS] document.

### 2.10.4 CPU Tjmax Sensors

These sensors are presenting Tjmax parameter for the processors. CPU Tjmax is the minimum temperature that the processor will start throttling due to TCC activation. For detailed description of the value refer to [EMTS] document.

The value of this sensor is constant and may only change upon HW configuration change.

### 2.10.5 Memory Throttling Status Sensors

These sensors provide information on memory throttling as a percentage (valid range is 0..200, value 1 means 0.5%), of memory cycles were throttled due to power limiting.





The value of the sensor is updated every 250 ms but the sensor returns the average over the last 6 seconds (24 samples).

## 2.11 IPMI Standard Completion Codes

Node Manager firmware IPMI commands use standard Completion Codes from table below and specific OEM commands codes if specified in command description. Unless specified otherwise, by a specific command description, fields following non-zero Completion Code are truncated<sup>2</sup>.

| Code   | Definition   |
|--|--|
| <b>Generic Completion Codes 00h, C0h-FFh</b> |  |
| 00h  | Command Completed Normally.  |
| C0h  | Node Busy. Command could not be processed because command processing resources are temporarily unavailable.<br><br>This Completion Code is returned when ME is erasing Flash and cannot handle IPMI requests – for example during firmware update procedure or when processing an NM configuration update request. |
| C1h  | Invalid Command. Used to indicate an unrecognized or unsupported command.<br><i>Note:</i> For Net Function = 2Eh – 2Fh for all unrecognized IANA Enterprise Numbers C1h is returned in response followed by up to 3 bytes containing unrecognized IANA Enterprise Number are copied from the original request.     |
| C2h  | Command invalid for given LUN.   |
| C3h  | Timeout while processing command. Response unavailable.  |
| C4h  | Out of space. Command could not be completed because of a lack of storage space required to execute the given command operation.   |
| C5h  | Reservation Canceled or Invalid Reservation ID.  |
| C6h  | Request data truncated.  |
| C7h  | Request data length invalid.   |
| C8h  | Request data field length limit exceeded.  |
| C9h  | Parameter out of range. One or more parameters in the data field of the Request are out of range. This is different from 'Invalid data field' (CCh) code in that it indicates that the erroneous field(s) has a contiguous range of possible values.   |
| CAh  | Cannot return number of requested data bytes.  |
| CBh  | Requested Sensor, data, or record not present.   |
| CCh  | Invalid data field in Request  |
| CDh  | Command illegal for specified sensor or record type.   |

---

<sup>2</sup> Exception: for Net Function = 2Eh-2Fh for all Completion Codes up to 3 bytes containing IANA Enterprise Number are copied from the original request.

| Code                                       | Definition  |
|--|---|
| CEh  | Command response could not be provided.   |
| CFh  | Cannot execute duplicated request. This Completion Code is for devices which cannot return the response that was returned for the original instance of the request. Such devices should provide separate commands that allow the completion status of the original request to be determined. An Event Receiver does not use this Completion Code, but returns the 00h Completion Code in the response to (valid) duplicated requests. |
| D0h  | Command response could not be provided. SDR Repository in update mode.  |
| D1h  | Command response could not be provided. Device in firmware update mode.   |
| D2h  | Command response could not be provided. BMC initialization or initialization agent in progress.   |
| D3h  | Destination unavailable. Cannot deliver request to selected destination. E.g. this code can be returned if a request message is targeted to SMS, but receive message queue reception is disabled for the particular channel.  |
| D4h  | Cannot execute command due to insufficient privilege level or other security based restriction (e.g. disabled for 'firmware firewall').   |
| D5h  | Cannot execute command. Command or request parameter(s) not supported in present state.   |
| D6h  | Cannot execute command. Parameter is illegal because command sub-function has been disabled or is unavailable (e.g. disabled for 'firmware firewall').  |
| FFh Unspecified                            | error.  |
| <b>Device Specific (OEM) codes 01h-7Eh</b> |   |
| 01h-7Eh                                    | Device specific (OEM) Completion Codes. This range is used for command specific codes that are also specific for a particular device and version. A-priori knowledge of the device command set is required for interpretation of these codes.   |
| <b>Command Specific codes 80h-Beh</b>      |   |
| 80h-BEh                                    | Standard command-specific codes. This range is reserved for command specific Completion Codes for commands specified in this document.  |

## 2.12 Generic Event/Reading Type Codes

| Generic Event/Reading Type Code | Event/Reading Class | Generic Offset | Description                        |
|---------------------------------|---------------------|----------------|------------------------------------|
| 01h                             | Threshold           | 00h            | Lower Non-critical – going low     |
|                                 |                     | 01h            | Lower Non-critical – going high    |
|                                 |                     | 02h            | Lower Critical – going low         |
|                                 |                     | 03h            | Lower Critical – going high        |
|                                 |                     | 04h            | Lower Non-recoverable – going low  |
|                                 |                     | 05h            | Lower Non-recoverable – going high |
|                                 |                     | 06h            | Upper Non-critical – going low     |
|                                 |                     | 07h            | Upper Non-critical – going high    |
|                                 |                     | 08h            | Upper Critical – going low         |
|                                 |                     | 09h            | Upper Critical – going high        |
|                                 |                     | 0Ah            | Upper Non-recoverable – going low  |
|                                 |                     | 0Bh            | Upper Non-recoverable – going high |



## 3. ME NM IPMI Interface

---

This chapter describes IPMI command interface used by the Intel® Intelligent Power Technology Node Manager 2.0 implementation:

- External NM IPMI commands that should be used by the external management console, see Chapter 3.1
- Discovery mechanism for the NM functionality, see Chapter 4.5
- Product specific NM IPMI commands used for low-level power management access and for debugging, see Chapter 3.1.1
- External IPMI sensors that should be monitored by external management console, see Chapter 0
- External DCMI Power Management Commands that should be used by the external DCMI management console, see Chapter 3.3

If NM SKU is disabled in firmware then firmware will respond with C1h Invalid Command Completion Code to the commands listed in this chapter. In addition, BMC should not expose the NM discovery SDR to the external console.

### 3.1 External Node Manager Configuration and Control Commands

Intel® Intelligent Power Technology Node Manager (NM) is a platform resident technology that enforces power and thermal policies for the platform. These policies are applied by exploiting subsystem knobs (such as processor P and T states) that can be used to control power consumption. Node Manager enables data center power and thermal management by exposing an external interface to management software through which platform policies can be specified. It also implements specific data center power management usage models such as power limiting.

The configuration and control commands are used by the external management software or BMC to configure and control the Node Manager feature. Since Node Manager firmware does not have any external interface, all these commands are first received by the BMC over LAN and then relayed to the Node Manager firmware over IPMB channel. The BMC merely acts as a relay and the transport conversion device for these commands using the standard IPMI bridging. In that case the privilege level to access to the ME SMLINK channel should be restricted to allow only the Admin level.

BMC provides the access point for remote commands from external management SW and generates alerts to them. In case Node Manager is on the ME, which is an IPMI satellite controller, there have to be mechanisms to forward commands to ME and send response back to originator. Similarly events from the Intel® ME have to be sent as alerts outside of BMC. It is the responsibility of BMC to implement these mechanisms for communication with Node Manager.

Note that all the below commands are not supported when Node Manager Feature Enabled is set to 'false' using Flash Image Tool.

| Net Function = 2Eh-2Fh |  |  |  |
|------------------------|--|--|--|
| Code                   | Command                                    | Request, Response Data   | Description  |
| C0h                    | Enable/Disable Node Manager Policy Control | <b>Request</b><br>Byte 1:3 - Intel Manufacturer ID – 000157h, LS byte first<br><br>Byte 4 – Flags<br>[0:2] – Policy Enable/Disable<br>=0h – Global Disable Node Manager policy control – disables policy control for all power domains regardless of the value set in Domain ID field (Byte 5).<br>=1h – Global Enable Node Manager policy control – enables policy control for all power domains regardless of the value set in Domain ID field (Byte 5).<br>=2h – Per Domain Disable Node Manager policies for the domain given by Byte 5.<br>=3h – Per Domain Enable Node Manager policies for the domain given by Byte 5.<br>=4h – Per Policy Disable Node Manager policy for the domain/policy given by Byte 5 and Byte 6.<br>=5h – Per Policy Enable Node Manager policy for the domain/policy given by Byte 5 and Byte 6.<br>[3:7] – Reserved. Write as 00000b.<br><br>Byte 5 – Domain ID<br>[0:3] – Domain ID<br>Identifies the domain that this Node Manager policy applies to. This field is valid if Per Policy Enable/Disable is set or if Per Domain Policy Enable/Disable is set)<br>=00h – Entire platform.<br>=01h – CPU subsystem.<br>=02h – Memory subsystem.<br>Other – Reserved.<br>[4:7] – Reserved. Write as 0000b.<br><br>Byte 6 – Policy ID.<br>This field is valid if Per Policy Enable/Disable is set. | Enable or Disable the Node Manager policy control feature.<br><br>Global enable/disable affects all policies for all domains.<br><br>Per Domain enable/disable affects all policies of the specified domain.<br><br>Per Policy enable/disable affects only the policy for the specified domain/policy combination.<br><br>After receiving the command it may take Node Manager up to 2s to stop limit power. |
|                        |  | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br>=80h – Policy ID Invalid.<br>=81h – Domain ID Invalid.<br><br>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.   |  |
| C1h                    | Set Node Manager Policy                    | <b>Request</b><br>Byte 1:3 - Intel Manufacturer ID – 000157h, LS byte first.<br><br>Byte 4 – Domain ID<br>[0:3] – Domain ID (Identifies the domain that this Node Manager policy applies to)<br>=00h – Entire platform.<br>=01h – CPU subsystem.<br>=02h – Memory subsystem.<br>Others – Reserved.<br>[4] – Policy Enabled (set to 1 if policy should be enabled by default during policy creation/modification). Policy will be enforced (enabled and evaluated in runtime) if the corresponding Per Domain control as well as Global control is already enabled see C0h command.<br>[5:7] – Reserved. Write as 000b.<br><br>Byte 5 – Policy ID.<br><br>Byte 6 – Policy Type   Policy Trigger Type  | User can specify any valid Policy ID. If already existing, this command will overwrite/modify the parameters for the existing policy, otherwise a new policy will be created with this policy Id. Modification is possible only if that policy for the specified Policy ID is disabled.<br><br><i>Note:</i> The Policy ID is unique over all domains. Set done for   |



| Net Function = 2Eh-2Fh |         |  |   |
|------------------------|---------|--|---|
| Code                   | Command | Request, Response Data   | Description   |
|                        |         | <p>[0:3] – Policy Trigger Type<br/>           =0 – No Policy Trigger, (In that case Policy Trigger Limit should be ignored).<br/>           =1 – Inlet Temperature Limit Policy Trigger in [Celsius].<br/>           =2 – Missing Power Reading Timeout in 1/10<sup>th</sup> of second.<br/>           =3 – Time After Host Reset Trigger in 1/10<sup>th</sup> of second.<br/>           =4 – Boot time policy. This policy will apply the power policy at boot time. This type of policy can be applied only to the Domain 00h and will be applied on each platform restart.</p> <p>[4] – Policy Configuration Action<br/>           =0 – Policy Pointed by Policy Id shall be removed (remaining bytes shall be ignored on read). Corresponding (with the same Policy Id) Alert Thresholds and Suspend Periods will be removed as well.<br/>           =1 – Add Power Policy. This command creates/modifies policy of type that will maintain Power limit.</p> <p>[5:6] – Aggressive CPU Power Correction<br/>           For policies with Domain ID 0 (Entire platform) and 1 (CPU subsystem) the flag indicates whether Node Manager can use CPU T-states to control CPU power consumption. This setting is ignored for some types of policies. I.e. a boot time policy is never aggressive and a policy with missing power reading timeout trigger is always aggressive.</p> <p>=00b – Automatic mode (default). Usage of T-states depends on Shutdown System bit in Policy Exception Actions field. When the bit is set to 1, NM shall use aggressive mode for power limiting (T-states and memory throttling). When the bit is cleared, NM does not use T-states and memory throttling. This behavior is backward compatible with NM 1.5.</p> <p>=01b – Force non-aggressive mode e.g. NM is not allowed to use T-states and memory throttling. User should use this setting if the Node Manager should use only performance-friendly controls.</p> <p>=10b – Force aggressive mode e.g. NM is allowed to use T-states and memory throttling. User should use this setting only if the target limit should be kept at any cost.</p> <p>=11b – Reserved.</p> <p>For policies with Domain ID 2 (Memory subsystem) the field shall be set to 0.</p> <p>[7] – policy storage option.<br/>           = 0b – persistent storage (default). Policy shall persist across a ME reset.<br/>           = 1b – volatile memory. Policy shall not persist across a ME reset.</p> <p>Byte 7 – Policy Exception Actions performed if policy cannot be maintained (if maintained policy power limit given by bytes 8-9 is exceeded over Correction Time Limit).<br/>           [0] – Send alert.<br/>           [1] – Shutdown system (hard shutdown via BMC).<br/>           [2:7] – reserved. Write as 000000b.</p> <p>Byte 8:9 – Policy Target Limit<br/>           For the following Policy Trigger Type value (Byte 6 bits [0:3]) this field contains the following data:<br/>           0, 1, 3 – Power Limit to be maintained in [Watts] as unsigned integer value. Zero value is treated in a special</p> | <p>existing Policy ID may move the policy to a different domain if different Domain ID is provided.</p> <p>The operator may define a special kind of Inlet Air Temperature policy called Minimum Power Consumption policy with the Power Limit set to 0. The policy does not have the power limit defined. When the inlet air temperature raises above the trigger value defined in the policy, the SPS firmware reduces the power consumption to minimum by requesting OSPM or SMM to set minimum P-state and T-state. The Minimum Power Consumption policy does not allow setting the correction action to "System Shutdown", but the operator can specify whether the NM shall minimize power consumption by only reducing P-state or the firmware shall use both P-state and T-state.</p> |

| Net Function = 2Eh-2Fh |         |  |             |
|------------------------|---------|--|-------------|
| Code                   | Command | Request, Response Data   | Description |
|                        |         | <p>way. If the limit is set to zero, NM sets highest throttling level regardless of the power consumption in the domain. NM does not send Power Limit Exceeded event. The zero value can be set even if the minimum power set for the domain is not zero.</p> <p>2 – Throttling level of the platform in %, where 100% enables maximum throttling of the system.</p> <p>4 – Power profile to be applied to the platform at boot time.</p> <p>[0] – Platform Booting mode<br/> =0 – Platform should boot (during BIOS POST) in power optimized mode. In this mode it's expected that BIOS will consume less power i.e. by running the CPU in LFM and using the least amount of threads.<br/> =1 – Platform should boot (during BIOS POST) in performance optimized mode.</p> <p>[1:7] – Cores Disabled – the number of physical CPU cores that should be disabled on each CPU socket. After disabling the cores BIOS POST should lock that value to the OS so that it cannot enable the cores. E.g. 1 passed on that field means than on each CPU package 1 core should be disabled by the BIOS.</p> <p>[8:15] – Reserved. Write as 00000000b.</p> <p>Byte 10:13 – Correction Time Limit – the maximum time in ms, in which the Node Manager must take corrective actions in order to bring the platform back to the specified power limit before taking the action specified in the "Policy Exception Action" parameter. This is an unsigned integer value.<br/> Correction Time does not apply to Boot Time Policy. If Trigger Type defines Boot Time Policy (4) the Correction Time Limit parameter should be set to zero.</p> <p>Byte 14:15 – Policy Trigger Limit<br/> For the following Policy Trigger Type value (Byte 6 bits [0:3]) this field contains the following data:</p> <p>0 – Policy Trigger Value will be ignored.<br/> 1 – Policy Trigger Value should define the Inlet temperature in Celsius. The inlet temperature value will be compared against this limit and if exceeded, cause a trigger to start enforcing the Power Limit specified (Power limit will not be enforced until the trigger happens).<br/> 2 – Policy Trigger should define time in 1/10 of second to perform an action if Missing Power Reading Timeout is detected.<br/> 3 – Policy Trigger should define time in 1/10 of second after Host reset or startup. If BMC does not send Set Event Receiver command within this time after Host reset or startup Node Manager will activate this policy. This policy could be defined in addition to the standard limiting policies. The policy is automatically disabled after next Host reset.<br/> 4 – Policy Trigger is not applicable for boot time policy and should be set to 0.</p> <p>Byte 16:17 – Statistics Reporting Period in seconds. The number of seconds that the measured power will be averaged over for the purpose of reporting statistics to external management SW. This is a moving window length. Note that this value is different from the period that Node Manager uses for maintaining an average for the purpose of power control. This is unsigned integer value.</p> |             |



| Net Function = 2Eh-2Fh |                         |  |  |
|------------------------|-------------------------|--|--|
| Code                   | Command                 | Request, Response Data   | Description  |
|                        |                         | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br>=80h – Policy ID Invalid.<br>=81h – Domain ID Invalid.<br>=82h – Unknown or unsupported Policy Trigger Type.<br>=84h – Power Limit out of range.<br>=85h – Correction Time out of range.<br>=86h – Policy Trigger value out of range.<br>=89h – Statistics Reporting Period out of range.<br>=8Bh – Invalid value of Aggressive CPU Power Correction field.<br>=D5h – Policy could not be updated since Policy ID already exists and is enabled.<br><br>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.   |  |
| C2h                    | Get Node Manager Policy | <b>Request</b><br>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.<br>Byte 4 – Domain ID<br>[0:3] - Domain ID (Identifies the domain that this Node Manager policy applies to).<br>=00h – Entire platform.<br>=01h – CPU subsystem.<br>=02h – Memory subsystem.<br>Others – Reserved.<br>[4:7] – Reserved. Write as 0000b.<br>Byte 5 – Policy ID.  | Gets the Node Manager policy parameters.<br><br>To allow for faster enumeration of all defined policies the error code 80h returns extended error information. |
|                        |                         | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br>=80h – Policy ID Invalid. In addition to bytes 2 to 4 extended error information is returned for this error code.<br>=81h – Domain ID Invalid. In addition to bytes 2 to 4 extended error information is returned for this error code.<br><br>For <b>Completion Code 00h</b> (Success) response bytes 2 to 17 are defined as follows:<br>Byte 2:4 - Intel Manufacturer ID – 000157h, LS byte first.<br>Byte 5 – Domain ID<br>[0:3] - Domain ID (Identifies the domain that this Node Manager policy applies to).<br>=00h – Entire platform.<br>=01h – CPU subsystem.<br>=02h – Memory subsystem.<br>Others – Reserved.<br>[4] – Policy enabled.<br>[5] – Per Domain Node Manager policy control enabled.<br>[6] – Global Node Manager policy control enabled.<br>[7] – Set to 1 if policy is created and managed by other management client e.g. DCMI management API.<br>Byte 6 – Policy Type   Policy Trigger Type<br>[0:3] – Policy Trigger Type |  |

| Net Function = 2Eh-2Fh |         |   |             |
|------------------------|---------|---|-------------|
| Code                   | Command | Request, Response Data  | Description |
|                        |         | <p>=0 – No Policy Trigger, Policy will maintain Power limit (in that case Policy Trigger Value will be equal to the Power Limit).<br/>           =1 – Inlet Temperature Limit Policy Trigger in [Celsius].<br/>           =2 – Missing Power Reading Timeout in 1/10<sup>th</sup> of second.<br/>           =3 – Time After Host Reset Trigger in 1/10<sup>th</sup> of second.<br/>           =4 – Boot time policy.<br/>           [4] – Policy Type<br/>           =1 – Power Control Policy. Policy will maintain Power limit.</p> <p>[5:6] – Aggressive CPU Power Correction<br/>           For policies with Domain ID 0 (Entire platform) and 1 (CPU subsystem) the flag indicates whether Node Manager can use CPU T-states to control CPU power consumption.</p> <p>=00b – Usage of T-states depends on Shutdown System bit in Policy Exception Actions field. When the bit is set to 1, NM shall use T-states. When the bit is cleared, NM does not use T-states. This behavior is backward compatible with NM 1.5.<br/>           =01b – NM is not allowed to use T-states.<br/>           =10b – NM is allowed to use T-states.<br/>           For policies with Domain ID 2 (Memory subsystem) the field shall be set to 0.</p> <p>[7] - policy storage option.<br/>           = 0b – persistent storage<br/>           = 1b – volatile memory.</p> <p>Byte 7 – Policy Exception Actions (if maintained policy power limit given by bytes 8-9 is exceeded over Correction Time Limit).<br/>           [0] – Send alert.<br/>           [1] – Shutdown system.<br/>           [2:7] – Reserved. Write as 000000b.</p> <p>Byte 8:9 – Power Limit.</p> <p>Byte 10:13 – Correction Time Limit - the maximum time in ms, in which Node Manager must take corrective actions in order to bring the platform back within the specified power limit before taking the action specified in the "Policy Exception Action" parameter. This is unsigned integer value.</p> <p>The time is counted from the moment when the average power consumption exceeds the power limit. The average power is calculated as arithmetic moving average with the time period equal to the half of Correction Time Limit. It means that Node Manager may take the exception action after the time period equal to 1.5 of Correction Time Limit parameter starting from the moment when instantaneous power crossed the power limit.</p> <p>Bytes 14:15 – Policy Trigger Limit</p> <p>For the following returned Policy Trigger Type value (Byte 6 bits [0:3]) this field contains the following data:<br/>           0 – The same value as Power Limit (i.e. it does not contain the trigger value passed to Node Manager using Set Node Manager Policy). This is unsigned integer value.<br/>           1 – Inlet temperature in Celsius. The inlet temperature value will be compared against this limit and if exceeded, cause a trigger to start enforcing the Power Limit specified (Power limit will not be enforced until the trigger happens).<br/>           2 – Time After Host Reset Trigger in 1/10<sup>th</sup> of second.<br/>           3 – Policy Trigger defines time in 1/10 of second after Host reset or startup. If BMC does not send Set Event Receiver command within this time after Host reset or startup Node</p> |             |





| Net Function = 2Eh-2Fh |  |  |  |
|------------------------|--|--|--|
| Code                   | Command                                  | Request, Response Data   | Description  |
|                        |  | <p>Manager will activate this policy. This policy could be defined in addition to the standard limiting policies. The policy is automatically disabled after next Host reset.</p> <p>4 – Boot time policy. This field is set to 0.</p> <p>Byte 16:17 – Statistics Reporting Period<br/>The number of seconds that the measured power will be averaged over for the purpose of reporting statistics to external management SW. This is a moving window length. Note that this value is different from the period that Node Manager uses for maintaining an average for the purpose of power control. This is unsigned integer value.</p> <p>For <b>Completion Code 80h</b> (Policy ID Invalid) response bytes 2 to 6 are defined as follows:</p> <p>Byte 2:4 - Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 5 – Next valid Policy ID.</p> <p>The field contains lowest valid Policy ID that is higher than Policy ID specified in the request for the Domain ID specified in the request. If no such Policy ID exists, zero value is returned.</p> <p>Byte 6 – Number of defined policies for the specified in request Domain ID.</p> <p><b>Note</b> – This information can be used to query all existing policies within specified domain. Start with Domain ID and Policy ID set to 0. Increment Policy ID treated as an unsigned integer value by one on success and set Policy ID to Byte 5 on reception of Completion Code 80h.</p> <p>For <b>Completion Code 81h</b> (Domain ID Invalid ) response bytes 2 to 6 are defined as follows:</p> <p>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 5 – Next valid Domain ID.</p> <p>[0:3] – Next valid Domain ID. The field contains lowest valid Domain ID that is higher than Domain ID specified in the request. If no such Domain ID exists, zero value is returned.</p> <p>[4:7] – Reserved. Write as 0000b.</p> <p>Byte 6 – Number of available domains.</p> <p><b>Note</b> – This information can be used to query all available domains. Start with Domain ID and Policy ID both set to 0. Increment Domain ID treated as an unsigned numerical value by one on success and set Domain ID to Byte 5 bits [0:3] on reception of Completion Code 81h.</p> |  |
| C3h                    | Set Node Manager Policy Alert Thresholds | <p><b>Request</b></p> <p>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 4 – Domain ID.</p> <p>[0:3] – Domain ID (Identifies the domain that this Node Manager policy applies to.)</p> <p>=00h – Entire platform.</p> <p>=01h – CPU subsystem.</p> <p>=02h – Memory subsystem.</p> <p>Others – Reserved.</p> <p>[4:7] – Reserved. Write as 0000b.</p>   | Sets the Node Manager Policy alert thresholds. This is part of the Node Manager Policy described earlier and applies to the same policy as specified by Policy ID. |

| Net Function = 2Eh-2Fh |  |   |  |
|------------------------|--|---|--|
| Code                   | Command                                  | Request, Response Data  | Description                                    |
|                        |  | <p>Byte 5 – Policy ID.</p> <p>Byte 6 – Number of alert thresholds as unsigned integer value.</p> <p>Byte 7:N – Alert threshold array (the array length is based on the number of thresholds given in the byte 6). The interpretation of the content of the array depends on the Trigger Type of the corresponding policy. For a policy without trigger the thresholds array contains average power consumption in Watts. For a policy with Inlet Temperature Trigger the array contains temperature in Celsius degrees. For Missing Power Reading Timeout and Time After Host Reset triggers the array contains time 1/10<sup>th</sup> of second. For Boot Time policy the thresholds cannot be defined. Node Manager will generate an event if the trigger value or the average power (computed over an averaging period derived based on correction time limit) exceeds any of the configured alert thresholds. Assertion is generated for exceeding the threshold (going high) and deassertion for going low. The hysteresis value for avoiding jitters around the threshold will be OEM configurable using factory-preset values.</p> <p><b>Note</b> – Max 3 alert thresholds are supported per policy. Each alert threshold is 2 bytes in length (LSB first). If number of alert thresholds is 0 then the previously set alert thresholds (if present) are removed from the policy. All the thresholds are unsigned integer values.</p> <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>           =00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>           =80h – Policy ID Invalid.<br/>           =81h – Domain ID Invalid.<br/>           =84h – Limit in one of thresholds is invalid.<br/>           =87h – Number of thresholds is too large or power limits are invalid.<br/>           =D5h – Alert thresholds cannot be changed for enabled policy, disable it first.</p> <p>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.</p> |  |
| C4h                    | Get Node Manager Policy Alert Thresholds | <p><b>Request</b></p> <p>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 4 – Domain ID<br/>           [0:3] – Domain ID (Identifies the domain that this Node Manager policy applies to).<br/>           =00h – Entire platform.<br/>           =01h – CPU subsystem.<br/>           =02h – Memory subsystem.<br/>           Others – Reserved.<br/>           [4:7] – Reserved. Write as 0000b.</p> <p>Byte 5 – Policy ID.</p> <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>           =00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>           =80h – Policy ID Invalid.</p>   | Gets the Node Manager Policy alert thresholds. |



| Net Function = 2Eh-2Fh |   |  |  |
|------------------------|---|--|--|
| Code                   | Command                                 | Request, Response Data   | Description  |
|                        |   | <p>=81h – Domain ID Invalid.</p> <p>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 5 – Number of alert thresholds as unsigned integer value.</p> <p>Byte 6:N – Alert threshold array (the array length is based on the number of threshold given in the byte 5). If number of alert thresholds is 0 then the array length is 0 bytes. The interpretation of the content of the array depends on the Trigger Type of the corresponding policy. For a policy without trigger the thresholds array contains average power consumption in Watts. For a policy with Inlet Temperature Trigger the array contains temperature in Celsius degrees. For Missing Power Reading Timeout and Time After Host Reset triggers the array contains time 1/10<sup>th</sup> of second. For Boot Time policy the thresholds cannot be defined.</p> <p><b>Note</b> – Max 3 alert thresholds are supported per policy. Each alert threshold is 2 bytes in length (LSB first). All alert thresholds are unsigned integer values.</p>  |  |
| C5h                    | Set Node Manager Policy Suspend Periods | <p><b>Request</b></p> <p>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 4 – Domain ID</p> <p>[0:3] - Domain ID (Identifies the domain that this Node Manager policy applies to).</p> <p>=00h – Entire platform.</p> <p>=01h – CPU subsystem.</p> <p>=02h – Memory subsystem.</p> <p>Others – Reserved.</p> <p>[4:7] – Reserved. Write as 0000b.</p> <p>Byte 5 – Policy ID.</p> <p>Byte 6 – Number of policy suspend periods. This value should be specified as 0 if all the suspend periods are to be removed (if previously set). This is an unsigned integer value.</p> <p>Byte 7:N – array of policy suspend periods (following information is repeated for each suspend period). Each suspend period is defined by 3 bytes:</p> <p>1<sup>st</sup> byte – Policy suspend start time. It is 0 – 239 number of minutes from mid-night divided by 6. Values 240 – 255 are reserved.</p> <p>2<sup>nd</sup> byte – Policy suspend stop time. It is 1 – 240 number of minutes from mid-night divided by 6. Values 0 and 241 – 255 are reserved.</p> <p>3<sup>rd</sup> byte – Suspend period recurrence pattern:</p> <p>[0] – Repeat the suspend period every Monday.</p> <p>[1] – Repeat the suspend period every Tuesday.</p> <p>[2] – Repeat the suspend period every Wednesday.</p> <p>[3] – Repeat the suspend period every Thursday.</p> <p>[4] – Repeat the suspend period every Friday.</p> <p>[5] – Repeat the suspend period every Saturday.</p> <p>[6] – Repeat the suspend period every Sunday.</p> <p>[7] – Reserved. Write as 0b.</p> <p><b>Note</b> – Policy suspend start and stop time is 1 byte in length each. Max 5 suspend periods can be specified per policy. If the number of policy suspend period (i.e. byte 6) is 0 then the rest of the bytes in the request message are not</p> | <p>Sets the Node Manager Policy suspend period (during which no platform power policy control will be enforced).</p> <p>The suspend periods are applied only if ME has valid RTC time. If RTC synchronization fails new suspend periods cannot be configured and existing periods are ignored.</p> <p>Note that RTC synchronization failure situation is signaled with Node Manager Health Event (see 3.4.7)</p> |

| Net Function = 2Eh-2Fh |   |  |   |
|------------------------|---|--|---|
| Code                   | Command                                 | Request, Response Data   | Description                                 |
|                        |   | <p>required and previously configured suspend periods are removed from the system for the specified policy Id.</p> <p>The suspend periods are specified as an array. For e.g. if policy suspend start time is in byte 7 then byte 8 will contain the policy suspend stop time and byte 9 will contain suspend period recurrence pattern. Similarly, if the 2<sup>nd</sup> set of suspend periods are to be specified then they will be present in bytes 10:12.</p> <p>The suspend times are encoded on one byte each as number of minutes from mid-night divided by 6 to fit into one byte. If there is a need to specify an end-time that is beyond midnight, use two suspend periods, one ending at midnight (suspend stop time byte set to 240) and one from midnight (suspend start time set to 0) until the necessary end-time of the next day.</p> <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>           =00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>           =80h – Policy ID Invalid.<br/>           =81h – Domain ID Invalid.<br/>           =85h – One of periods in the table is inconsistent. Start time is greater than or equal to stop time or stop time sets time beyond 1 day.<br/>           =87h – Number of policy suspend periods invalid.<br/>           =D5h – Suspend periods cannot be changed for enabled policy, disable it first.<br/>           =D6h – Command disabled/unavailable due to lack of RTC synchronization.</p> <p>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.</p> |   |
| C6h                    | Get Node Manager Policy Suspend Periods | <p><b>Request</b></p> <p>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 4 – Domain ID<br/>           [0:3] – Domain ID (Identifies the domain that this Node Manager policy applies to).<br/>           =00h – Entire platform.<br/>           =01h – CPU subsystem.<br/>           =02h – Memory subsystem.<br/>           Others – Reserved.<br/>           [4:7] – Reserved. Write as 0000b.</p> <p>Byte 5 – Policy ID.</p> <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>           =00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>           =80h – Policy ID Invalid.<br/>           =81h – Domain ID Invalid.</p> <p>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 5 – Number of policy suspend periods.</p> <p>Byte 6:N – array of suspend periods. Each suspend period is defined by 3 bytes:<br/>           1<sup>st</sup> byte – Policy suspend start time encoded as a number of minutes from mid-night divided by 6.<br/>           2<sup>nd</sup> byte – Policy suspend stop time encoded as a number of</p>  | Get the Node Manager Policy suspend periods |



| Net Function = 2Eh-2Fh |                               |   |  |
|------------------------|-------------------------------|---|--|
| Code                   | Command                       | Request, Response Data  | Description  |
|                        |                               | <p>minutes from mid-night divided by 6.<br/> 3<sup>rd</sup> byte – Suspend period recurrence pattern:<br/> [7] – Reserved. Write as 0b.<br/> [6] – Repeat the suspend period every Sunday.<br/> [5] – Repeat the suspend period every Saturday.<br/> [4] – Repeat the suspend period every Friday.<br/> [3] – Repeat the suspend period every Thursday.<br/> [2] – Repeat the suspend period every Wednesday.<br/> [1] – Repeat the suspend period every Tuesday.<br/> [0] – Repeat the suspend period every Monday.</p> <p><b>Note</b> – If byte 5 is 00h then no subsequent bytes will be present in the response. This means that there are no suspend periods configured for the specified policy Id.</p> <p><b>Note</b> – The suspend periods are specified as an array. For e.g. if 1<sup>st</sup> suspend period is in bytes 6:8 then bytes 9:11 will contain the 2<sup>nd</sup> policy suspend period.</p>  |  |
| C7h                    | Reset Node Manager Statistics | <p><b>Request</b></p> <p>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 4 – Mode<br/> [0:4] – Mode<br/> =00h – Reset global statistics including power statistics, throttling statistics and inlet temperature statistics.<br/> =01h – Reset per policy statistics including power, throttling statistics and trigger statistics.<br/> =1Bh – Reset global Host Unhandled Requests statistics.<br/> =1Ch – Reset global Host Response Time statistics.<br/> =1Dh – Reset global CPU throttling statistics.<br/> =1Eh – Reset global memory throttling statistics.<br/> =1Fh – Reset global Host Communication Failure statistics.<br/> [5-7] – Reserved. Write as 000b.</p> <p>Byte 5 – Domain ID<br/> [0:3] – Domain ID (Identifies the domain that this Node Manager policy applies to)<br/> =00h – Entire platform.<br/> =01h – CPU subsystem.<br/> =02h – Memory subsystem.<br/> Others – Reserved.<br/> [4:7] – Reserved. Write as 0000b.</p> <p>For Global Host Communication Failure, CPU throttling and memory throttling statistics Domain ID must be set to 00h.</p> <p>Byte 6 – Policy ID (ignored if field Mode is set to 00h).</p> <p><b>Response</b></p> <p>Byte 1 – Completion code<br/> =00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/> =80h – Policy ID Invalid.<br/> =81h – Domain ID Invalid.<br/> =88h – Invalid Mode.</p> <p>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.</p> | <p>This command clears Node Manager statistics of given type.</p> <p>Note that all statistics get cleared at ME restart. Thus in a system where ME works in S0/S1 Only power mode the statistics get cleared every time the system enters any sleeping states.</p> |
| C8h                    | Get Node Manager Statistics   | <p><b>Request</b></p> <p>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 4 – Mode<br/> [0:4] – Mode</p>  | <p>This command provides statistics of requested type. The statistics are collected since last ME restart</p>  |

| Net Function = 2Eh-2Fh |         |   |  |
|------------------------|---------|---|--|
| Code                   | Command | Request, Response Data  | Description  |
|                        |         | <p>=01h – Global power statistics in [Watts].<br/>           =02h – Global inlet temperature statistics in [Celsius].<br/>           =03h – 10h – Reserved.<br/>           =11h – Per policy power statistics in [Watts].<br/>           =12h – Per policy trigger statistics in [Celsius].<br/>           =13h – Per policy throttling statistics in [%].<br/>           =1Ah – Reserved.<br/>           =1Bh – Global Host Unhandled Requests statistics.<br/>           =1Ch – Global Host Response Time statistics.<br/>           =1Dh – Global CPU throttling statistics.<br/>           =1Eh – Global memory throttling statistics.<br/>           =1Fh – Global Host Communication Failure statistics.<br/>           [5:7] – Reserved. Write as 000b.</p> <p>Byte 5 – Domain ID<br/>           [0:3] – Domain ID (Identifies the domain that this Node Manager policy applies to)<br/>           =00h – Entire platform.<br/>           =01h – CPU subsystem.<br/>           =02h – Memory subsystem.<br/>           Others – Reserved.<br/>           [4:7] – Reserved. Write as 0000b.</p> <p><b>Note</b> – For Mode (byte 4) in a range 1Bh – 1Fh Domain ID must be set to 00h.</p> <p>Byte 6 – Policy ID</p> <p><b>Note</b> – If Mode field equals 11h or 12h, this field indicates Policy ID for which the statistics are requested. Otherwise the field shall be set to 0.</p> | <p>or since it was cleared with Reset NM Statistics command.</p> <p>Note that the average values provided here may be different from the averaged values used by Node Manager for taking corrective action or triggering alerts based a 'Set Node Manager Alert Threshold' because the averaging periods could be different.</p> <p>Modes for per policy statistics are correlated to policy trigger type. A request for mode 0x11 can be only issued for policies with trigger type 0 (no trigger defined) and 3 (Time After Host Reset). Mode 0x12 is available only for policies with trigger type set to 1 (Inlet Temperature). Mode 0x13 can be issued for policies with trigger type 2 (Missing Power Readings).</p> <p>Note that Global CPU throttling statistics (1Dh) and Global memory throttling statistics (1Eh) provide information about actual available performance of the platform that is adjusted to the defined power cap limit and the load that runs on the platform. Current Value of those statistics will always be greater than 0 if the policy is triggered and is actively limiting to the defined power limit (restore Byte 21 [7]-Policy activation state bit set to 1). Because NM is continuously adjusting available platform performance</p> |
|                        |         | <p><b>Response</b></p> <p>Byte 1 – Completion code<br/>           =00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>           =80h – Policy ID Invalid.<br/>           =81h – Domain ID Invalid.<br/>           =88h – Invalid Mode.</p> <p>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 5:6 – Current Value (unsigned integer) see 3.1.1</p> <p>Byte 7:8 – Minimum Value (unsigned integer) see 3.1.2</p> <p>Byte 9:10 – Maximum Value (unsigned integer) see 3.1.3</p> <p>Byte 11:12 – Average Value (unsigned integer) see 3.1.4</p> <p>Byte 13:16 – Timestamp as defined by the IPMI v2.0 specification indicating when the response message was sent. If NM cannot obtain valid time, the timestamp is set to FFFFFFFFh as defined in the IPMI v2.0 specification.</p> <p>Byte 17:20 – Statistics Reporting Period (the timeframe in seconds, over which the firmware collects statistics). This is unsigned integer value. For all global statistics this field contains the time after the last statistics reset.</p> <p>Byte 21 – Domain ID   Policy State<br/>           [0:3] – Domain ID (Identifies the domain that this Node Manager policy applies to)<br/>           =00h – Entire platform.<br/>           =01h – CPU subsystem.<br/>           =02h – Memory subsystem.<br/>           Others – Reserved.</p>   |  |



| Net Function = 2Eh-2Fh |                               |   |  |
|------------------------|-------------------------------|---|--|
| Code                   | Command                       | Request, Response Data  | Description  |
|                        |                               | <p>[4] – Policy/Global Administrative state.</p> <p>If request Byte 4 is in range 11h – 13h state<br/>           =1 – If policy is enabled by user and Node Manager Policy Control is Globally Enabled (see C0h command) and Node Manager Domain control is also Enabled (see C0h command).<br/>           =0 – Otherwise.</p> <p>If request Byte 4 equals 01h or 02h state<br/>           =1 – if Node Manager Policy Control is Globally Enabled (see C0h command).<br/>           =0 – Otherwise.</p> <p>[5] – Policy Operational state</p> <p>If request Byte 4 is in range 11h – 13h state<br/>           =1 – Policy is actively monitoring defined trigger (power or thermal) and will start enforcing the power limit if defined trigger is exceeded.<br/>           =0 – Policy is suspended so it cannot actively limit to defined power limit. It may happen if one of the defined below events happens:<br/>           – Suspend period is enforced.<br/>           – There is a problem with trigger readings.<br/>           – There is a host communication problem.<br/>           – Host is in Sx state.<br/>           – Host did not send End Of POST notification.<br/>           – Policy is administratively disabled.</p> <p>If request Byte 4 equals 01h or 02h state this bit is always set to 0 as it applies only for per policy statistics.</p> <p>[6] – Measurements state<br/>           =1 – Measurements in progress (host CPU is in S0 state and there are no problems with the readings reported to the remote console).<br/>           =0 – Measurements are suspended (host CPU in Sx state or in the S0 state problem with the readings reported to the remote console).</p> <p>[7] – Policy activation state</p> <p>If request Byte 4 is in range 11h – 13h state<br/>           =1 – Policy is triggered and is actively limiting to the defined power limit.<br/>           =0 – Policy is not triggered.</p> <p>If request Byte 4 equals 01h or 02h state this bit is always set to 0 as it applies only for per policy statistics.</p> <p>For Host Response Time, Host Unhandled Requests, Host Communication Failure, CPU and memory throttling statistics Byte 21 is always set to 0.</p> | <p>this will be true even if the power consumption will be below the policy power limit.</p> |
| C9h                    | Get Node Manager Capabilities | <p><b>Request</b></p> <p>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 4 – Domain ID<br/>           [0:3] – Domain ID (Identifies the domain that this Node Manager policy applies to)<br/>           =00h – Entire platform.<br/>           =01h – CPU subsystem.<br/>           =02h – Memory subsystem.<br/>           Others – Reserved.<br/>           [4:7] – Reserved. Write as 0000b.</p>  | <p>Get Node Manager capabilities.</p>  |

| Net Function = 2Eh-2Fh |         |  |             |
|------------------------|---------|--|-------------|
| Code                   | Command | Request, Response Data   | Description |
|                        |         | <p>Byte 5 – Policy Type   Policy Trigger Type<br/>           [0:3] – Policy Trigger Type<br/>           =0 – No Policy Trigger.<br/>           =1 – Inlet Temperature Policy Trigger value in [Celsius].<br/>           =2 – Missing Power Reading Timeout in 1/10<sup>th</sup> of second.<br/>           =3 – Time After Host Reset Trigger in 1/10<sup>th</sup> of second.<br/>           =4 – Boot time policy.<br/>           Others – Reserved.<br/>           [4:7] – Policy Type<br/>           =1 – Power Control Policy.<br/>           Others – Reserved.</p> <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>           =00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>           =81h – Domain ID Invalid.<br/>           =82h – Unknown or unsupported Policy Trigger Type.<br/>           =83h – Unknown Policy Type.</p> <p>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 5 – Max Concurrent Settings – number of policies supported for the given policy trigger type, policy type and Domain ID.</p> <p>Byte 6:7 Max Power/Thermal/Time After Reset</p> <p>If Policy Trigger Type in the request equals 0 (No Policy Trigger) then this field contains maximum power Limit to be maintained. The power value is expressed in [Watts] as unsigned integer value. It can be either received from BIOS (for total power limit in domain 0) or read by NM from CPU (for CPU and memory power limit in domains 1 and 2). It can also be set manually using Set Power Draw Range command. If the field is equal to zero, it means that NM does not impose any limit.</p> <p>If Policy Trigger Type in the request equals 1 (Inlet Temperature Policy Trigger) then this field contains maximum temperature value to be settable as trigger. The temperature is expressed in [Celsius] as unsigned integer value.</p> <p>If Policy Trigger Type in the request equals 2 (Missing Power Reading Timeout) or 3 (Time After Host Reset) then this field contains maximum time to be settable as trigger. The time is expressed in 1/10 seconds as unsigned integer value.</p> <p>If Policy Trigger Type in the request equals 4 (Boot time policy) then this field is not applicable and is set to 0.</p> <p>Byte 8:9 – Min Power/Thermal/Time After Reset</p> <p>If Policy Trigger Type in the request equals 0 (No Policy Trigger) then this field contains minimum power Limit to be maintained. The power value is expressed in [Watts] as unsigned integer value. It can be either received from BIOS (for total power in domain 0) or read by NM from CPU (for CPU and memory power limit in domains 1 and 2). It can also be set manually using Set Power Draw Range command. If this field is equal to zero, it means that NM does not impose any limit.</p> <p>If Policy Trigger Type in the request equals 1 (Inlet Temperature Policy Trigger) then this field contains</p> |             |





| Net Function = 2Eh-2Fh |                          |   |   |
|------------------------|--------------------------|---|---|
| Code                   | Command                  | Request, Response Data  | Description   |
|                        |                          | <p>minimum temperature value to be settable as trigger. The temperature is expressed in [Celsius] as unsigned integer value.</p> <p>If Policy Trigger Type in the request equals 2 (Missing Power Reading Timeout) or 3 (Time After Host Reset) than this field contains minimum time to be settable as trigger. The time is expressed in 1/10 seconds as unsigned integer value.</p> <p>If Policy Trigger Type in the request equals 4 (Boot time policy) than this field is not applicable and is set to 0.</p> <p>Byte 10:13 – Min Correction Time settable in milliseconds (ms) as unsigned integer value.</p> <p>Byte 14:17 – Max Correction Time settable in milliseconds (ms) as unsigned integer value.</p> <p>Byte 18:19 – Min Statistics Reporting Period in seconds as unsigned integer value.</p> <p>Byte 20:21 – Max Statistics Reporting Period in seconds as unsigned integer value.</p> <p>Byte 22 – Domain limiting scope<br/>           [0:3] – Domain ID (Identifies the domain that this Node Manager policy applies to)<br/>           =00h – Entire platform.<br/>           =01h – CPU subsystem.<br/>           =02h – Memory subsystem.<br/>           Others – Reserved.<br/>           [4:6] – Reserved. Write as 000b.<br/>           [7] – Limiting based on<br/>           =0 – Wall input power – PSU input power.<br/>           =1 – DC power – PSU output power or bladed system or direct DC reading from the CPU.</p> |   |
| CAh                    | Get Node Manager Version | <p><b>Request</b></p> <p>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.</p>   | <p>Get Node Manager firmware version.</p> <p>Major Firmware revision and Minor Firmware revision unambiguously identify firmware release. For every release, at least one of these numbers changes.</p> |
|                        |                          | <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>           =00h – Success (Remaining standard Completion Codes are shown in section 2.11).</p> <p>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 5 – Node Manager version<br/>           =01h – Supported Node Manager 1.0 – one power policy.<br/>           =02h – Supported Node Manager 1.5 – multiple policies and thermal triggers for power policy.<br/>           =03h – Supported Node Manager 2.0 – multiple policies and thermal triggers for power policy.<br/>           =04h – Supported Node Manager 2.5<br/>           =05h – FFh – Reserved for future use.</p> <p>Byte 6 – IPMI interface version<br/>           =01h – Node Manager IPMI version 1.0.<br/>           =02h – Node Manager IPMI version 2.0 (version defined in this document).</p> <p>Byte 7 – Patch version (binary encoded).</p>   |   |

| Net Function = 2Eh-2Fh |                                   |  |   |
|------------------------|-----------------------------------|--|---|
| Code                   | Command                           | Request, Response Data   | Description   |
|                        |                                   | <p><b>Note</b> – Change on this byte does not impact IPMI interface (Byte 6) nor Node Manager version (Byte 5). Should be set to 0h if patch version is not used by the firmware.</p> <p>Byte 8 – Major Firmware revision (binary encoded) – identifies current build of the code –and should contain the same value as the "Get Device Id" command response byte 4 – Major firmware revision.</p> <p><b>Note</b> – Change on this byte does not impact IPMI interface (Byte 6) nor Node Manager version (Byte 5).</p> <p>Byte 9 – Minor Firmware revision (BCD encoded) – identifies current build of the code and should contain the same value as the "Get Device Id" command response byte 5 Minor firmware revision.</p> <p><b>Note</b> – Change on this byte does not impact IPMI interface (Byte 6) nor Node Manager version (Byte 5).</p>  |   |
| CBh                    | Set Node Manager Power Draw Range | <p><b>Request</b></p> <p>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 4 – Domain ID<br/>[0:3] – Domain ID (Identifies the domain that this setting applies to)<br/>=00h – Entire platform.<br/>=01h – CPU subsystem.<br/>=02h – Memory subsystem.<br/>Others – Reserved.<br/>[4:7] – Reserved. Write as 0000b.</p> <p>Byte 5:6 – Minimum Power Draw in [Watts]. For domain 00h, if set to 0 the minimum power draw value will be invalidated and no validation of policy parameters against minimum power consumption will be performed. For domain 01h and 02h, if set to zero, minimum power draw will be obtained from CPU via PECI. This is an unsigned integer value.</p> <p>Byte 7:8 – Maximum Power Draw in [Watts]. For domain 00h, if set to 0 the maximum power draw value will be invalidated and no validation of policy parameters against maximum power consumption will be performed. For domain 01h and 02h, if set to zero, minimum power draw will be obtained from CPU via PECI. This is an unsigned integer value.</p> | <p>Set the Min/Max power consumption ranges. This information is preserved in the persistent storage.</p> <p>After receiving the request NM validates whether there are any policies with limit not fitting into the new power consumption range. If NM detects such policies, it sends NM Health Event with Policy Misconfiguration flag set. Additionally Node Manager disables all the policies with power limit below the Minimum Power Draw. The policies are disabled permanently – NM does not enable them until it receives a Set Node Manager Policy IPMI request.</p> <p>The same action is taken, when Node Manager receives a new power draw range from CPU via PECI at POST.</p> |
|                        |                                   | <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>=81h – Invalid Domain ID.</p> <p>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.</p>   |   |



| Net Function = 2Eh-2Fh |                                    |   |   |
|------------------------|------------------------------------|---|---|
| Code                   | Command                            | Request, Response Data  | Description   |
| CEh                    | Set Node Manager Alert Destination | <b>Request</b><br>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.<br>Byte 4 – Channel number<br>[0:3] – BMC channel number over which to send the alert from BMC to management console. Alerts can be sent to only one console.<br>[4:6] – Reserved. Write as 000b.<br>[7] – Destination information operation<br>=0 – Register alert receiver.<br>=1 – Unregister alert receiver. Use this bit to invalidate the current destination configuration. Alerts will be blocked.<br>Byte 5 – Destination Information<br><b>For channel medium – IPMB</b><br>[0] – reserved. Write as 0b.<br>[1:7] – 7-bit I2C Slave Address.<br><b>For channel medium - 802.3 LAN</b><br>Destination Selector/ Operation<br>[0:3] – Destination selector. Selects which alert destination should go to<br>=0h – Use volatile destination info.<br>=1h – Fh – Use non-volatile destination info.<br>Destination Selector definition is the same as in the “Set/Get LAN Configuration Parameters” command.<br>[4:7] – Reserved. Write as 0000b.<br>Byte 6 – Alert String Selector. Selects which Alert String, if any, to use with the alert.<br>[0:6] – String selector.<br>=00h – Use volatile Alert String.<br>=01h – 7Fh – Use non-volatile string selector.<br>Alert String Selector definition is the same as in the “Set/Get PEF Configuration Parameters” command.<br>[7]<br>=0b – Do not send Alert String.<br>=1b – Send Alert String identified by following string selector. | Provide alert destination information for Node Manager to send direct alerts that bypass the BMC SEL.<br>Destination Selector/Operation and Alert String Selector fields correspond to the associated LAN configuration parameters applicable to the BMC channel number over which to send the alert.<br>Note that the NM will determine the channel medium type in order to resolve the destination. If NM is implemented in an entity separate from the BMC, then this is done by querying the BMC using the Get Channel Info IPMI command. |
|                        |                                    | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.  |   |
| CFh                    | Get Node Manager Alert Destination | <b>Request</b><br>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.  | Provides alert destination information that is used to send alerts from for Node Manager.   |
|                        |                                    | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.<br>Byte 5 – Channel number<br>[0:3] – BMC channel number over which alert from BMC to  |   |

| Net Function = 2Eh-2Fh |         |   |             |
|------------------------|---------|---|-------------|
| Code                   | Command | Request, Response Data  | Description |
|                        |         | <p>management console will be sent.<br/> [4:6] – Reserved. Write as 000b.<br/> [7] – Destination information operation<br/> =0 – Configuration valid. Alert receiver registered.<br/> =1 – Configuration invalid. Alert receiver not registered.<br/> Alerts are blocked.</p> <p>Byte 6 – Destination Selector/ Operation<br/> [0:3] – Destination selector. Selects which alert destination should go to.<br/> =0h – Use volatile destination info.<br/> =1h – Fh – Use non-volatile destination info.<br/> [4:7] – reserved. Write as 00000b.</p> <p>Byte 7 – Alert String Selector. Selects which Alert String, if any, to use with the alert.<br/> [0:6] – String selector<br/> =00h – Use volatile Alert String.<br/> =01h – 7Fh – Use non-volatile string selector.<br/> [7]<br/> =0b – Do not send Alert String.<br/> =1b – Send Alert String identified by following string selector.</p> |             |

### 3.1.1 Get Node Manager Statistics Current Value Field

For Host Communication Failure statistics if Host Communication Failure is detected, the field contains the time in 1/100 of second between the moment NM switched to Host Communication Failure and the moment of sending the response. If OSPM responds to NM requests the field contains 0.

For Host Response Time statistics the field contains the recently measured time in 1/100 of second between Node Manager requesting P/T-state change and ASL code acknowledging the change.

For CPU throttling and memory throttling statistics the field contains the CPU or memory throttling level applied by Node Manager at the moment of receiving the command. The throttling level is expressed in %. 100% means that the CPU/memory is throttled to maximum level that Node Manager can enforce. 0% means that the CPU/memory is not throttled at all (i.e. they are running on maximum performance level).

For Host Unhandled Requests statistics the field contains the number of P/T/PUR-state change requests not handled properly by the OSPM.

### 3.1.2 Get Node Manager Statistics Minimum Value Field

For Host Communication Failure statistics, the field contains the shortest time during which NM stayed in Host Communication Failure mode. It does not include the time indicated in Current field. The time is expressed in 1/100 of second.

For Host Response Time statistics, the field contains the minimum time measured between Node Manager requesting P/T-state change and ASL code acknowledging the change. The time is expressed in 1/100 of second.

For CPU throttling and memory throttling statistics the field contains the minimum CPU or memory throttling level applied by Node Manager after the last reset of statistics. The throttling level is expressed in %. 100% means that the CPU/memory is throttled to maximum level that Node Manager can enforce. 0% means that the CPU/memory is not throttled at all (i.e. they are running on maximum performance level).



For Host Unhandled Requests statistics the field always contains zero.

### 3.1.3 Get Node Manager Statistics Maximum Value Field

For Host Communication Failure statistics, the field contains the longest time during which NM stayed in Host Communication Failure mode. It does not include the time indicated in Current field. The time is expressed in 1/100 of second.

For Host Response Time statistics, the field contains the maximum time measured between Node Manager requesting P/T-state change and ASL code acknowledging the change. The time is expressed in 1/100 of second.

For CPU throttling and memory throttling statistics the field contains the maximum CPU or memory throttling level applied by Node Manager after the last reset of statistics. The throttling level is expressed in %. 100% means that the CPU/memory is throttled to maximum level that Node Manager can enforce. 0% means that the CPU/memory is not throttled at all (i.e. they are running on maximum performance level).

For Host Unhandled Requests statistics the field contains the number of P/T/PUR-state change requests not handled properly by the OSPM.

### 3.1.4 Get Node Manager Statistics Average Value Field

For Host Communication Failure statistics, the field contains the percentage of time NM worked in Host Communication Failure mode from the recent statistics reset.

For Host Response Time statistics, the field contains the arithmetic average of all the Host response time measurements. The time is expressed in 1/100 of second.

For CPU throttling and memory throttling statistics the field contains the average CPU or memory throttling level applied by Node Manager after the last reset of statistics. The average is calculated as arithmetic average after last statistics reset. The throttling level is expressed in %. 100% means that the CPU/memory is throttled to maximum level that Node Manager can enforce. 0% means that the CPU/memory is not throttled at all (i.e. they are running on maximum performance level).

For Host Unhandled Requests statistics the field always contains zero.

## 3.2 Local Platform Node Manager Configuration and Control Commands

The following commands should not be exposed to the external software. Only BMC may use the following commands.

Note that all the below commands are not supported when Node Manager Feature Enabled is set to 'false' using Flash Image Tool.

| Net Function = 2Eh-2Fh |                        |  |   |
|------------------------|------------------------|--|---|
| Code                   | Command                | Request, Response Data   | Description   |
| D0h                    | Set Total Power Budget | <b>Request</b><br>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.<br>Byte 4 – Domain ID<br>[0:3] – Domain ID (Identifies the domain that this setting applies to)<br>=00h – Entire platform.<br>=01h – CPU subsystem. | Set total power budget for the CPUs. This command is optional and may be unavailable on certain implementations. This command controls the platform power limit using |

| Net Function = 2Eh-2Fh |                                     |   |   |
|------------------------|-------------------------------------|---|---|
| Code                   | Command                             | Request, Response Data  | Description   |
|                        |                                     | <p>=02h – Memory subsystem.<br/>Others – Reserved.<br/>[4:7] – Reserved. Write as 0000b.</p> <p>Byte 5:6 – Target power budget in [Watts] that should be maintained by the Power Budget Control Service.</p>  | <p>aggressive settings.</p> <p><b>Note</b> – “Set Total Power Budget” function is only accessible if the Node Manager policy control feature is disabled.</p>   |
|                        |                                     | <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>=81h – Invalid Domain ID.<br/>=84h – Power Budget out of range.</p> <p>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p><b>Note</b> – When the Node Manager policy control is enabled globally (see C0h command) the command should respond with Completion Code D5h – Cannot execute command – command not supported in present state.</p>   |   |
| D2h                    | Set Max Allowed CPU P-state/T-state | <p><b>Request</b></p> <p>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 4 – Domain ID<br/>[0:3] – Domain ID – Identifies the domain that this Node Manager setting applies to<br/>=00h – Entire platform – for compatibility with previous NM versions P/T state settings are applied to CPU subsystem.<br/>Others – Reserved.<br/>[4:5] – Control Knob<br/>=00b – Set max allowed CPU P-state/T-state.<br/>=01b – Set max allowed logical processors.<br/>=10b – Reserved.<br/>=11b – Reserved.<br/>[6:7] – Reserved. Write as 00b.</p> <p>For Control Knob set to 00b:<br/>Byte 5 – P-state number to be set.<br/>Byte 6 – T-state number to be set.</p> <p>For Control Knob set to 01b:<br/>Byte 5:6 – Set max allowed logical processors.</p> <p><b>Note</b> – If any of the fields is set to FFh, it should be omitted when setting the value.</p> <p><b>Note</b> – This setting is volatile. It is cleared after reset of either host or ME side.</p> | <p>This command is optional and may be unavailable on certain implementations.</p> <p>This command imposes additional limit on the host side apart of the limit that may be applied by NM Policy Control, or Total Power Budget. If the limit applied by Policy Control or Power Budget is lower than the effect of P-state/T-state limit then that power limit is used. If P-state/T-state limit imposes lower power limit than the P-state/T-state limit is kept.</p> |
|                        |                                     | <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>=81h – Invalid Domain ID.<br/>=8Ah – P-state or T-state out of range.</p> <p>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.</p>  |   |



| Net Function = 2Eh-2Fh |                                     |  |   |
|------------------------|-------------------------------------|--|---|
| Code                   | Command                             | Request, Response Data   | Description   |
| D1h                    | Get Total Power Budget              | <b>Request</b><br>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.<br>Byte 4 – Domain ID<br>[0:3] – Domain ID (Identifies the domain that this setting applies to)<br>00h – Entire platform.<br>01h – CPU subsystem.<br>02h – Memory subsystem.<br>Others – Reserved.<br>[4:7] – Reserved. Write as 0000b.   | This command is optional and may be unavailable on certain implementations. |
|                        |                                     | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br>=81h – Invalid Domain ID.<br>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.<br>Byte 5:6 – Target power budget in [Watts] that should be maintained by the Power Budget Control Service.  |   |
| D3h                    | Get Max Allowed CPU P-state/T-state | <b>Request</b><br>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.<br>Byte 4 – Domain ID<br>[0:3] – Domain ID – Identifies the domain that this Node Manager setting applies to<br>=00h – Entire platform – for compatibility with previous NM versions P/T state settings are applied to CPU subsystem.<br>Others – Reserved.<br>[4:5] – Control Knob<br>=00b – get max allowed CPU P-state/T-state.=01b – get max allowed logical processors.<br>=10b – Reserved.<br>=11b – Reserved.<br>[6:7] – Reserved. Write as 00b.   | This command is optional and may be unavailable on certain implementations. |
|                        |                                     | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br>=81h – Invalid Domain ID.<br>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.<br>For Control Knob set to 00b:<br>Byte 5 – Current maximum P-state.<br>Byte 6 – Current maximum T-state.<br>For Control Knob set to 01b:<br>Byte 5:6 – Total requested by ME number of allowed logical processors on a system. This is a number requested by ME and OSPM is not required to fulfill this request.<br><b>Note</b> – If any of the fields is set to FFh, it means that value is unavailable.<br><b>Note</b> – This command returns value recently confirmed by the host side. For a short while it may be differ from the |   |

| Net Function = 2Eh-2Fh |                                 |   |   |
|------------------------|---------------------------------|---|---|
| Code                   | Command                         | Request, Response Data  | Description   |
|                        |                                 | value set with D2h. If the difference persists it indicates an issue on the host side.  |   |
| D4h                    | Get Number Of P-states/T-states | <b>Request</b><br>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.<br>Byte 4 – Domain ID<br>[0:3] – Domain ID – Identifies the domain that this setting applies to<br>=00h – Entire platform – for compatibility with previous NM versions state settings are applied to CPU subsystem.<br>Others – Reserved.<br>[4:5] – Control Knob<br>=00b – Max allowed processor P-states/T-states.<br>=01b – Max allowed logical processors.<br>=10b – Reserved.<br>=11b – Reserved.<br>[6:7] – Reserved. Write as 00b.   | This command is optional and may be unavailable on certain implementations.   |
|                        |                                 | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br>=81h – Invalid Domain ID.<br>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.<br>For Control Knob set to 00b:<br>Byte 5 – Number of P-states available on the platform. This number will be always be 1 or more, even if BIOS will pass information that 0 P-states are supported.<br>Byte 6 – Current Number of T-states available on the platform. This number will be always 1 or more, even if BIOS will pass information that T-states are supported.<br>For Control Knob set to 01b:<br>Byte 5:6 – Number of logical processors on the platform.                |   |
| D7h                    | Set PSU Configuration           | <b>Request</b><br>Byte 1:3 – Intel manufacturer ID – 0x000157, LS byte first.<br>Byte 4 – Domain ID<br>[0:3] – Domain ID (Identifies Domain which uses the defined PSU set). Currently, FW supports only one domain – Domain 0).<br>[4:7] – Reserved. Write as 0000b.<br>Byte 5 – PMBUS PSU address 1. Node Manager will monitor the presence of the defined PSU.<br>[0] – PSU mode<br>=1 – PSU is installed and lack of power readings should be reported to Management Console.<br>=0 (default) – PSU is installed or may be attached in the future.<br>[1:7] – 7-bit PSU SMBUS address. Set to 00h if address is not used.<br>Byte 6:12 – PMBUS PSU address 2 to PMBUS PSU address 8 encoded as in the Byte 5. | This command may override the supported set of PSU's by defining a set of all supported PSU's. Only the PSU SMBUS addresses are stored in the persistent storage.<br><br>This command should be send to Node Manager by BMC if the lack of reading from the defined PSU should be reported to the Management Console using node Manager Health Event. Otherwise, Node Manager will send a notification only if all PSUs will disappear and there will be no power readings available. |
|                        |                                 | <b>Response</b>   |   |





| Net Function = 2Eh-2Fh |                        |   |  |
|------------------------|------------------------|---|--|
| Code                   | Command                | Request, Response Data  | Description  |
|                        |                        | <p>Byte 1 – Completion Code<br/>           =00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>           =81h – Invalid Domain ID.</p> <p>Byte 2:4 – Intel manufacturer ID – 0x000157, LS byte first.</p>  |  |
| D8h                    | Get PSU Configuration  | <p><b>Request</b></p> <p>Byte 2:4 – Intel manufacturer ID – 0x000157, LS byte first.</p> <p>Byte 4 – Domain ID<br/>           [0:3] – Domain Id (Identifies Domain which uses the defined PSU set). Currently, FW supports only one domain – Domain 0).<br/>           [4:7] – Reserved. Write as 0000b.</p> <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>           =00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>           =81h – Invalid Domain ID.</p> <p>Byte 2:4 – Intel manufacturer ID – 0x000157, LS byte first.</p> <p>Byte 5 – Domain ID<br/>           [0:3] – Domain Id (Identifies Domain which uses the defined PSU set). Currently, FW supports only one domain – Domain 0).<br/>           [4:7] – Reserved. Return as 0000b.</p> <p>Byte 6 – PMBUS PSU address 1. Node Manager will monitor the presence of the defined PSU<br/>           [7:1] – 7-bit PSU SMBUS address. Set to 00h if address is not used.<br/>           [0] – PSU mode<br/>           =1 – PSU is installed and lack of power readings should be reported to Management Console.<br/>           =0 (default) – PSU is installed or may be attached in the future.</p> <p>Bytes 7:13 – PMBUS PSU address 2 to PMBUS PSU address 8 encoded as in the Byte 6.</p> |  |
| D9h                    | Send Raw PMBUS command | <p><b>Request</b></p> <p>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 4 – Flags<br/>           [7] = 1b – Enable PEC.<br/>           [6] = 1b – Do not report PEC errors in Completion Code. If the bit is set, NM firmware does not return “bad PEC” Completion Codes. In this case BMC can learn that the transaction failed by checking PEC in the PMBUS response message. The flag does not disable PMBUS command retries.<br/>           [5:4] – Reserved. Write as 00b.<br/>           [3:1] – SMBUS message transaction type<br/>           =0h – SEND_BYTE.<br/>           =1h – READ_BYTE.<br/>           =2h – WRITE_BYTE.<br/>           =3h – READ_WORD.<br/>           =4h – WRITE_WORD.<br/>           =5h – BLOCK_READ.<br/>           =6h – BLOCK_WRITE.</p>  | This command sends one PMBUS command to the specified address. Address is validated against factory presets. |

| Net Function = 2Eh-2Fh |                   |  |  |
|------------------------|-------------------|--|--|
| Code                   | Command           | Request, Response Data   | Description  |
|                        |                   | <p>=7h – BLOCK_WRITE_READ_PROC_CALL.<br/>[0] – Reserved. Write as 0b.</p> <p>Byte 5 – Target PSU Address<br/>[7:1] – 7-bit PSU SMBUS address.<br/>[0] – Reserved. Write as 0b.</p> <p>Byte 6 – Write Length.</p> <p>Byte 7 – Read Length. This field is used to validate if the slave returns proper number of bytes. Read Length does not include PEC byte.</p> <p>Byte 8:M – PMBUS command.</p> <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>=80h – Command response timeout. SMBUS device was not present.<br/>=81h – Command not serviced. Not able to allocate the resources for serving this command at this time. Retry needed.<br/>=82h – Command not executed due to conflict with PSU Optimization feature.<br/>=A1h – Illegal SMBUS PSU Slave Target Address.<br/>=A2h – PEC error.<br/>=A3h – Number of bytes returned by the Slave different from Read Length see byte 7 of request.<br/>=A5h – Unsupported Write Length.<br/>=A6h – Unsupported Read Length.<br/>=AAh – SMBUS timeout.</p> <p>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 5:N – PMBUS response data received from PSU during Read transaction phase. Response from the slave is returned for Completion Codes: 00h, A2h, A3h.</p> |  |
| EAh                    | Get Host CPU data | <p><b>Request</b></p> <p>Byte 1:3 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 4 – Domain ID<br/>[0:3] – Domain ID (Identifies the set of processors supported by the domain. Currently FW supports only one domain – Domain 0).<br/>[4:7] – Reserved. Write as 0000b.</p> <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>=81h – Invalid Domain ID.</p> <p>Byte 2:4 – Intel Manufacturer ID – 000157h, LS byte first.</p> <p>Byte 5 – Host CPU data<br/>[7] – Set to 1 if End of POST notification was received.<br/>[6:5] – Reserved. Write as 00b.<br/>[4] – Set to 1 if Host CPU discovery data provided with that command is valid.<br/>[3] – Set to 1 if Node Manager already activated regular power limiting policies after Host startup.<br/>[2:0] – Reserved. Write as 000b.</p>   | The command returns the CPU configuration data passed from BIOS to ME using HECI messages. |



| Net Function = 2Eh-2Fh |         |  |             |
|------------------------|---------|--|-------------|
| Code                   | Command | Request, Response Data   | Description |
|                        |         | <p><b>Note</b> – Bytes 6:25 are ignored if Byte 5 bit [4] is set to 0. If Byte 5 bit [4] is set to 1 Bytes 6:25 should describe the actual Host CPU data of the platform. Additionally bytes 6:24 should be set to 0 if the CPU discovery data is passed to NM directly by the BIOS. Per processor discovery data will be provided only for the lowest number processor that is installed. In the multiprocessor environment all other processors installed on board should match the number of performance states and each processor performance state must have identical performance and power consumption parameters.</p> <p>Byte 6 – Number of P-states supported by the current platform CPU configuration<br/>           =0 – If P-states are disabled by the user.<br/>           =1 – If CPU does not support more P-states or in the multiprocessor environment some processors installed on board do not match the lowest number processor power consumption parameters.<br/>           = 2 – 255 – Actual number of supported P-states by the lowest number processor.</p> <p><b>Note</b> – Other processors should match the number of performance states of lowest number processor.</p> <p>Byte 7 – Number of T-states supported by the current platform CPU configuration<br/>           =0 – If T-states are disabled by the user.<br/>           =1 – 255 – Actual number of supported T-states by the lowest number processor.</p> <p><b>Note</b> – Other processors should match the number of throttling states of lowest number processor.</p> <p>Byte 8 – Number of installed processor packages. This value is calculated as a number of all sockets with CPU package present.</p> <p>Bytes 9:16 – Processor Discovery Data for the lowest number processor in LSByte-first order. Turbo power current Limit MSR 1ACh for the lowest number processor passed by BIOS.</p> <p>Bytes 17:24 – Processor Discovery Data 2 for the lowest number processor in LSByte-first order. Platform Info MSR 0Ceh for the lowest number processor passed by BIOS.</p> <p>Byte 25 – Reserved. Write as 00000000b.</p> |             |

### 3.3 External DCMI Power Management Commands

The standard DCMI Power Management commands defined in [DCMI] offer functionality that is a subset of the Intel® Intelligent Power Technology Node Manager (see 3.1). In DM and DNM firmware, and on Romley or Denlow platforms also in NM firmware the following commands are supported:

| Net Function = DCGRP (0x2C) |                          |  |  |
|-----------------------------|--------------------------|--|--|
| Code                        | Command                  | Request, Response Data   | Description  |
| 01h                         | Get DCMI Capability Info | <b>Request</b><br>Byte 1 – Group Extension Identification = DCh.<br>Byte 2 – Parameter Selector<br>=5 – Enhanced Power Statistics attributes.<br>Other parameter selector values are not supported.  | This command is intended for BMC or Remote Console to provide information about DCMI Power Manager capabilities. |
|                             |                          | <b>Response</b><br>Byte 1 – Completion Code<br>=C1h – Returned if DCMI mode is not present.<br>Byte 2 – Group Extension Identification =DCh.<br>Byte 3:4 – DCMI Specification Conformance<br>Byte 3 – Major Version =1.<br>Byte 4 – Minor Version = 5.<br>Byte 5 – Parameter Revision<br>=02h – Enhanced Power Statistics attributes.<br>Byte 6 – The number of supported rolling average time periods = 09h.<br>Byte 7:15 – Rolling Average Time periods<br>[7:6] – Time duration units<br>=00b – Seconds.<br>=01b – Minutes.<br>=10b – Hours.<br>=11b – Days.<br>[5:0] – Time duration<br>The following periods are supported<br>=05h – 5 seconds.<br>=0Fh – 15 seconds.<br>=1Eh – 30 seconds.<br>=41h – 1 minute.<br>=43h – 3 minutes.<br>=47h – 7 minutes.<br>=4Fh – 15 minutes.<br>=5Eh – 30 minutes.<br>=81h – 1 hour. |  |
| 02h                         | Get Power Reading        | <b>Request</b><br>Byte 1 – Group Extension Identification = DCh.<br>Byte 2 – Mode<br>=1 – System Power Statistics.<br>=2 – Enhanced System Power Statistics.<br>Byte 3 – Rolling Average Time periods<br>For Mode = 1 – Reserved. Write 00h.<br>For Mode = 2 – One of periods reported in bytes 7:14 of the response to the Get DCMI Capability Info command.<br>Byte 4 – Reserved.  |  |



| Net Function = DCGRP (0x2C) |                 |  |             |
|-----------------------------|-----------------|--|-------------|
| Code                        | Command         | Request, Response Data   | Description |
|                             |                 | <b>Response</b><br>Byte 1 – Completion Code<br>=C1h – Returned if DCMI mode is not present.<br>Byte 2 – Group Extension Identification = DCh.<br>Byte 3:4 – Current Power in [Watts].<br>Byte 5:6 – Minimum Power over sampling duration in [Watts].<br>Byte 7:8 – Maximum Power over sampling duration in [Watts].<br>Byte 9:10 – Average Power over sampling duration in [Watts].<br>Byte 11:14 – IPMI Specification based Time Stamp based on SEL. For Mode = 2 – The time stamp specifies the end of the averaging window.<br>Byte 15:18 – Statistics reporting time period.<br>For Mode = 1 – Timeframe in milliseconds, over which the controller collects statistics.<br>For Mode = 2 – Timeframe reflects the Averaging Time period in units.<br>Byte 19 – Power Reading State<br>[0:5] – Reserved.<br>[6]<br>=1b – Power Measurement active.<br>=0b – No Power Measurement is available.<br>[7] – Reserved.<br><b>Note</b> – Sampling duration depends on Mode selection. For Mode 1 global system statistics collected over system up time are returned. |             |
| 03h                         | Get Power Limit | <b>Request</b><br>Byte 1 – Group Extension Identification = DCh.<br>Byte 2:3 – Reserved for future use. Write 0000h.   |             |

| Net Function = DCGRP (0x2C) |                 |   |   |
|-----------------------------|-----------------|---|---|
| Code                        | Command         | Request, Response Data  | Description   |
|                             |                 | <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>           =00h – Power Limit Active.<br/>           =80h – No Set Power Limit.<br/>           =C1h – Returned if DCMI mode is not present.</p> <p>Byte 2 – Group Extension Identification = DCh.</p> <p>Byte 3:4 – Reserved for future use.</p> <p>Byte 5 – Exception actions. Actions taken if the power limit is exceeded and cannot be controlled within the correction time limit.<br/>           =00h – No action<br/>           =01h –Hard Power Off system and log event to SEL.<br/>           =11h – Log event to SEL.</p> <p>Byte 6:7 – Power Limit Requested in [Watts].</p> <p>Byte 8:11 – Correction time limit in milliseconds. Maximum time taken to limit the power, otherwise exception action will be taken as configured.</p> <p>Byte 12:13 – Reserved for future use.</p> <p>Byte 14:15 – Management application Statistics Sampling period in seconds.</p>   |   |
| 04h                         | Set Power Limit | <p><b>Request</b></p> <p>Byte 1 – Group Extension Identification = DCh.</p> <p>Byte 2:4 – Reserved for future use.</p> <p>Byte 5 – Exception actions. Actions taken if the power limit exceeded and cannot be controlled within the correction time limit.<br/>           =00h – No Action.<br/>           =01h –Hard Power Off system and log event to SEL.<br/>           =11h – Log event to SEL.</p> <p>Byte 6:7 – Power Limit Requested in [Watts].</p> <p>Byte 8:11 – Correction time limit in milliseconds. Maximum time taken to limit the power, otherwise exception action will be taken as configured. The Exception Action shall be taken if the system power usage constantly exceeds the specified power limit for more than the Correction Time Limit interval. The Correction Time Limit timeout automatically restarts if the system power meets or drops below the Power Limit.</p> <p>Byte 12:13 – Reserved for future use.</p> <p>Byte 14:15 – Management application Statistics Sampling period in seconds.</p> <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>           =84h – Power Limit out of range.<br/>           =85h – Correction Time out of range.<br/>           =89h – Statistics Reporting Period out of range.<br/>           =C1h – Returned if DCMI mode is not present.</p> <p>Byte 2 – Group Extension Identification = DCh.</p> | The following defaults are used by node Manager during policy creation (see command "Set Node Manager Policy"):<br>Domain=0<br>Is Dcml=TRUE<br>Policy Trigger Type=0<br>Aggressive Cpu Power<br>Correction=0<br>Trigger Limit=0 |



| Net Function = DCGRP (0x2C) |                                 |   |             |
|-----------------------------|---------------------------------|---|-------------|
| Code                        | Command                         | Request, Response Data  | Description |
| 05h                         | Activate/Deactivate Power Limit | <b>Request</b><br>Byte 1 – Group Extension Identification = DCh.<br>Byte 2 – Power Limit Activation<br>=00h – Deactivate Power Limit.<br>=01h – Activate Power Limit.<br>Byte 3:4 – Reserved. |             |
|                             |                                 | <b>Response</b><br>Byte 1 – Completion Code<br>=C1h – Returned if DCMI mode is not present.<br>Byte 2 – Group Extension Identification = DCh.   |             |

### 3.4 IPMI Sensors Implemented by Node Manager FW

The below table summarizes the sensors exposed by Node Manager FW.

| Event  | Sensor Type        | Event Dir                        | Event Type                   | Alert Immediate<br>(See section 4.5) |
|--|--------------------|----------------------------------|------------------------------|--------------------------------------|
| Node Manager Exception Event                       | DCh – OEM          | 0 – assertion                    | 72h – OEM                    | No                                   |
| Node Manager Health Event                          | DCh – OEM          | 0 – assertion                    | 73h – OEM                    | Yes                                  |
| Node Manager Operational Capabilities Change Event | DCh – OEM          | 0 – assertion<br>1 – deassertion | 74h – OEM                    | Yes                                  |
| Node manager Alert Threshold Exceeded              | DCh – OEM          | 0 – assertion<br>1 – deassertion | 72h – OEM                    | Yes                                  |
| Typical Platform power consumption in Sx state     | DCh – OEM          | N/A                              | N/A                          | N/A                                  |
| PSU Status   | 08h – PSU          | 0 – assertion<br>1 – deassertion | 6Fh – Sensor Spec            | No                                   |
| DCMI Power Threshold Event                         | 09h – Power Unit   | 0 – assertion<br>1 – deassertion | 05h – Generic – limit status | No                                   |
| DCMI Power Off Event                               | 12h – System Event | 0 – assertion                    | 0Ah – Generic - system state | No                                   |

*Firmware SKU Availability* column specifies whether the sensor is supported only in some SKU:

- A – sensor available in all FW SKUs,
- NM – sensor available only when Node Manager Feature Enabled parameter is set to 'true' using Flash Image Tool,

- PECI – sensor available only when PECI is attached to the PCH.

*Reading Availability* column specifies when the sensor reading is available:

- A – always when ME is On,
- H – when HOST CPU is On,
- O – after reception of END\_OF\_POST notification,
- E – No reading available (Event Only).

*Defaults Configurable in FIT* column defines whether the default configuration of the sensors can be set using Flash Image Tool. The default configuration includes:

- Thresholds
- Event Enable Mask
- Scanning Periods
- Scanning Enable Flag
- Per-sensor Event Enable Flag

The default configuration is applied by NM Firmware at first management engine startup after G3 condition on Global Platform Reset (see definition in PCH EDS).

| Sensor Number | Description                                       | Firmware SKU Availability | Reading Availability | Defaults Configurable in Flash Image Tool | Notes   |
|---------------|---|---------------------------|----------------------|---|---|
| 21 T          | Typical Platform power consumption in Sx state    | NM/DNM/DM                 | A                    | Yes                                       | OEM sensor that presents Typical Power consumption in Sx state. The value of this sensor is constant and may only change upon HW configuration change. Node Manager uses this data to estimate the power consumed in Sx state.  |
| 24 D          | Dynamic Power Node Manager Exception event Sensor | NM/DNM                    | E                    | No  | <p>OEM Event only sensor.</p> <p>Event will be sent each time when maintained policy power limit is exceeded over Correction Time Limit.</p> <p>Events are sent no faster than every 300 milliseconds.</p> <p>First occurrence of not acknowledged event will be retransmitted according to IPMB specification after ~230 milliseconds.</p> <p>“Command illegal for specified sensor or record type (CDh)” error code is returned in response to the following commands: Get Sensor Reading, Set/Get Sensor Thresholds,</p> |





| Sensor Number | Description  | Firmware SKU Availability | Reading Availability | Defaults Configurable in Flash Image Tool | Notes   |
|---------------|--|---------------------------|----------------------|---|---|
|               |  |                           |                      |   | Re-Arm Sensor Events, Set/Get Sensor Event Enable.  |
| 25 D          | Dynamic Power Node Manager Health event Sensor             | NM/DNM                    | E                    | No  | <p>OEM Event only sensor used to send events about integrity of Node Manager policy or necessary sensor readings.</p> <p>Events are sent no faster than every 300 milliseconds.</p> <p>First occurrence of not acknowledged event will be retransmitted according to IPMB specification after ~230 milliseconds.</p> <p>“Command illegal for specified sensor or record type (CDh)” error code is returned in response to the following commands: Get Sensor Reading, Set/Get Sensor Thresholds, Re-Arm Sensor Events, Set/Get Sensor Event Enable.</p>   |
| 26 D          | Dynamic Power Node Manager Operational Capabilities sensor | NM/DNM                    | A                    | No  | <p>OEM sensor, that value will indicate the operational capabilities of the sensor. Whenever the sensor value changes, an immediate alert is also sent. Please see the event description for the description of the values of the sensor.</p> <p>“Command illegal for specified sensor or record type (CDh)” error code is returned in response to the following commands: Set/Get Sensor Thresholds.</p> <p>Events are sent no faster than every 300 milliseconds.</p> <p>Current value of not acknowledged capabilities sensor will be retransmitted according to IPMB specification after ~230 milliseconds.</p> |
| 27            | Node manager Alert Threshold Exceeded sensor               | NM/DNM                    | E                    | No  | <p>OEM Event only sensor used to send events when Node Manager detects that a specified alert threshold for one of the policies is exceeded.</p> <p>Events are sent no faster</p>   |

| Sensor Number | Description            | Firmware SKU Availability | Reading Availability | Defaults Configurable in Flash Image Tool | Notes   |
|---------------|------------------------|---------------------------|----------------------|---|---|
|               |                        |                           |                      |   | <p>than every 300 milliseconds.</p> <p>First occurrence of Threshold exceeded event assertion/deassertion will be retransmitted <a href="#">according to IPMB specification</a> after ~230 milliseconds.</p> <p>"Command illegal for specified sensor or record type (CDh)" error code is returned in response to the following commands: Get Sensor Reading, Set/Get Sensor Thresholds, Re-Arm Sensor Events, Set/Get Sensor Event Enable.</p> |
| 58            | DCMI Power Off Sensor  | NM/DNM/DM                 | A                    | Yes                                       | Sensor is defined in [DCMI].  |
| 59 DCMI       | Power Threshold Sensor | NM/DNM/DM                 | A                    | Yes                                       | Sensor is defined in [DCMI].  |
| 102-109       | PSU#0 to PSU#7 Status  | NM/DNM/DM A               |                      | Yes                                       | <p>Discrete sensor.</p> <p>This sensor is used for getting the Power Supply status.</p>   |

### 3.4.1 Typical Platform power consumption in Sx state

This is a sensor that does not generate any events and supports only Get Sensor Reading command. The sensor is used in Platform Event messages to BMC when policy limit is exceeded. This Sensor presents Typical Power consumption in Sx state in Watts. The value of this sensor is constant and may only change upon HW configuration change. Node Manager uses this data to estimate the power consumed by the platform in Sx state.

### 3.4.2 NM Exception Event Sensor

This is an Event-Only sensor that does not support Get Sensor Reading command nor re-arm IPMI command. The sensor is used in Platform Event messages to BMC when policy limit is exceeded.

### 3.4.3 NM Alert Threshold Exceeded Sensor

This is an Event-Only sensor that does not support Get Sensor Reading command nor re-arm IPMI command. The sensor is used in Immediate Alert Event messages to remote console when policy threshold or policy limit is exceeded.

### 3.4.4 NM Health Sensor

OEM Event only sensor used to send events about integrity of Node Manager Policy or necessary sensor readings.



### 3.4.5 NM Operational Capabilities Change Sensor

OEM sensor indicating the operational capabilities of the Node Manager service. Whenever the sensor value changes, an immediate alert is also sent. Please see the event description for the description of the values of the sensor.

| NM Operational Capability | Description   | Enable Condition   |
|---------------------------|---|--|
| Policy Interface          | Node Manager accepts policy configuration change requests.                                | <ol style="list-style-type: none"><li>1. Node Manager operational image is running.</li><li>2. Flash wear-out protection mechanism did not lock access to SPI Flash.</li></ol>   |
| Monitoring Capability     | Node Manager is monitoring total system power consumption and CPU power and Memory power. | <ol style="list-style-type: none"><li>1. Platform is in S0/S1 state.</li><li>2. Node Manager did not detect power reading failure conditions.</li></ol>  |
| Power Limiting Capability | Node Manager is able to limit power consumption.  | <ol style="list-style-type: none"><li>1. Node Manager has CPU information from BIOS.</li><li>2. Total Power Monitoring Capability is enabled or Node Manager enforces non-zero CPU or memory throttling level.</li><li>3. OSPM driver is responding to requests from Node Manager.</li><li>4. PECI interface is working properly.</li></ol> <p><b>Note</b> – When Node Manager does actively limits power, either because no power limit is configured or power consumption is below all the configured limits, NM may not detect a change in this capability.</p> |

### 3.4.6 PSU Status Sensors

This sensor is used for getting the Power Supply status. There is a separate sensor per PSU. Sensors are numbered from 102 to 109.

PSU Status Sensors are created only for PSUs that have responded to ME Discovery commands. This means that no events will be sent to BMC for PSUs that were never populated, even if such PSUs were configured in the xml file. However such events will be sent for PSUs that had been removed during platform operation as once created PSU Status Sensors will remain in the SDR.

Each sensor supports the following bits:

| PSU Status Assertion bit      | Description   |
|-------------------------------|---|
| Presence detected             | Asserted if power supply module is present. Events are only logged for power supply presence upon changes in the presence status after AC power is applied (no events logged for initial state).  |
| Power supply failure detected | Asserted if power supply module has failed.<br><br>This is a "logical OR" of all the faults such as IOut OC Fault, Pout OP fault (STATUS_IOUT), OT fault (STATUS_TEMPERATURE), Fan Fault (STATUS_FANS_1_2).   |
| Predictive failure            | Asserted if a condition that is likely to lead to a power supply module failure has been detected, such as a failing fan.<br><br>This is a logical OR of all the status_xxx bits. such as Iout OC warning, Pout OP warning (STATUS_IOUT), Vin UV warning, Iin over current warning, Pin over power warning (STATUS_INPUT), OT warning (STATUS_TEMPERATURE), Fan 1 warning, Fan 2 warning (STATUS_FANS_1_2). |
| AC lost                       | Asserted if there is no AC power input to a power supply module.  |

| PSU Status Assertion bit   | Description   |
|----------------------------|---|
| Configuration Error        | Asserted for not supported PSU, when other information cannot be read.  |
| PSU optimization triggered | <p>Asserted if PSU optimization feature is throttling the total system power to prevent the system failure. PSU optimization could be activated by one of the following warnings Iout OC warning (STATUS_IOUT), Vin UV warning (STATUS_INPUT) or OT warning (STATUS_TEMPERATURE).</p> <p>"Re-arm Sensor Events" IPMI command will reset CLST events (over-current and over-temperature) in the PSU. This allows in the manual CLST mode to reset the platform throttling by the BMC when the corrective action on the platform was performed.</p> |

### 3.4.7 Event Messages Definition

| Net Function = S/E (04h) |  |  |   |
|--------------------------|--|--|---|
| Code                     | Command  | Request, Response Data   | Description   |
| 02h                      | Platform Event Message<br>Node Manager Exception Event | <b>Request</b><br>Byte 1 – EvMRev<br>=04h – IPMI2.0 format.<br>Byte 2 – Sensor Type<br>=DCh – OEM.<br>Byte 3 – Sensor Number<br>=24 – Node Manager Event Sensor.<br>Byte 4 – Event Dir   Event Type<br>[0:6] – Event Type<br>=72h – OEM.<br>[7] – Event Dir<br>=0 – Assertion Event.<br>=1 – Deassertion Event.<br>Byte 5 – Event Data 1<br>[0:2] – Reserved.<br>[3] – Node Manager Policy event<br>=0 – Reserved.<br>=1 – Policy Correction Time Exceeded – policy did not meet the contract for the defined policy. The policy will continue to limit the power or shutdown the platform based on the defined policy action.<br>[4:5] = 10b – OEM code in byte 3.<br>[6:7] = 10b – OEM code in byte 2.<br>Byte 6 – Event Data 2<br>[0:3] – Domain ID (Identifies the domain that this Node Manager policy applies to)<br>=00h – Entire platform.<br>=01h – CPU subsystem.<br>=02h – Memory subsystem.<br>Others – Reserved.<br>[4:7] – Reserved.<br>Byte 7 – Event Data 3<br>= <Policy ID> | Event will be sent each time when maintained policy power limit is exceeded over Correction Time Limit.<br>First occurrence of not acknowledged event will be retransmitted no faster than every 300 milliseconds.<br><b>Note</b> – In DNM SKU this event is logged to SEL. |
|                          |  | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).   |   |



| Net Function = S/E (04h) |  |   |   |
|--------------------------|--|---|---|
| Code                     | Command  | Request, Response Data  | Description   |
|                          |  | Others – Error.   |   |
| 16h                      | Alert Immediate Message<br>Node Manager Health Event | <p><b>Request</b></p> <p>Bytes 1:3 – As set in the “Set Node Manager Alert Destination” IPMI command.</p> <p>Byte 4 – Generator ID<br/>[7:1] – 7-bit IPMB address of the Firmware.<br/>[0] = 0b – Generator ID is IPMB address.</p> <p>Byte 5 – EvMRev<br/>=04h – IPMI2.0 format.</p> <p>Byte 6 – Sensor Type<br/>=DCh – OEM.</p> <p>Byte 7 – Sensor Number<br/>=25 – Node Manager Health sensor.</p> <p>Byte 8 – Event Dir   Event Type<br/>[0:6] – Event Type<br/>=73h – OEM.</p> <p>[7] – Event Dir<br/>=0 – Assertion Event.</p> <p>Byte 9 – Event Data 1<br/>[0:3] – Health Event Type<br/>=02h – Sensor Node Manager.<br/>[4:5] = 10b – OEM code in byte 3.<br/>[6:7] = 10b – OEM code in byte 2.</p> <p>Byte 10 – Event Data 2<br/>[0:3] – Domain ID.<br/>[4:7] – Error Type<br/>0 – 9 – Reserved.<br/>=10 – Policy Misconfiguration.<br/>=11 – Power Sensor Reading Failure.<br/>=12 – Inlet Temperature Reading Failure.<br/>=13 – Host Communication Error.<br/>=14 – Real-time clock synchronization failure. This is sent 10 minutes after NM reads invalid time from system RTC.<br/>=15 – Platform shutdown initiated by NM policy due to execution of action defined by Policy Exception Action see “Set NM Policy” command Byte 7 bit [1].<br/>16 – Reserved.</p> <p>Byte 11 – Event Data 3<br/>For Error Type = 10 or 15 – &lt;Policy ID&gt;.<br/>For Error Type = 11 &lt;PowerSensorAddress&gt;.<br/>For Error Type = 12 &lt;Inlet Sensor Address&gt;.</p> <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>Others – Error.</p> | <p>This message provides a run-time error indication about Node Manager's health.</p> <p><b>Note</b> – Misconfigured policy error can happen when the max/min power consumption set using Set Power Draw Range exceeds the values in any of the configured policy.</p> <p>Real-time clock synchronization failure alert is sent when NM is enabled and capable of limiting power, but within 10 minutes the firmware cannot obtain valid calendar time from system RTC.</p> <p>Host Communication Error is only detected when in NM there is a policy actively limiting power or when policy-based power limiting is disabled and NM enforces limit configured using Set Total Power Budget for power domain 0 or 1.</p> <p>Power Sensor Reading Failure may indicate single sensor failure, and sensor address is provided in Event Data 3 byte, or it may indicate lack of reading from all sensors. In this case Event Data 3 byte is set to zero.</p> <p>First occurrence of not acknowledged event will be retransmitted no faster than every 300 milliseconds.</p> <p><b>Note</b> – In DNM SKU this message is logged to SEL.</p> |

| Net Function = S/E (04h) |   |  |   |
|--------------------------|---|--|---|
| Code                     | Command   | Request, Response Data   | Description   |
| 16h                      | <b>Alert Immediate Message</b><br><br><b>Node Manager Operational Capabilities Change</b> | <b>Request</b><br><br>Bytes 1:3 – As set in the “Set Node Manager Alert Destination” IPMI command.<br><br>Byte 4 – Generator ID<br>[7:1] – 7-bit IPMB address of the Firmware.<br>[0] = 0b – Generator ID is IPMB address.<br><br>Byte 5 – EvMRev<br>=04h – IPMI2.0 format.<br><br>Byte 6 – Sensor Type<br>=DCh – OEM.<br><br>Byte 7 – Sensor Number<br>=26 – Node Manager Operational Capabilities sensor.<br><br>Byte 8 – Event Dir   Event Type<br>[0:6] – Event Type<br>=74h – OEM.<br>[7] – Event Dir<br>=0 – Assertion Event.<br>=1 – Deassertion Event.<br><br>Byte 9 – Event Data 1<br>[0:3] – Current state of Operational Capabilities. The same value is also returned by the Get Sensor Reading command invoked for Operational Capabilities sensor.<br>[0] – Policy interface capability<br>=0 – Not Available.<br>=1 – Available.<br>[1] – Monitoring capability<br>=0 – Not Available.<br>=1 – Available.<br>[2] – Power limiting capability<br>=0 – Not Available.<br>=1 – Available.<br>[4:7] – Reserved. | This message provides a run-time error indication about Node Manager's operational capabilities. This applies to all domains.<br><br>Assertion and deassertion of these events are supported.<br><br>Refer to section 3.4.5 for detailed description of the conditions triggering changes of the capabilities bitmask.<br><br>Current value of not acknowledged capabilities sensor will be retransmitted no faster than every 300 milliseconds.<br><br><b>Note</b> – In DNM SKU this message is logged to SEL. |
|                          |   | <b>Response</b><br><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br>Others – Error.  |   |
| 16h                      | <b>Alert Immediate Message</b><br><br><b>Node manager Alert Threshold Exceeded</b>        | <b>Request</b><br><br>Bytes 1:3 – As set in the “Set Node Manager Alert Destination” IPMI command.<br><br>Byte 4 – Generator ID<br>[7:1] – 7-bit IPMB address of the Firmware.<br>[0] = 0b – Generator ID is IPMB address.<br><br>Byte 5 – EvMRev<br>=04h – IPMI2.0 format.<br><br>Byte 6 – Sensor Type<br>=DCh – OEM.<br><br>Byte 7 – Sensor Number<br>=27 – Node Manager Event Sensor.<br><br>Byte 8 – Event Dir   Event Type<br>[0:6] – Event Type  | Policy Correction Time Exceeded Event will be sent each time when maintained policy power limit is exceeded over Correction Time Limit.<br><br>First occurrence of not acknowledged event will be retransmitted no faster than every 300 milliseconds.<br><br>First occurrence of Threshold exceeded event assertion/ deassertion will be retransmitted no faster   |



| Net Function = S/E (04h) |   |  |   |
|--------------------------|---|--|---|
| Code                     | Command   | Request, Response Data   | Description   |
|                          |   | <p>=72h – OEM.<br/>           [7] – Event Dir<br/>           =0 – Assertion Event.<br/>           =1 – Deassertion Event.</p> <p>Byte 9 – Event Data 1<br/>           [0:1] – Threshold Number<br/>           =0 – 2 – Threshold index.<br/>           [2] – Reserved.<br/>           [3] – Node Manager Policy Event<br/>           =0 – Threshold exceeded.<br/>           =1 – Policy Correction Time Exceeded – policy did not meet the contract for the defined policy. The policy will continue to limit the power or shutdown the platform based on the defined policy action.<br/>           [4:5] =10b – OEM code in byte 3.<br/>           [6:7] =10b – OEM code in byte 2.</p> <p>Byte 10 – Event Data 2<br/>           [0:3] – Domain ID (Identifies the domain that this Node Manager policy applies to)<br/>           =00h – Entire platform.<br/>           =01h – CPU subsystem.<br/>           =02h – Memory subsystem.<br/>           Others – Reserved<br/>           [4:7] – Reserved.</p> <p>Byte 11 – Event Data 3<br/>           – &lt;Policy ID&gt;.</p> <p><b>Response</b><br/>           Byte 1 – Completion Code<br/>           =00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>           Others – Error.</p> | <p>than every 300 milliseconds.</p> <p><b>Note</b> – In DNM SKU this message is logged to SEL.</p>  |
| 02h                      | <b>Platform Event Message</b><br><b>Node Manager PSU Status Event</b> | <p><b>Request</b></p> <p>Byte 1 – EvMRev<br/>           =04h – IPMI2.0 format.</p> <p>Byte 2 – Sensor Type<br/>           =08h – Power Supply.</p> <p>Byte 3 – Sensor Number<br/>           =102 – 109 – PSU Status Sensor.</p> <p>Byte 4 – Event Dir   Event Type<br/>           [0:6] – Event Type<br/>           =6Fh – Sensor Specific.<br/>           [7] – Event Dir<br/>           =0 – Assertion Event.<br/>           =1 – Deassertion Event.</p> <p>Byte 5 – Event Data 1<br/>           [7:6] =00b – Unspecified byte 2.<br/>           [5:4] =00b – Unspecified byte 3.<br/>           [3:0] – Offset from Event/Reading Code for discrete event state. (Assertion and Deassertion refer to Byte 4, bit [7] - Event Dir.)=0 – Presence detected<br/>           Assertion: Power supply module is present.<br/>           Deassertion: Power supply module is not present.<br/>           =1 – Power supply failure<br/>           Assertion: No Power Supply failure detected.<br/>           Deassertion: One of faults detected: IOOut OC Fault, Pout OP fault (STATUS_IOUT), OT fault</p>   | <p>Event will be sent each time when monitored PSU Status changes.</p> <p>First occurrence of not acknowledged event will be retransmitted no faster than every 300 milliseconds.</p> <p>“Re-arm Sensor Events” IPMI command on bit [5] will reset CLST events (over-current and over-temperature) in the PSU. This allows in the manual CLST mode to reset the platform throttling by the BMC when the corrective action on the platform was performed.</p> <p><b>Note</b> – In DNM SKU this event is logged to SEL.</p> |

| Net Function = S/E (04h) |         |  |             |
|--------------------------|---------|--|-------------|
| Code                     | Command | Request, Response Data   | Description |
|                          |         | <p>(STATUS_TEMPERATURE), Fan Fault (STATUS_FANS_1_2).<br/>           =2 – Predictive failure<br/>           Assertion: No Power Supply warning detected.<br/>           Deassertion: One of warnings detected: lout OC warning, Pout OP warning (STATUS_IOUT), Vin UV warning, Iin over current warning, Pin over power warning (STATUS_INPUT), OT warning (STATUS_TEMPERATURE), Fan 1 warning, Fan 2 warning (STATUS_FANS_1_2).<br/>           =3 – AC Lost<br/>           Assertion: AC Power OK.<br/>           Deassertion: No AC power input to a power supply module (Unit Off For Low Input Voltage bit in STATUS_INPUT is set).<br/>           =4 – Configuration Error<br/>           Assertion: PSU configuration OK.<br/>           Deassertion: – Not supported PSU, when other information cannot be read (PSU not responded for PMBus requests).<br/>           =5 – PSU optimization feature triggered. Throttling to the system power applied to prevent the system failure. PSU optimization could be activated by one of the following warnings lout OC warning (STATUS_IOUT), Vin UV warning (STATUS_INPUT) or OT warning (STATUS_TEMPERATURE)<br/>           Assertion: PSU optimization feature activated to prevent the system failure.<br/>           Deassertion: No Power Supply warning detected.<br/>           =6--15 – Reserved.</p> <p>Byte 6 – Event Data 2<br/>           [7:4] - Optional offset from 'Severity' Event/Reading Code. (0Fh if unspecified).<br/>           [3:0] - Optional offset from Event/Reading Type Code for previous discrete event state. (0Fh if unspecified.)</p> <p>Byte 7 – Event Data 3<br/>           Optional OEM code. FFh or not present if unspecified.</p> |             |
|                          |         | <p><b>Response</b></p> <p>Byte 1 – Completion Code<br/>           =00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br/>           Others – Error.</p>   |             |

## 3.5 Error Conditions

There may be situations when the Node Manager is not available to respond to IPMI interface. This may happen when the management controller is not powered by stand-by power and the system is powered down or when there is a firmware update in progress. In addition, when using bridged commands, the BMC may not respond to the IPMI commands when it is unavailable for reasons such as flash update.

The external management SW needs to be aware of the fact there may be situations like these resulting in scenarios when responses to bridged command may not arrive or alerts





may not be generated. The external management software needs to be designed appropriately to account for these.

§

## 4. BMC IPMI Interface

This section contains IPMI commands and sensor devices which shall be provided by BMC in order to enable Node Manager firmware.

According to [Addr] specification ME expects BMC at address 10h in 7-bit format (20h in 8-bit). This address is fixed.

To support initialization of Management Engine owned sensors based on associated SDRs, OEM BMC must be able to use both the slave address and BMC channel number fields of the type1 and type2 SDRs for the purpose of directing the IPMI sensor commands to the management engine.

For proper ME initialization BMC should send "Set Event Receiver" IPMI command. If BMC does not send "Set Event Receiver" command within this time after Host reset or startup Node Manager will activate "Time After Host Reset" policy.

### 4.1 IPMI Device "Global" Commands

| Net Function = Chassis (00h) |                         |   |   |
|------------------------------|-------------------------|---|---|
| Code                         | Command                 | Request, Response Data  | Description   |
| 02h                          | Chassis Control Command | <b>Request</b><br>Byte 1<br>[7:4] – Reserved. Write as 0000b.<br>[3:0] – Control Command<br>=0 – Power Down.<br>=5 – Soft Shutdown (via ACPI) (optional). | This is standard IPMI 2.0 command.<br><br><b>Note</b> – NM Firmware will use Soft Shutdown (optional) and Power Down. |
|                              |                         | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).                            |   |

### 4.2 Sensor Device Commands

| Net Function = S/E (04h) |                        |                                    |  |
|--------------------------|------------------------|------------------------------------|--|
| Code                     | Command                | Request, Response Data             | Description  |
| 02h                      | Platform Event Message | For command description see [IPMI] | This is general format of Platform Event Message. Detailed description of all messages generated by Node Manager firmware can be found in sections 2.6.4 and 3.4.7 |



## 4.3 Alert Immediate Command

In order to properly forward the alerts from the Node Manager to the remote management software BMC should support the Alert Immediate Command.

| Net Function = S/E (04h) |                 |                                    |  |
|--------------------------|-----------------|------------------------------------|--|
| Code                     | Command         | Request, Response Data             | Description  |
| 16h                      | Alert Immediate | For command description see [IPMI] | <p>This is standard IPMI 2.0 command.</p> <p><b>Note</b> – Alert Immediate command used by ME and supported by BMC must be compliant with IPMI Errata E358.</p> <p>It is expected that BMC will send to the remote console at least the contents of the bytes 5 to 11 of the IPMI message.</p> |

## 4.4 OEM Commands Implemented by BMC

Depending on the firmware factory settings SHOULD be supplied by the BMC for the associated Intel ME services to work correctly. If the referenced services are not activated, the BMC does not need to provide the sensor(s) and/or OEM commands.

### 4.4.1 Power Consumption Readings

If the firmware is configured to use BMC for power readings the sensors return Platform Power consumption. Single power reading, subtotal of per-rail readings from one PSU as well as total of single run of readings across all attached PSU's cannot exceed 32767 Watt. Such a power reading will be treated as a reading failure. That rule applies to any power reading source.

Depending on the configuration Firmware may use one of the following sources for platform power consumption readings:

1. PMBUS compliant PSU or voltage regulators attached directly to PCH SMLINK1 (recommended) or PCH SMLINK0. In that case there is no need to implement any support in the BMC.
2. BMC sensor read by management engine Firmware using OEM command implemented by the BMC. In that case BMC should implement OEM command to return the sensor value on the query from management engine Firmware. BMC should average the power over 1 second to allow management engine Firmware to read the power two to ten times per second<sup>3</sup>. This type of power reading allow for non-PMBUS compliant PSU or voltage regulator support. Additionally, the OEM command code may be configured using Factory presets:

---

<sup>3</sup> The frequency of power readings of 10, 5 or 2 per second should be defined by the OEM in the factory presets. For the fastest correction time of Node Manager Policy that is below one second the frequency of 10 readings per second should be supported by BMC.

| OEM command                        | Description  | Encoding  |
|------------------------------------|--|---|
| XXh<br>(default E2h <sup>4</sup> ) | 'OEM Get Reading' with type "Platform Power Consumption" | The value of the reading is encoded on 16-bit encoding 2s-complement signed integer. Values below 0 are ignored and treated as a power reading failure. |

#### 4.4.2 Inlet Air Temperature readings

Depending on the configuration management engine Firmware may use one of the following sources for inlet temperature readings:

1. SST sensors attached directly to PCH. In that case there is no need to implement any support in the BMC.
2. BMC sensor read by management engine Firmware using OEM command implemented by the BMC. In that case BMC should implement OEM command to return the sensor value on the query from management engine Firmware. Additionally, the OEM command code may be configured using Factory presets:

| OEM command                     | Description   | Encoding   |
|---------------------------------|---|--|
| XXh (default E2h <sup>5</sup> ) | 'OEM Get Reading' with type "Inlet Air Temperature" | The value of the sensor returned in "Get Inlet Air Temperature" is encoded on 16-bit encoding 2s-complement signed integer. Values below -128 degrees centigrade and above +127 degrees centigrade are ignored and treated as a temperature reading failure. |

*Note: OEM command code value may be different for Inlet Air Temperature readings and Power Consumption reading)*

3. With "Get Sensor Reading" IPMI command management engine Firmware expects the sensor reading to be temperature in range of 0 to 100 centigrade. Sensor readings outside of the specified range are ignored and treated as a temperature reading failure.
4. Inlet Air Temperature reading can be disabled. In that case the inlet temperature statistics and Node Manager Policies using temperature trigger will be not available.

#### 4.4.3 OEM Management Engine Power State Change

Optionally, instead of the event SPS Firmware may send an IPMI command with information as defined above. Using Factory presets OEM may choose to use an event or OEM command.

<sup>4</sup> Note: OEM command code value may be different for Power Consumption reading)

<sup>5</sup> Note: OEM command code value may be different for Inlet Air Temperature readings)



| OEM command          | Description                 | Encoding   |
|----------------------|-----------------------------|--|
| XXh<br>(default E3h) | "OEM ME Power State Change" | The values returned by the command:<br>00h – Transition to Running – ME is started.<br>02h – Transition to Power Off – ME is to be powered down. |

#### 4.4.4 OEM Command Definition

| Net Function = SDK General Application (0x30) |                           |   |   |
|---|---------------------------|---|---|
| Code  | Command                   | Request, Response Data  | Description   |
| E2h   | OEM Get Reading           | <b>Request</b><br><br>Byte 1 – Domain Id/Reading Type<br>[0:3] – Domain Id. Currently only domain 0 will be queried. Other domains are queried internally by ME using.<br>[4:7] – Reading Type<br>=00h – Platform Power Consumption. For platform power consumption the Domain Id will be set to 0. Values from 1 to 15 could be used to address power rails. Per-rail readings are optional and Firmware needs to be preconfigured in the factory settings.<br>=01h – Inlet Air Temperature. For Inlet Air Temperature the Domain Id will be set to 0.<br>=03h – 0Fh – Reserved.<br>Byte 2:3 – Reserved. Write as 0000h.   | This command is optional and may be implemented by the BMC. |
|   |                           | <b>Response</b><br><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11).<br><br>Byte 2 – Reading Type<br>[0:3] – Domain Id copied from request (depending on the Byte 2 identifies the processor which should be queried for I <sub>CC_TDC</sub> or power rail).<br>[4:7] – Reading Type<br>=00h – Platform Power Consumption in [Watts]. Values below 0 are ignored and treated as a power reading failure.<br>=01h – Inlet Air Temperature in degrees centigrade. Values below -128 degrees centigrade and above +127 degrees centigrade will be ignored and treated as a temperature reading failure.<br>=02h – I <sub>CC_TDC</sub> reading from PECL. Values below 0 and above 255 will be ignored and treated as a PECL reading failure.<br>=03h – FFh – Reserved.<br><br>Byte 3:4 – Reading value 16-bit encoding 2s-complement signed integer. |   |
| E3h   | OEM ME Power State Change | <b>Request</b><br><br>Byte 1 – Power State<br>=00h – Transition to Running – ME is started.<br>=02h – Transition to Power Off – ME is to be   | This command is optional and may be implemented by the BMC. |

| Net Function = SDK General Application (0x30) |         |  |             |
|---|---------|--|-------------|
| Code  | Command | Request, Response Data   | Description |
|   |         | powered down.  |             |
|   |         | <b>Response</b><br>Byte 1 – Completion Code<br>=00h – Success (Remaining standard Completion Codes are shown in section 2.11). |             |

#### 4.4.5 Summary of Options

| Summary of Implementation Options          |                 |   |
|--|-----------------|---|
| Functionality                              | Type            | Description   |
| 4.4.1 Power Consumption Readings           | Factory presets | Power readings via "OEM Get Reading" command are only needed if no PMBUS PSU is directly connected to the SMLINK.<br><i>Default in the Factory Presets is to use PMBUS PSU.</i> |
| 4.4.2 Inlet Air Temperature readings       | Factory presets | Required if inlet temperature statistics and Node Manager Policies using temperature trigger will be exposed by the platform.   |
| 4.4.3 Management Engine Power State Change | Factory presets | ME Firmware is able to send notification about the power state change using standard IPMI sensor or via OEM command.<br><i>Default is to use OEM command.</i>                   |

### 4.5 BMC requirements for Node Manager Discovery

The following discovery mechanism should be implemented by the BMC in order to allow external management software to properly configure the communication channel between Node Manager and the external management software.

For command routing purposes, the external SW needs to know which microcontroller implements the Node Manager functionality. Additionally, the external SW needs to know the IPMI sensor numbers associated with each Node Manager sensor of interest. This information is provided via a Node Manager OEM SDR<sup>6</sup>.

The first step in the Node Manager discovery process is for the software to search the SDR repository for this OEM. If the Device Slave Address found in this SDR matches that of the BMC (20h), then all of the NM-related IPMI commands are sent directly to the BMC. Otherwise, standard IPMI bridging is used to send these commands to the satellite Node Manager controller<sup>7</sup>. The SW application uses the sensor information in this SDR to

<sup>6</sup> BMC should not expose the NM discovery SDR to the external console if NM SKU is disabled in the firmware.

<sup>7</sup> Note: The privilege level to access to the ME SMLINK channel should be restricted to allow only the Admin level.



comprehend the mapping of the sensor numbers to the NM sensors of interest. Additional sensor information can be retrieved by searching for associated type1, type2, or type3 SDRs for the specific sensors.

OEM SDR records are of type C0h. They contain a manufacturer ID and OEM data in the record body. Intel OEM SDR records also have a sub-type field in them as the first byte of the OEM data indicates the type of record following.

| Byte<br>(beginning<br>after SDR<br>record<br>header) |      | Name   | Description   |
|--|------|--|---|
| 0:2  |      | OEM ID                                       | Intel Manufacturer ID = 000157h   |
| 3  |      | Record Subtype                               | NM Discovery = 0Dh  |
| NM Record  | 4    | Version number of this record subtype        | =01h – for the version specified in this document.  |
|  | 5    | NM Device Slave Address                      | [7:1] – 7-bit I2C Slave Address of NM controller on channel.<br>[0] – Reserved.   |
|  | 6    | Channel Number / Sensor Owner LUN            | [7:4] – Channel number for the channel that the NM management controller is on. Use 0h if the primary BMC is the NM controller.<br>[3:2] – Reserved. Write as 00b.<br>[1:0] – Sensor owner LUN used for accessing all NM sensors enumerated in this record. |
|  | 7    | Node Manager Health Event sensor             | =25   |
|  | 8    | NM Exception Event sensor                    | =24   |
|  | 9 NM | Operational Capabilities sensor              | =26   |
|  | 10   | Node manager Alert Threshold Exceeded sensor | =27   |

## 4.6 Alerts

Events mentioned in section 0 that are not marked as ‘Alert Immediate’ are sent as IPMI alerts to external software using PEF/PET. The BMC will perform the filtering based on filters set up by the external software.

In order to avoid excessive logging into the SEL due to NM “threshold exceeded” and “NM Health” events, a mechanism is provided to send PET alerts without use the PEF mechanism. These use the ‘Alert Immediate’ mechanism. This requires that the external SW application provide the NM with the alert destination and alert string information needed to

properly form and send the alert. The external SW must first properly configure the alert destination and string in the BMC LAN configuration using standard IPMI commands, then provide the associated selectors to the BMC using the “Set Node Manager Alert Destination” OEM command. Section 3.1 contains the description of this OEM command.

Setting alert destination using “Set Node Manager Alert Destination” will cause all events marked “Immediate Alert” in the table of events to be routed to that destination as PET alerts. It is not possible to have some types of events sent to one destination and others to another.

No provision will be accommodated at this time for an in-band agent to receive alerts. The identification of an in-band alerting method if required will be defined at a later time.

When Node Manager needs to generate multiple events at the same time, Node Manager will not generate a new Platform Event and Alert Immediate command until BMC does not send positive acknowledgment for the previous one (i.e. the Completion Code must be zero).

## 4.7 Command Passing via BMC

If Node Manager is implemented on BMC, then external SW will send the commands directly addressed to BMC (no bridging).

If Node Manager is implemented on a controller other than BMC, then external SW will send ‘bridged’ IPMI commands to BMC. Encapsulated bridged IPMI commands must follow the format for the channel that is being bridged to.

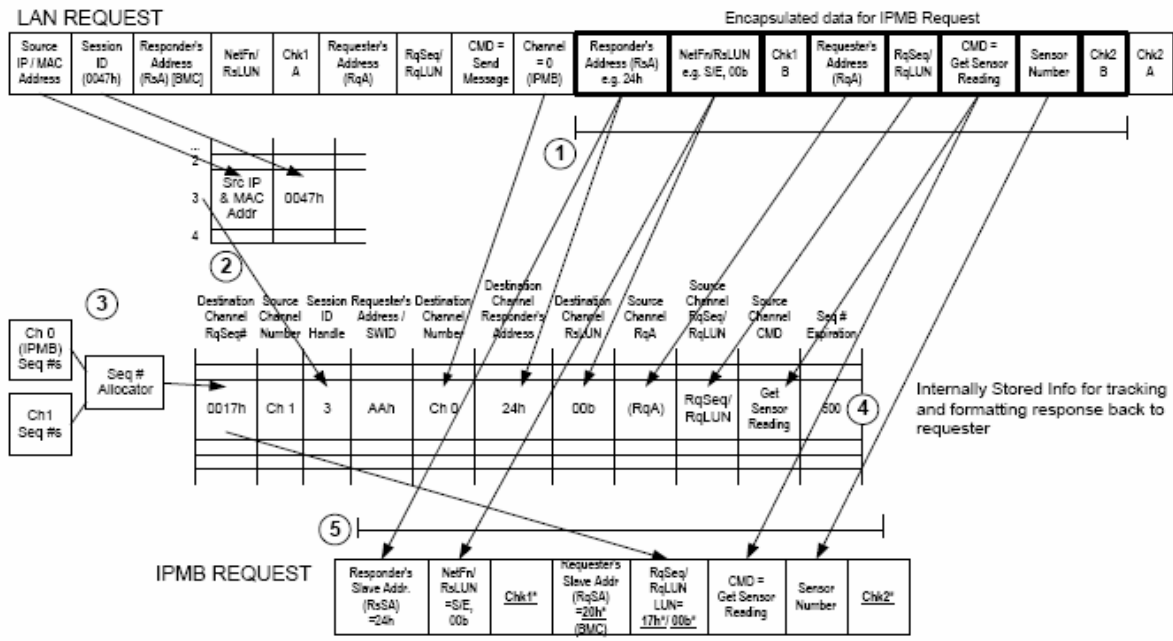
This will be in the form an IPMI packet encapsulated in another packet. BMC will need to examine the payload of the packet addressed to it, construct the proper IPMB packet, forward it to Node Manager and return the response from Node Manager to the sender. Please refer to IPMI 2.0 specification for details.

Because of the potentially performance-limiting and system shut-down effects of some of the commands that can be bridged, the BMC shall restrict the IPMI command bridging to the Administrator privilege level.





Figure 2 Example IPMI Command Bridging from LAN



For the purpose of constructing a bridged IPMI command, external software would need to know the below parameters to construct the proper IPMI packet:

- 'Responder's address'
- Destination channel #
- Network Function/LUN
- Channel protocol

It is the responsibility of the BMC to let external software know of the 'Responder's address'. This will be done as part of Node Manager discovery procedure. The responder in this case will be the actual management controller implementing Node Manager (management engine, BMC).

In addition, the actual medium to communicate the bridged packet on IPMB is also needed. This is done as part of its initial discovery by query of the BMC for channel protocol using the IPMI command "Get Channel Info". For example, if Node Manager is implemented in the management engine (satellite controller), BMC needs to forward the packet over IPMB. This information is used by external software to construct the encapsulated packet appropriate to the medium.

The length of the supported IPMI frames is extended to 80 bytes so up to 68 bytes of payload can be passed in one IPMI request.

## 4.8 IPMB Reset Scenarios

BMC and management engine Firmware share the same SMBUS link. The management engine Firmware will perform SMBUS hang recovery only when the hang happens during a transaction initiated by the management engine. The recovery method is by „bit banging” 9 clock pulses. The ME FW does not reset SMLINK on startup.

§



## 5. IPMI OEM Command Summary

The table below summarizes the OEM commands used by NM firmware for management engine-BMC communication. All the commands use Network Function 2Eh/2Fh and LUN 00b with Intel IANA Number 000157h.

| NET Function                  | Command Code | Command Name  | Implemented by | SKU                        |
|-------------------------------|--------------|---|----------------|----------------------------|
| 2Eh + 000157h                 | A0-A8h       | Online Firmware Update Commands                                     | ME             | SiEn/NM                    |
| 2Eh + 000157h                 | C0-CDh       | Node Manager Commands   | ME             | NM/DNM                     |
| 2Eh + 000157h                 | CEh          | Set Node Manager Alert Destination                                  | ME             | NM                         |
| 2Eh + 000157h                 | CFh          | Get Node Manager Alert Destination                                  | ME             | NM                         |
| 2Eh + 000157h                 | D0h          | Set Total Power Budget  | ME             | NM                         |
| 2Eh + 000157h                 | D1h          | Get Total Power Budget  | ME             | NM                         |
| 2Eh + 000157h                 | D2h          | Set Max Allowed CPU P-state/T-state                                 | ME             | NM                         |
| 2Eh + 000157h                 | D3h          | Get Max Allowed CPU P-state/T-state                                 | ME             | NM                         |
| 2Eh + 000157h                 | D4h          | Get Number Of P-states/T-states                                     | ME             | NM                         |
| 2Eh + 000157h                 | D7h          | Set PSU Configuration   | ME             | NM/DNM                     |
| 2Eh + 000157h                 | D8h          | Get PSU Configuration   | ME             | NM/DNM                     |
| 2Eh + 000157h                 | D9h          | Send Raw PMBus  | ME             | SiEn/NM                    |
| 2Eh + 000157h                 | DCh          | Set ME Power State  | ME             | SiEn/NM                    |
| 2Eh + 000157h                 | 40h-5Fh      | PECI Proxy and PECI related commands in NM 2.0                      | ME/BMC         | SiEn/NM                    |
| 2Eh + 000157h                 | DFh          | Force ME Recovery   | ME             | SiEn/NM/DM/DNM             |
| 2Eh + 000157h                 | E0h          | Get ME Factory Presets Signature                                    | ME             | SiEn/NM/DM/DNM             |
| 2Eh + 000157h                 | EAh          | Get Host CPU Data   | ME             | NM/DNM                     |
| 2Ch (DCGRP) + DCh             | 01-05h       | DCMI Power management Commands                                      | ME             | DNM/DM/On Romley/Denlow NM |
| 30h (SDK General Application) | A0-A7h       | Online Firmware Update Commands (backward compatibility with NM1.5) | ME             | SiEn/NM                    |
| 30h (SDK General Application) | XXh          | OEM Management Engine Power State Change                            | ME             | SiEn/NM                    |

For a list of commands supported by recovery boot loader see chapter: 2.8 IPMI Commands Supported by Recovery Boot Loader.

### S





## 6. Appendix NM2.0 Platform Changes

This appendix describes changes between Intel® Intelligent Power Technology Node Manager External Interface Specification using IPMI for NM1.5 and NM2.0.

### 6.1 IPMI Commands Not Supported on NM2.0 Platform

The table below presents IPMI commands not supported by Node Manager 2.0 firmware in comparison with NM1.5 implementations on Tylersburg-EP and Boxboro-EX. Node Manager responds to the commands with C1h – Invalid Command Completion Code.

| NET Function               | Command Code | Command Name            | Implemented by ME or BMC |
|----------------------------|--------------|-------------------------|--------------------------|
| 30h SDK General App        | DAh          | Set Cooling Coefficient | ME                       |
| 2Eh + 000157h (Intel IANA) | D6h          | Set Host CPU data       | ME                       |

### 6.2 IPMI Sensors Implemented by Node Manager FW

List of sensors implemented by the ME Firmware has changed see Chapter 2.10 IPMI OEM PECI Proxy Sensors for details.

Average CPU Temperature sensors are not implemented.

#### 6.2.1 Event Generation Control

The below table summarizes how event generation can be enabled/disabled for each sensor implemented by NM Firmware.

| Sensor Types          | IPMI command                                       | Event Generation Control Methods   |
|-----------------------|--|--|
| PECI Proxy Sensors    | Platform Event                                     | Event generation can be enabled/disabled using the IPMI commands: <ul style="list-style-type: none"><li>Set Event Receiver.</li><li>Set Event Enable – both per sensor and per threshold control support.</li></ul> Default event enable flags can be set using Flash Image Tool for each threshold.                               |
| ME Power State Sensor | Platform Event or OEM command (settable using FIT) | When the notification is sent using Platform Event message, event generation can be enabled/disabled using the IPMI command: <ul style="list-style-type: none"><li>Set Event Receiver.</li></ul> If the event is sent using OEM message, the notification is sent always to address 20h regardless of the Event Receiver settings. |

| Sensor Types   | IPMI command    | Event Generation Control Methods  |
|--|-----------------|---|
| Node Manager Exception Sensor  | Platform Event  | Event generation can be enabled/disabled using the IPMI command: <ul style="list-style-type: none"> <li>Set Event Receiver</li> </ul>   |
| Node Manager Operational Capabilities Sensor                               | Alert Immediate | Event generation can be enabled/disabled using the IPMI commands: <ul style="list-style-type: none"> <li>Set Node Manager Alert Destination – clearing alert destination parameters disables generation of events sent using Alert Immediate IPMI command.</li> <li>Set Event Enable – Scanning Enabled bit disables/enables the whole sensor. Changing per sensor Event Enable flag is supported.</li> </ul> |
| Node Manager Health Sensor<br>Node Manager Alert Threshold Exceeded Sensor | Alert Immediate | Event generation can be enabled/disabled using the IPMI command: <ul style="list-style-type: none"> <li>Set Node Manager Alert Destination – clearing alert destination parameters disables generation of events sent using Alert Immediate IPMI command.</li> </ul>  |
| NM FW Health Sensor  | Platform Event  | Not possible to enable/disable event generation from this sensor.   |
| PSU Status Sensors   | Platform Event  | Event generation can be enabled/disabled using the IPMI commands: <ul style="list-style-type: none"> <li>Set Event Receiver.</li> <li>Set Event Enable – both per sensor and per threshold control support.</li> </ul> Default event enable flags can be set using Flash Image Tool for each threshold.   |

### 6.3 IPMI Platform Event Messages Generated by NM FW

ME Firmware will try to deliver the all the events by resending them if event reception is not acknowledged by the BMC.

| Event   | Sensor Type           | Event Dir                        | Event Type         | Event command   |
|---|-----------------------|----------------------------------|--------------------|-----------------|
| <b>Memory Throttling Status</b>                           | 0Ch – Memory          | 0 – Assertion<br>1 – Deassertion | 01h – Threshold    | Platform Event  |
| <b>ME Power State</b>                                     | 16h – Microcontroller | 0 – Assertion                    | 0Ah – Availability | Platform Event  |
| <b>Node Manager Exception Event</b>                       | DCh – OEM             | 0 – Assertion                    | 72h – OEM          | Platform Event  |
| <b>Node manager Alert Threshold Exceeded sensor</b>       | DCh – OEM             | 0 – Assertion<br>1 – Deassertion | 72h – OEM          | Alert Immediate |
| <b>Node Manager Health Event</b>                          | DCh – OEM             | 0 – Assertion                    | 73h – OEM          | Alert Immediate |
| <b>Node Manager Operational Capabilities Change Event</b> | DCh – OEM             | 0 – Assertion<br>1 – Deassertion | 74h – OEM          | Alert Immediate |



| Event                                  | Sensor Type | Event Dir                        | Event Type      | Event command  |
|--|-------------|----------------------------------|-----------------|----------------|
| ME Firmware Health Event               | DCh – OEM   | 0 – Assertion                    | 75h – OEM       | Platform Event |
| CPU Thermal Status                     | 07h – CPU   | 0 – Assertion<br>1 – Deassertion | 76h – OEM       | Platform Event |
| CPU Thermal Control Circuit Activation | 07h – CPU   | 0 – Assertion<br>1 – Deassertion | 01h – Threshold | Platform Event |

## §